

# Intro to Distributed Memory Parallel Computing with MPI

**Bruno Abreu, Ph.D.**

Research Scientist

[babreua@illinois.edu](mailto:babreua@illinois.edu)

October 24, 2023



National Center for  
Supercomputing Applications

UNIVERSITY OF ILLINOIS URBANA-CHAMPAIGN

# Outline

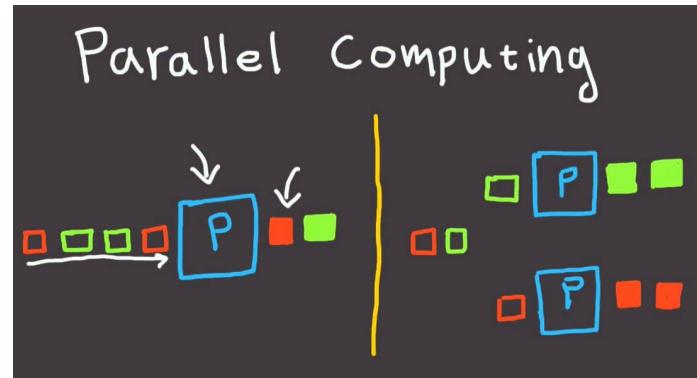
## Part A: Parallel Computing Quick Review (10 min)

- Task and Data parallelism
- Parallel computers
- Threads, processes and more



## Part B: MPI Intro (60 min)

- Core concepts
- Message passing
- Execution



## Part C: Exercises (45+ min)

Ex 1: Hello World

Ex 2: Domain Decomposition

Ex 3: Parallel Linear Regression

# Tools and resources we will cover

During the next 90 minutes, you will be directly or indirectly using/learning how to use:

- Git and GitHub
- SLURM
- OpenMP
  - Loop scheduling
  - Data declarations
  - Reduction clauses



- Parallel computing
  - Measuring performance
  - Bottlenecks and common issues



- Illinois Campus Cluster



# HPC Training Moodle



[www.hpc-training.org](http://www.hpc-training.org)

- Self-paced Tutorials
- Training Workshops
- Digital Badges

≡  NCSA HPC-Moodle

## Getting Started with MPI

[Home](#) / [My courses](#) / [Getting Started with MPI](#) / [MPI Overview](#) / [MPI Overview](#)

### MPI Overview

In this book, you will learn about some general concepts of the Message Passing Interface standard and how to identify applications that could benefit from MPI parallelization. We will present and briefly discuss:

- The core components of an MPI message
- The general MPI program structure
- MPI distributions
- Compilation and execution of MPI applications



HPC-Moodle

- Intro to OpenMP
- Intro to MPI
- Intro to Visualization
- Getting Started on the Illinois Campus Cluster
- Developing Effective Training Webinars

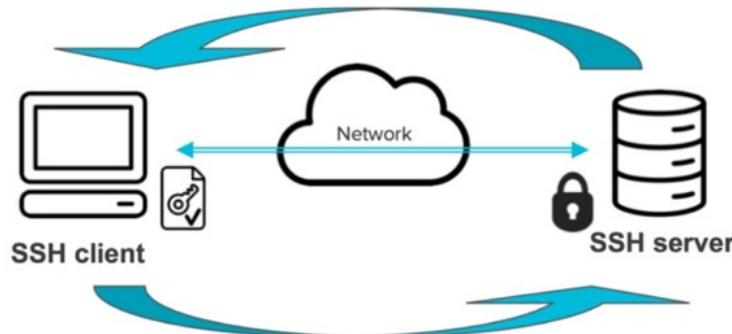
We develop live training sessions (like this one!) on top of existing self-paced courses. If you want to revisit today's contents and/or learn more about other HPC tools and frameworks, this is the place!



NCSA | NATIONAL CENTER FOR SUPERCOMPUTING APPLICATIONS

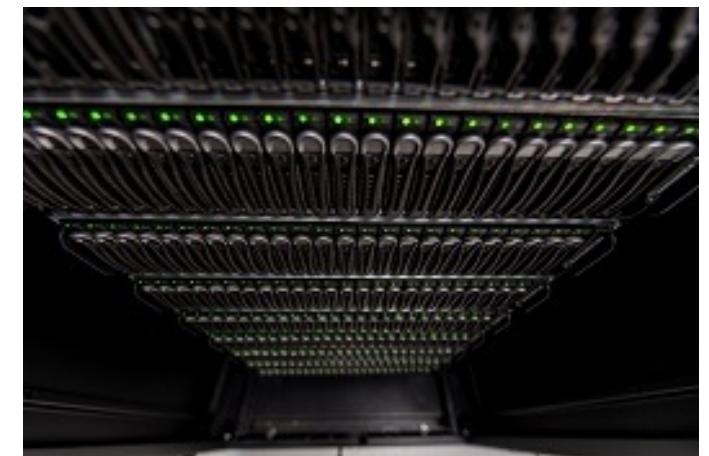


# Environment access



ssh

Illinois Campus Cluster Program  
AT THE UNIVERSITY OF ILLINOIS AT URBANA-CHAMPAIGN



"Terminal" on *Spotlight Search or Finder*



Start -> Windows System  
-> Command Prompt



Ctrl+Alt+T

- `ssh -l <NetID> cc-login1.campuscluster.illinois.edu`

# Cloning the repository

The exercises for today  
can be found here:

<https://github.com/babreu-ncsa/IntroToMPI/tree/uiuc-icc>

To copy these files to your ICC home directory, use:

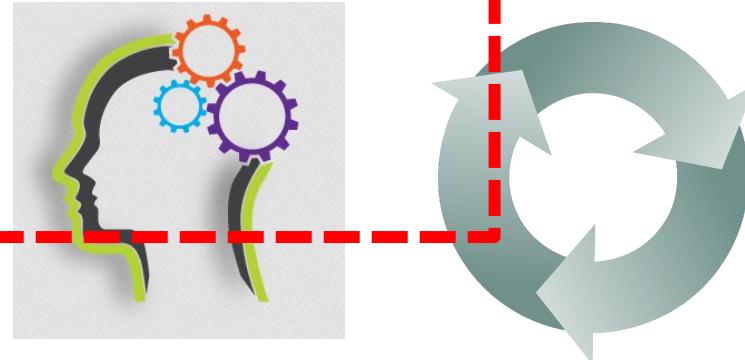
- module load git
- git clone --branch uiuc-icc --single-branch  
<https://github.com/babreu-ncsa/IntroToMPI.git>



# Outline

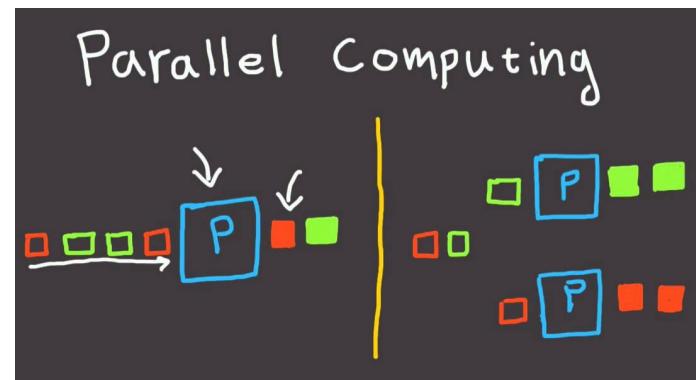
## Part A: Parallel Computing Quick Review (10 min)

- Task and Data parallelism
- Parallel computers
- Threads, processes and more



## Part B: MPI Intro (60 min)

- Core concepts
- Message passing
- Execution



## Part C: Exercises (45+ min)

Ex 1: Hello World

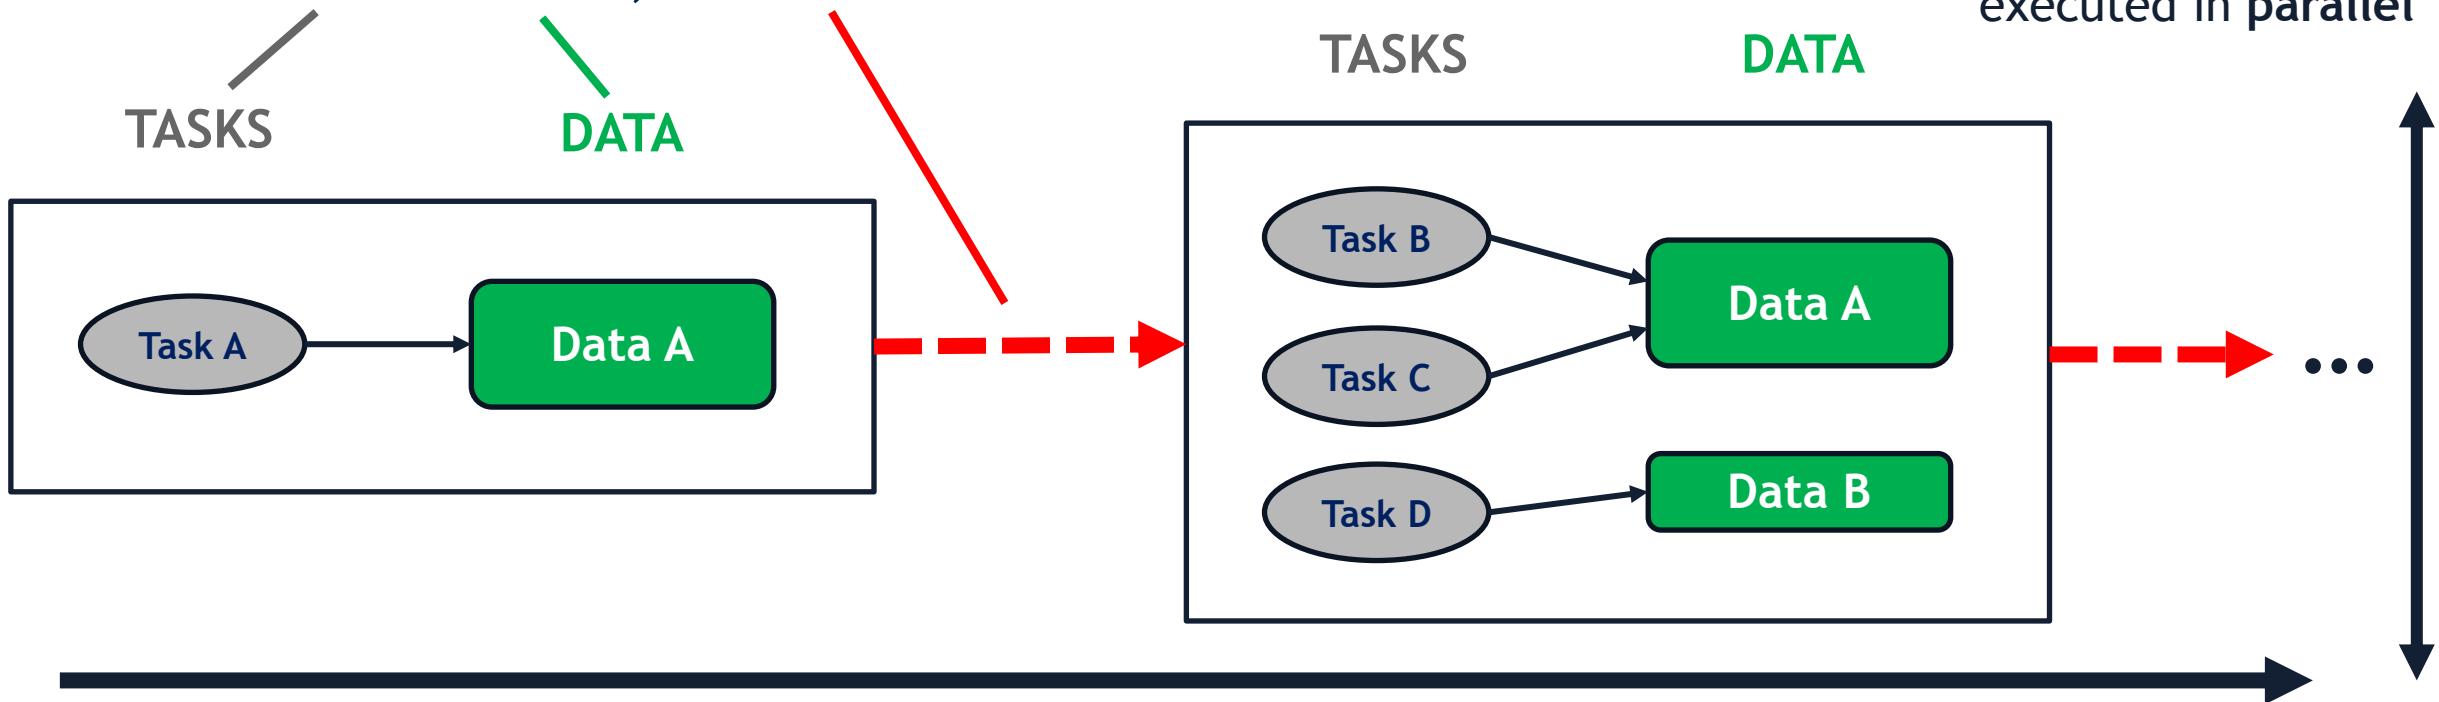
Ex 2: Domain Decomposition

Ex 3: Parallel Linear Regression

# Instruction blocks

A computer program is a sequence of **blocks of instructions**:

“Do this on **that**, and **then...**”

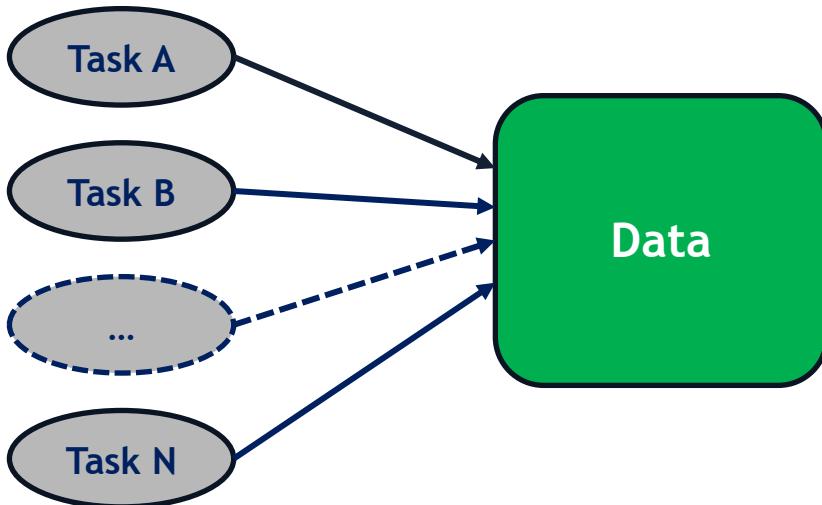


Different blocks are executed **sequentially**

# Data & Task Parallelism

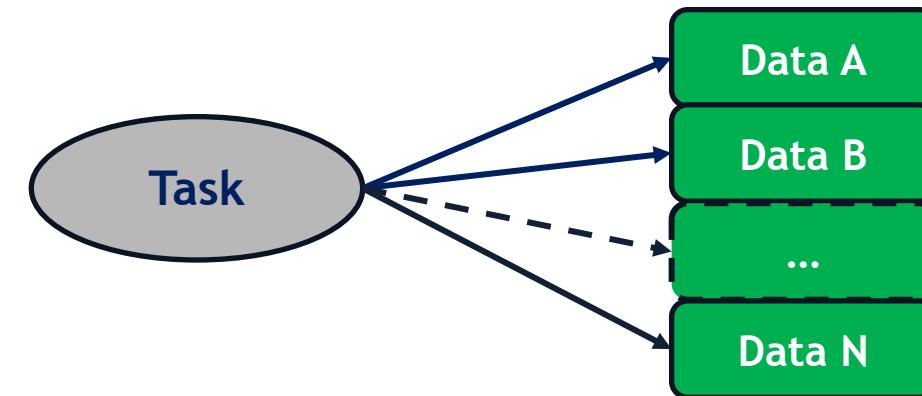
- An instruction is a pair (TASK, DATA)
- There are two “pure” ways of executing instructions in parallel:

Task Parallelism



Execute different tasks over the same data

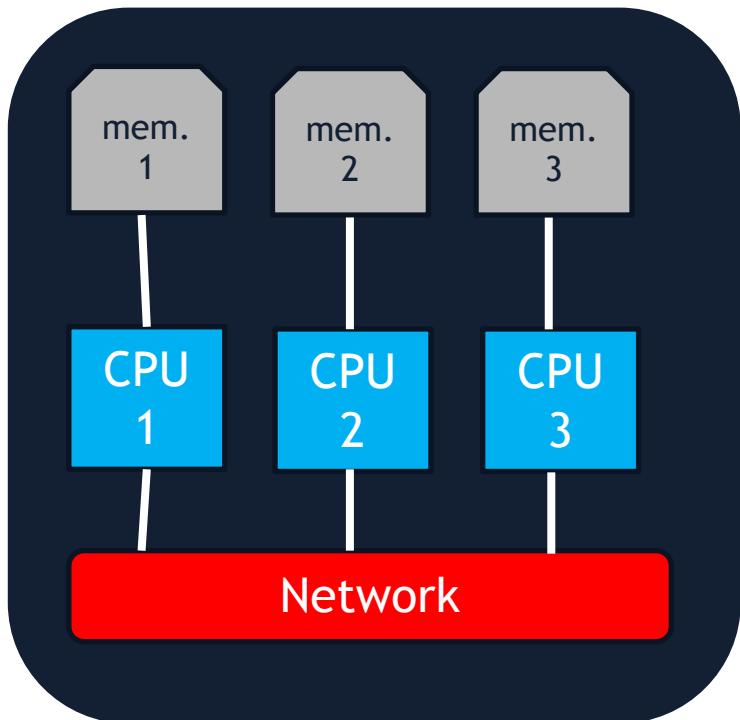
Data Parallelism



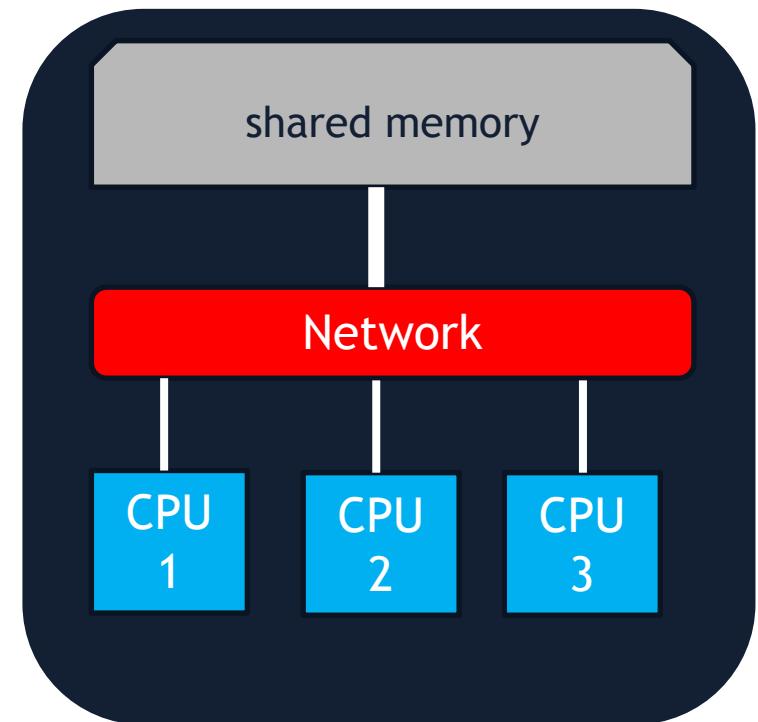
Execute the same task over different chunks of data

# Parallel computers: memory

Distributed Memory

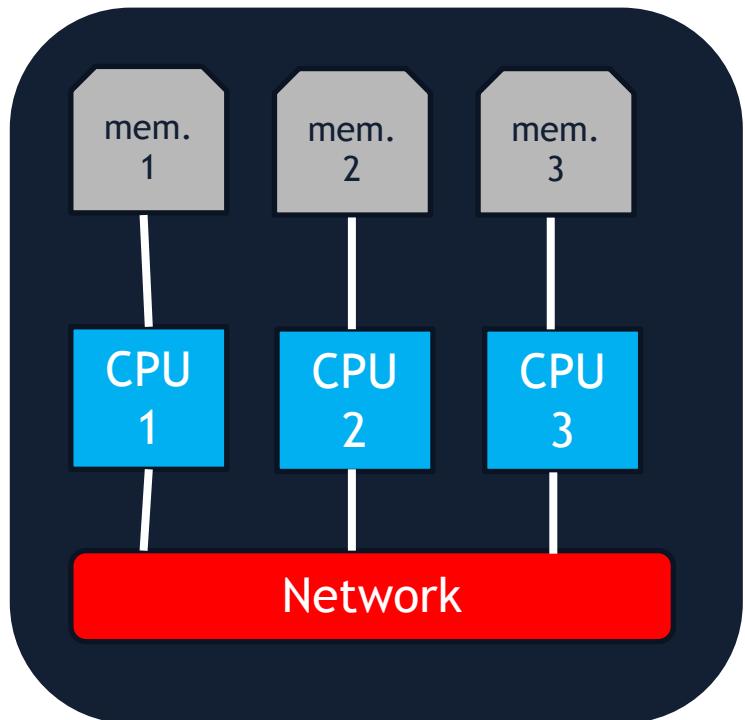


Shared Memory

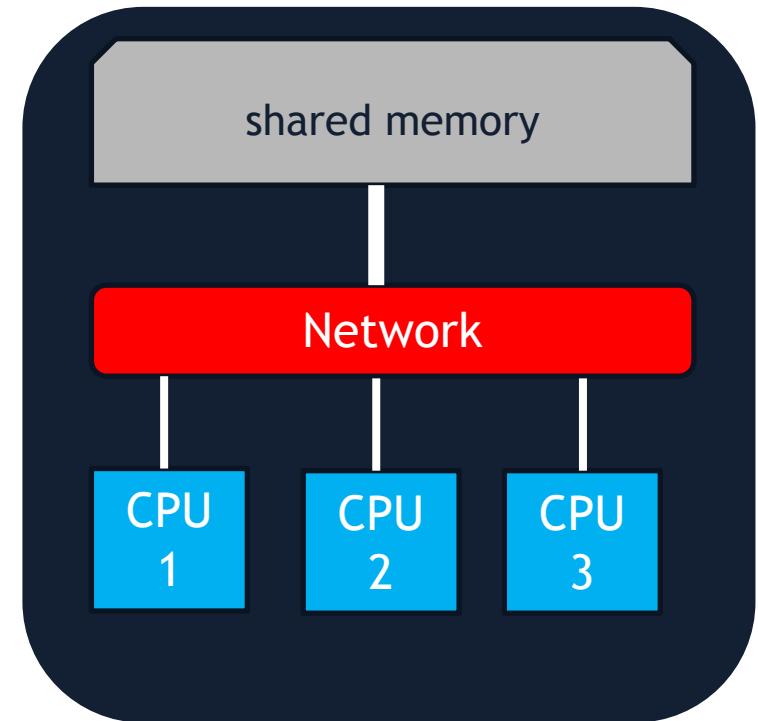


# Parallel computers: memory

Distributed Memory

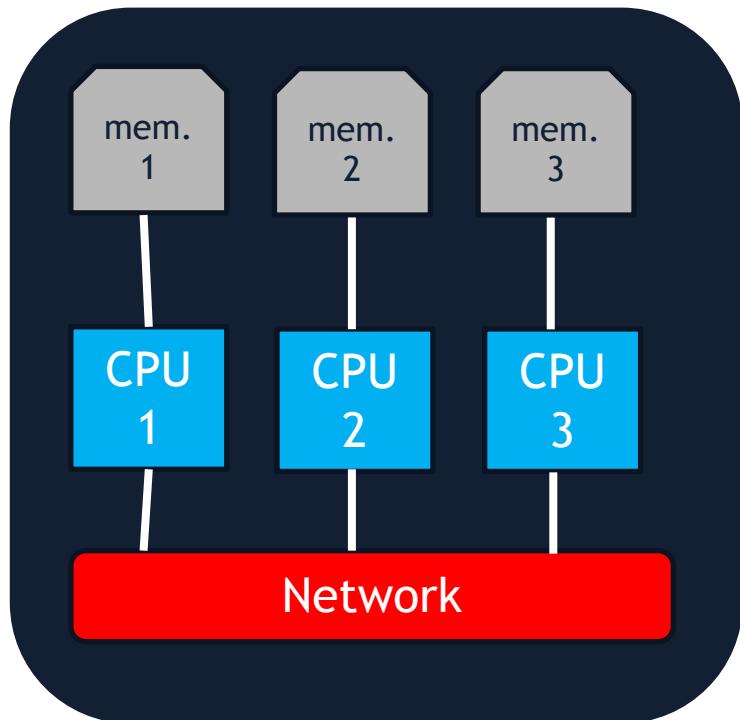


Shared Memory



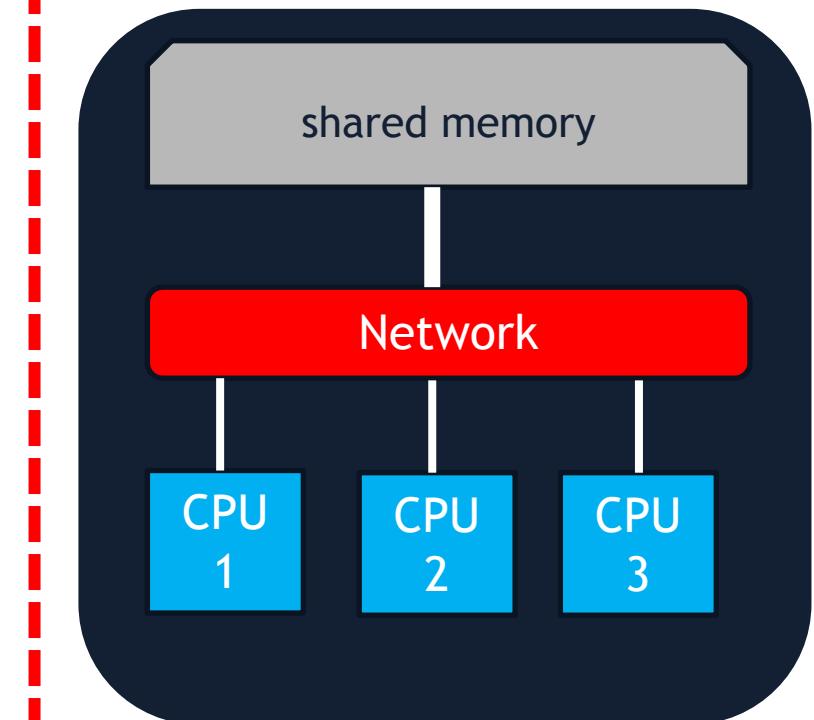
# Parallel computers: memory

Distributed Memory

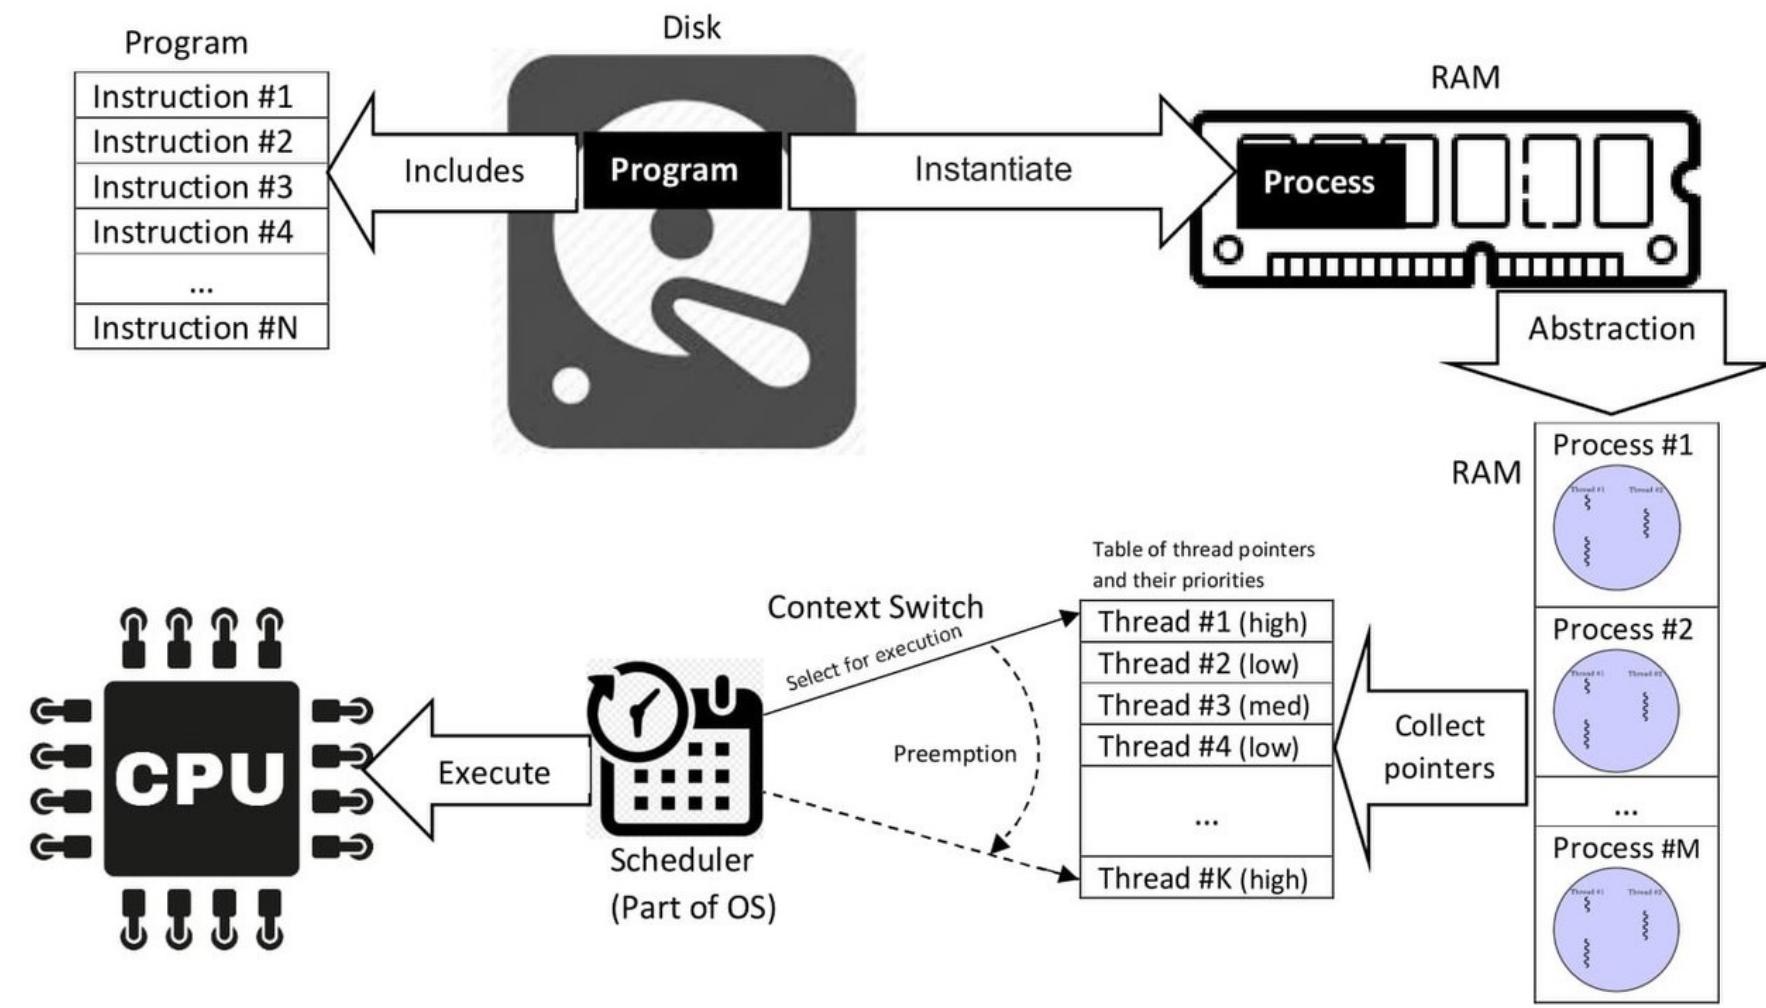


**OpenMP™**

Shared Memory

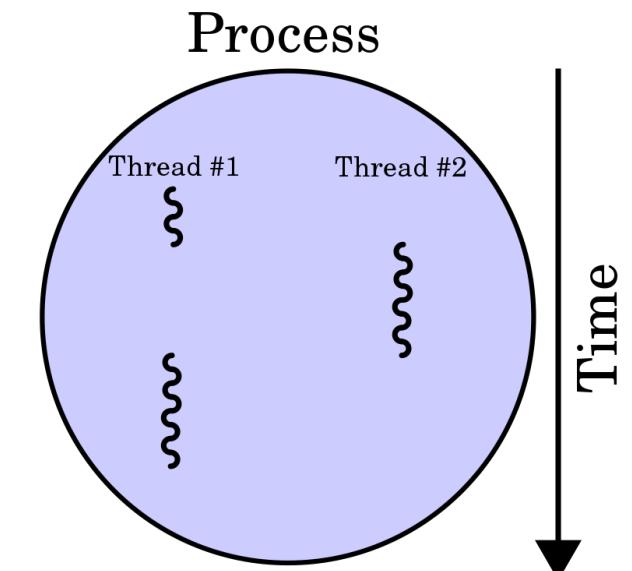


# Processes, threads, and more

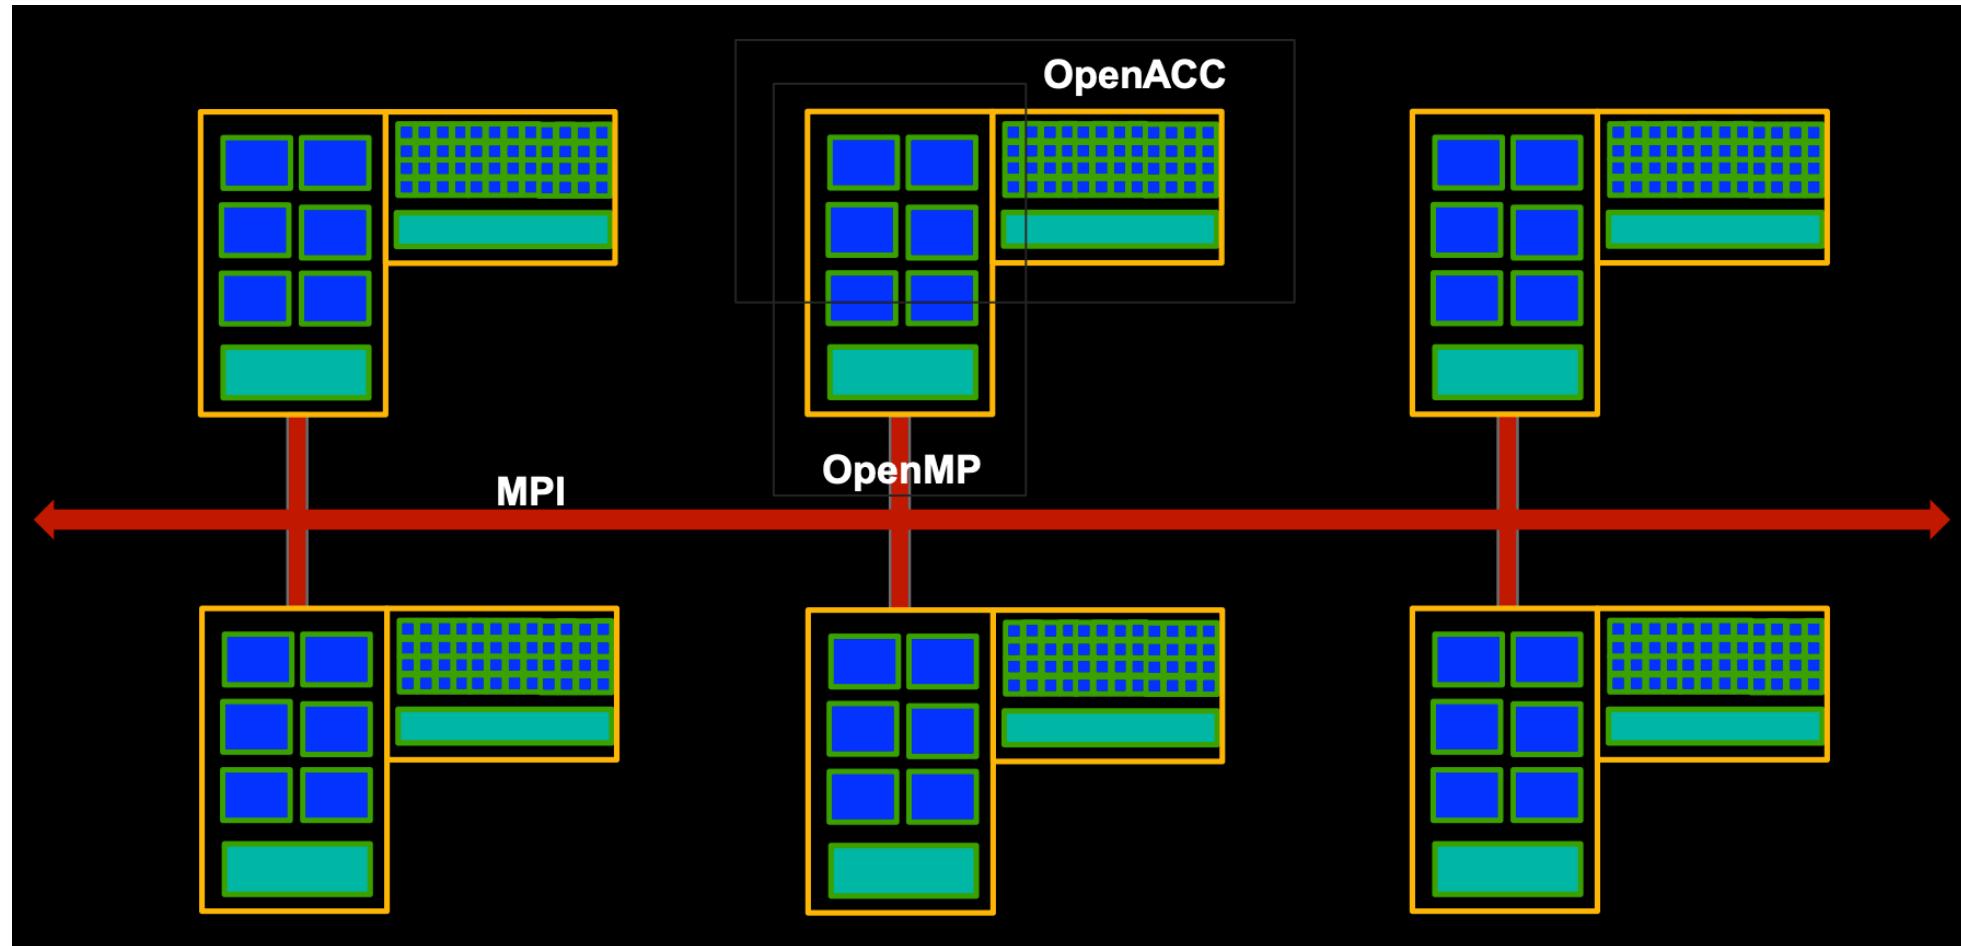


**Thread:** smallest sequence of instructions that the OS scheduler can manage

**Process:** composed of threads. Memory space allocated when program is loaded from disk.



# Frameworks



Courtesy of John Urbanic @ PSC



# Outline

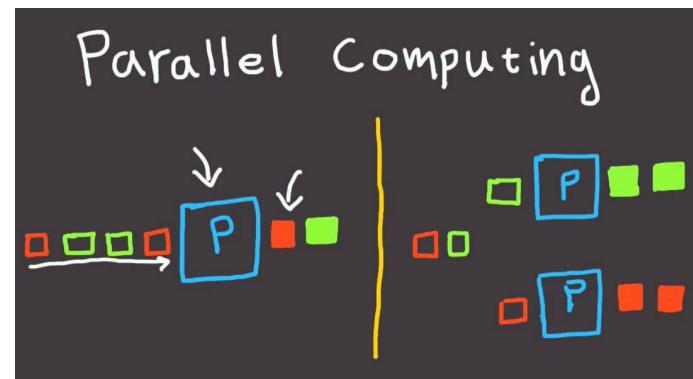
## Part A: Parallel Computing Quick Review (10 min)

- Task and Data parallelism
- Parallel computers
- Threads, processes and more



## Part B: MPI Intro (60 min)

- Core concepts
- Message passing
- Execution



## Part C: Exercises (45+ min)

**Ex 1: Hello World**

**Ex 2: Domain Decomposition**

**Ex 3: Parallel Linear Regression**

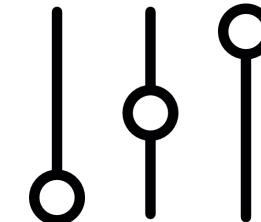
# What is MPI?

*“Message Passing Interface (MPI) is a standardized and portable message-passing framework designed to function on parallel computing architectures.”*

<https://www mpi-forum.org/>



- Library function/routines
- Different implementations



Application runtime behavior

- Established in the early 1990s by a group of researchers interested in standardizing message-passing
- Wide portability, applicable to multiple settings, independent of network speed or memory architecture
- Provides essential virtual topology, synchronization, and communication functionality between a set of processes in a language-independent way

# MPI Pros & Cons



- Robust, mature, widely portable framework
- Fully scalable parallelism
- Allows for strong scaling, massive speedups
- Present in every supercomputing system and on the roadmap of the Exascale generation
- Integration with a variety of accelerators and their frameworks (MPI+X, e.g. CUDA)

- Initial learning barrier can be discouraging
- No incremental parallelism
- May require re-designing the entire application
- Local installation can be tricky. Several implementations available (open-source and vendor specific)
- Library routines are implemented either on C or Fortran only. Python binders are available.

# When to (not) use MPI



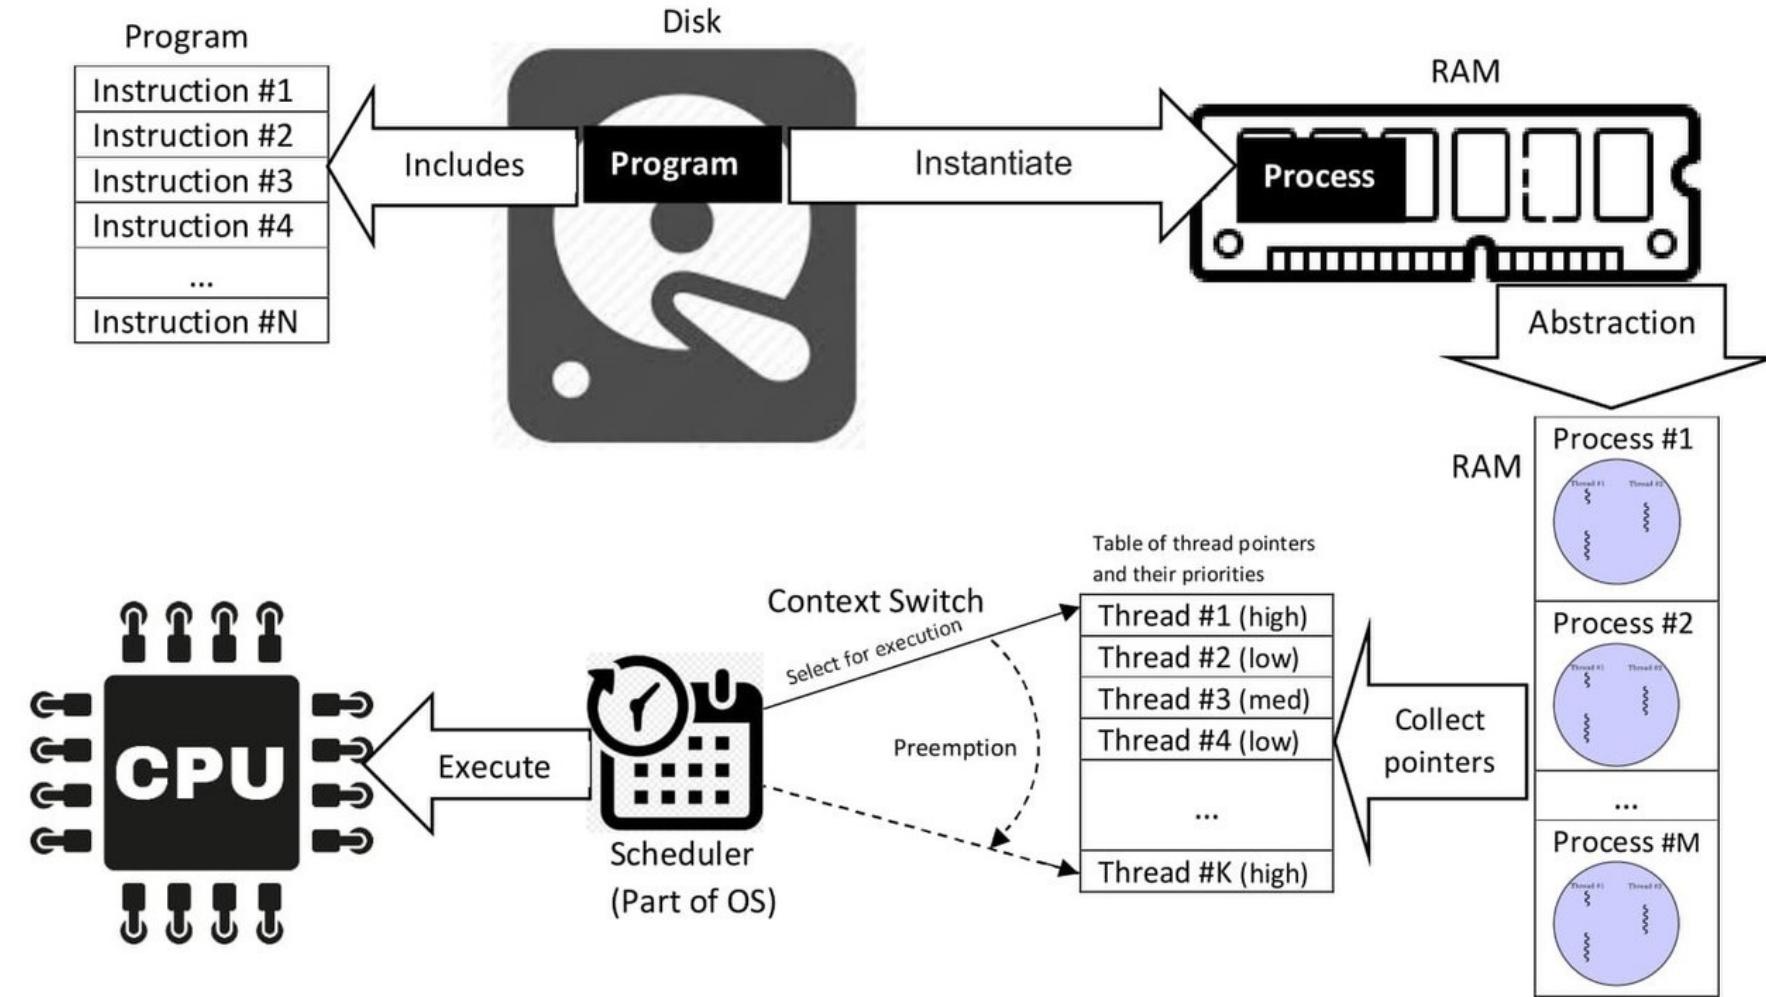
Parallel code needs to be portable across platforms

Large performance gain and/or strong scaling is required across the entire application

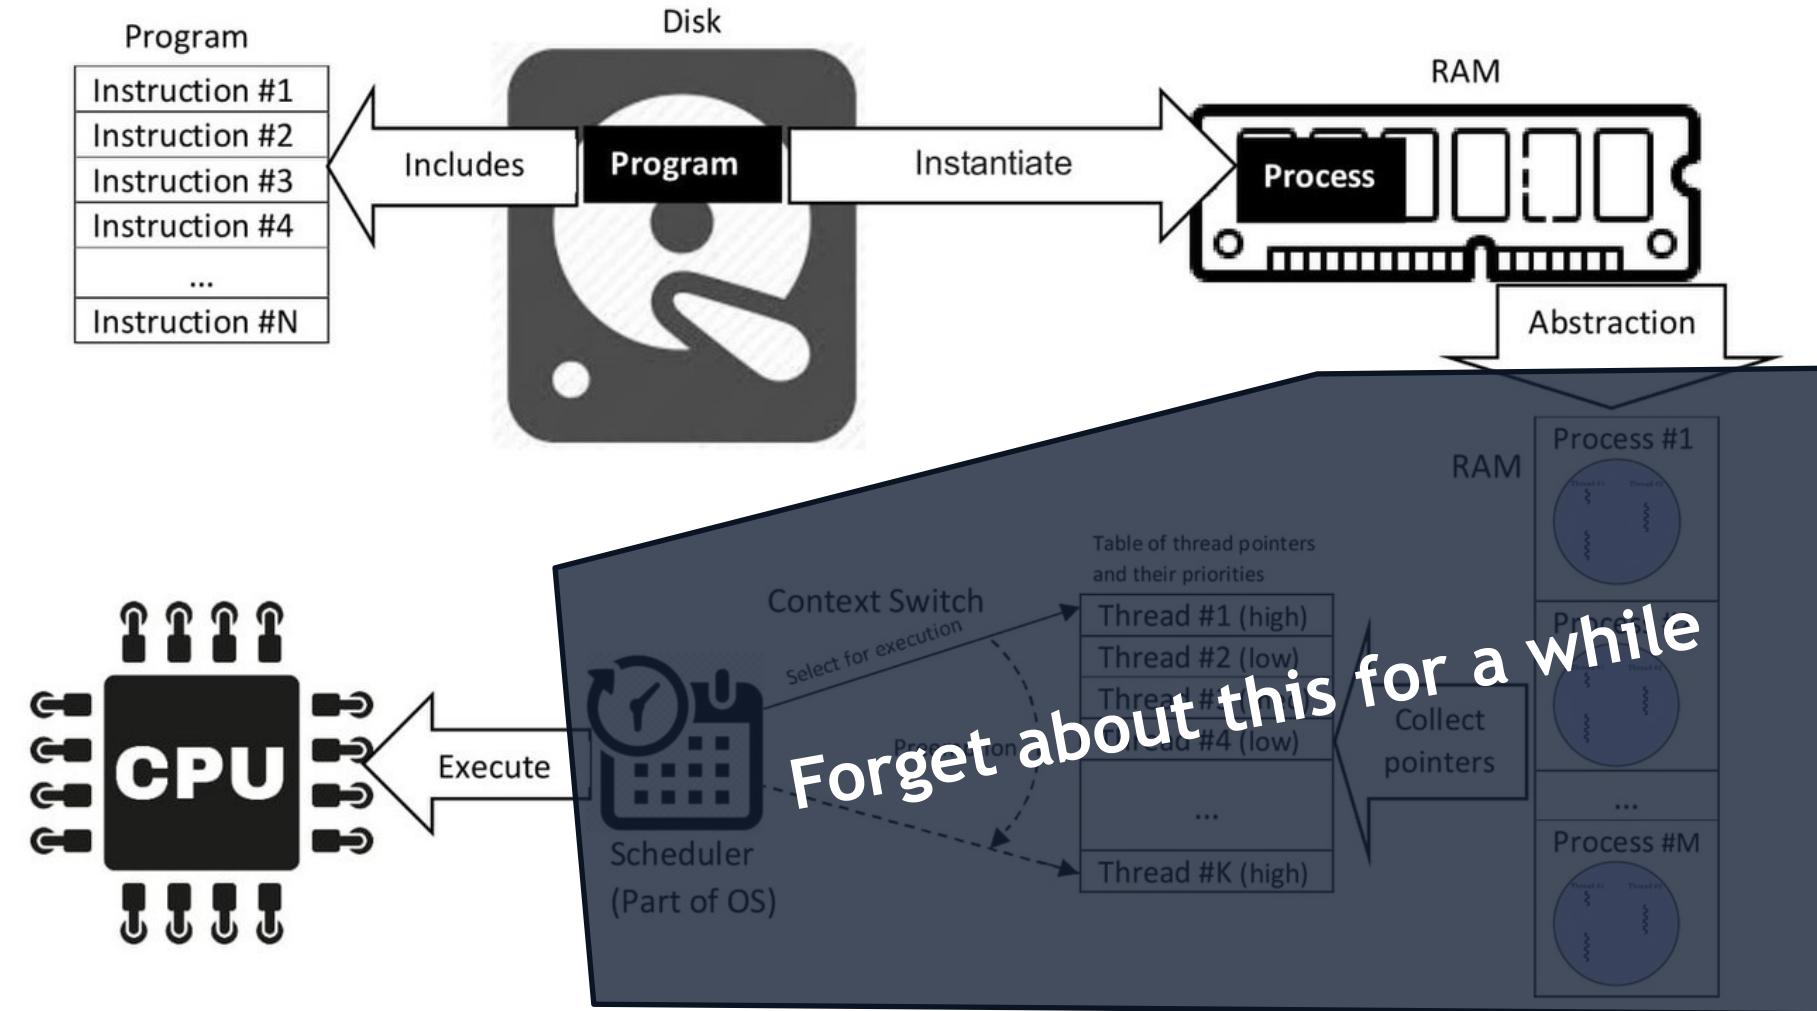
It's possible to achieve desired performance gain with only loop-level parallelism

Can use pre-existing parallel numerical libraries (e.g. ScaLAPACK)

# How MPI works

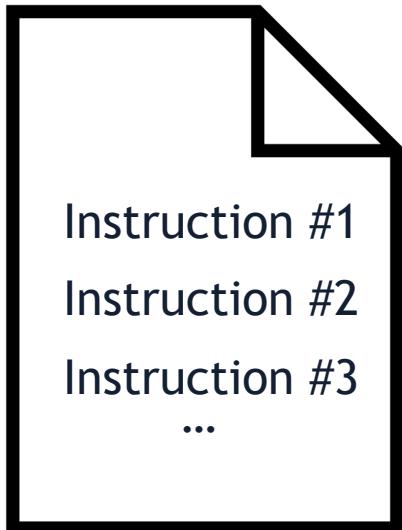


# How MPI works



# Serial code compilation and execution

`YourCode.<c/cpp/f90>`



Instruction #1

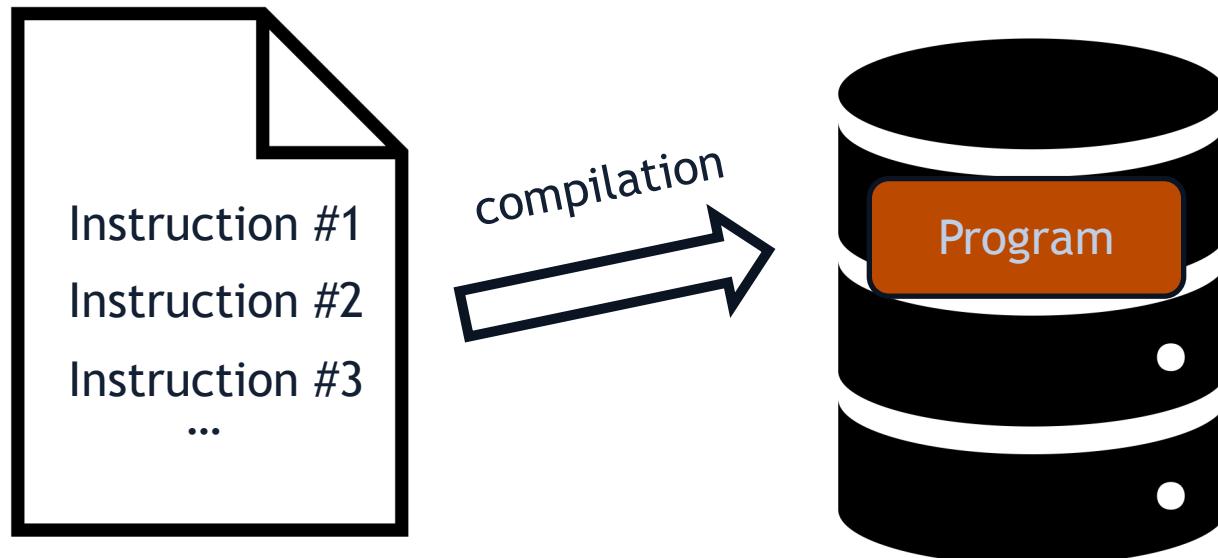
Instruction #2

Instruction #3

...

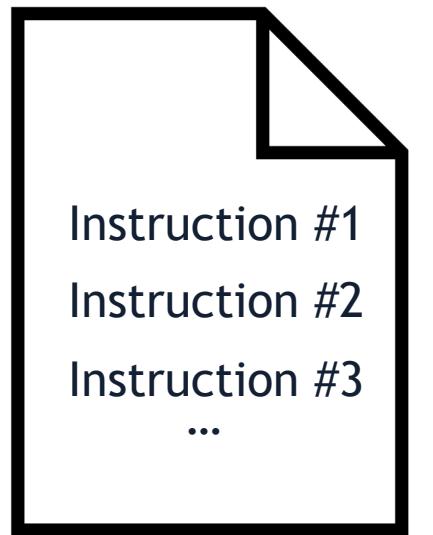
# Serial code compilation and execution

YourCode.<c/cpp/f90>



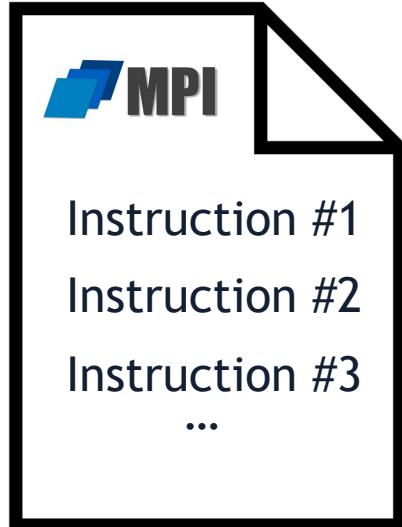
# Serial code compilation and execution

YourCode.<c/cpp/f90>



# MPI code compilation and execution

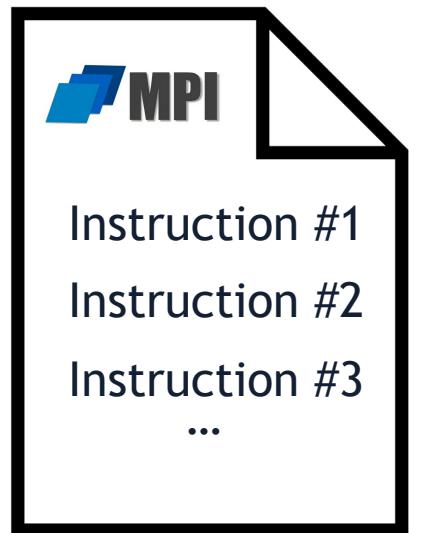
YourCode.<c/cpp/f90>



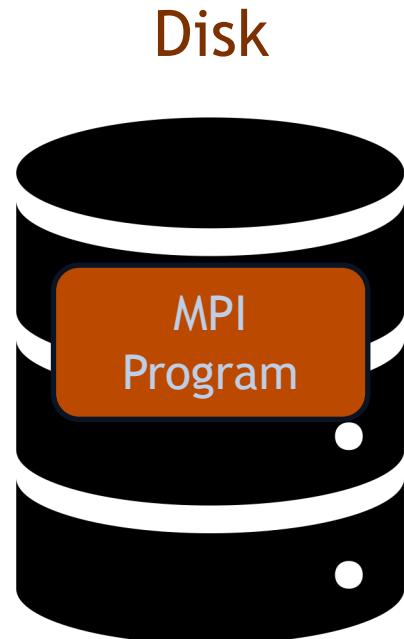
NCSA | NATIONAL CENTER FOR SUPERCOMPUTING APPLICATIONS

# MPI code compilation and execution

YourCode.<c/cpp/f90>



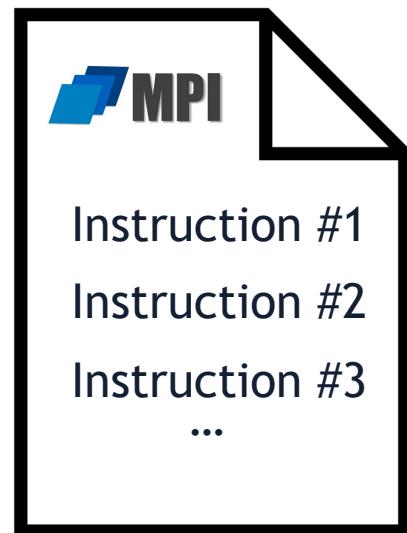
MPI  
compilation



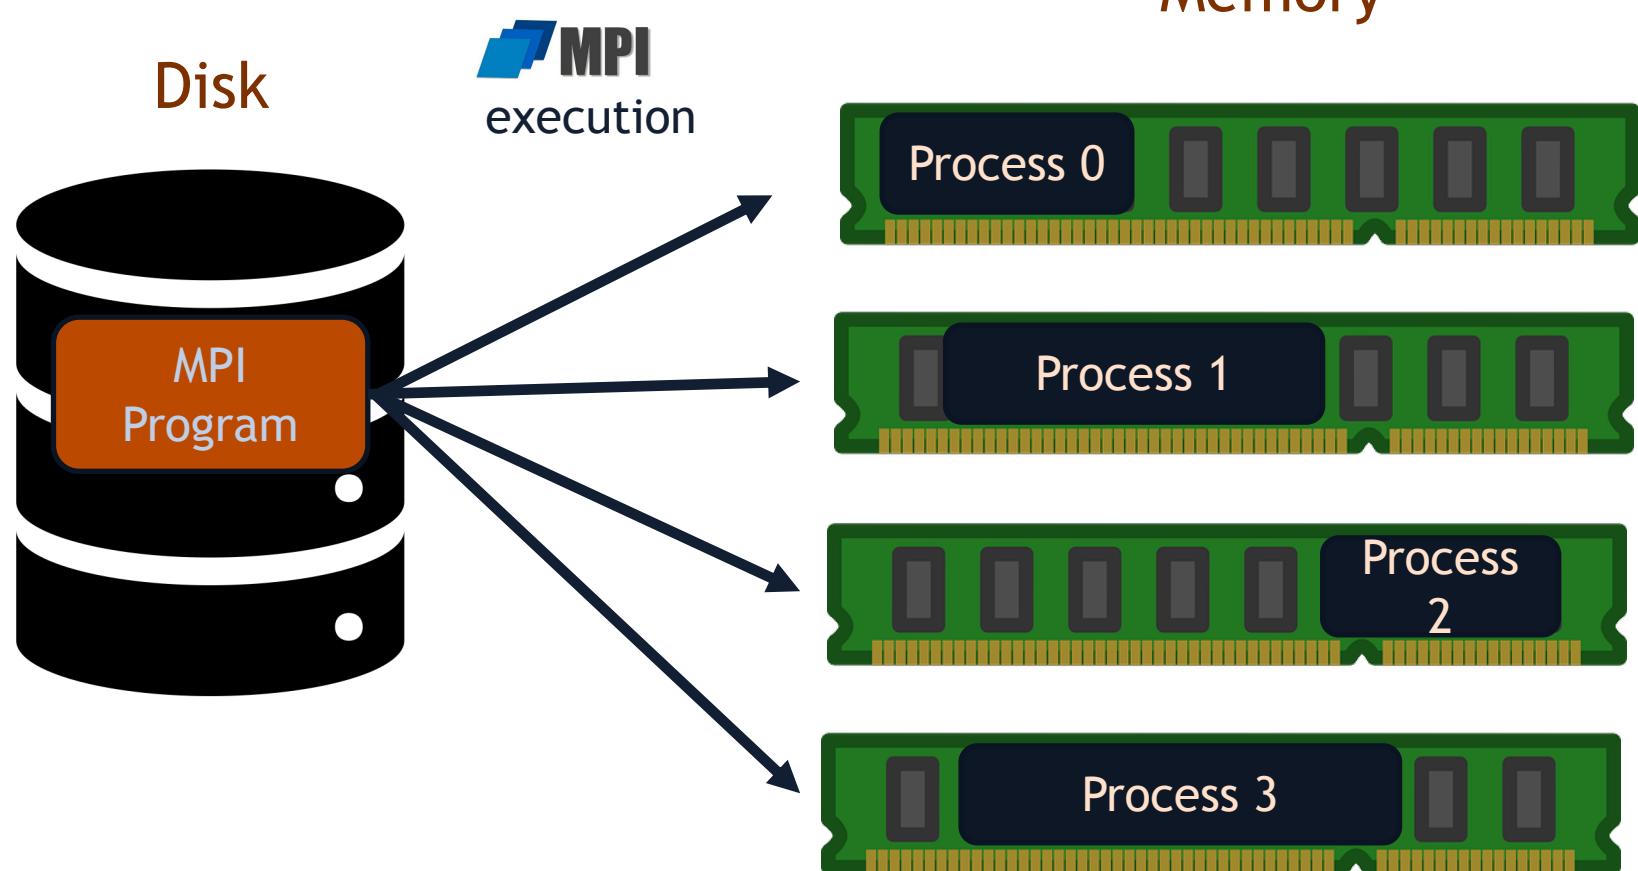
Disk

# MPI code compilation and execution

YourCode.<c/cpp/f90>

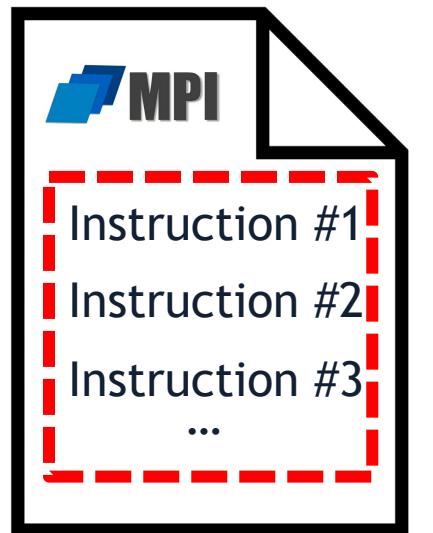


**MPI**  
compilation

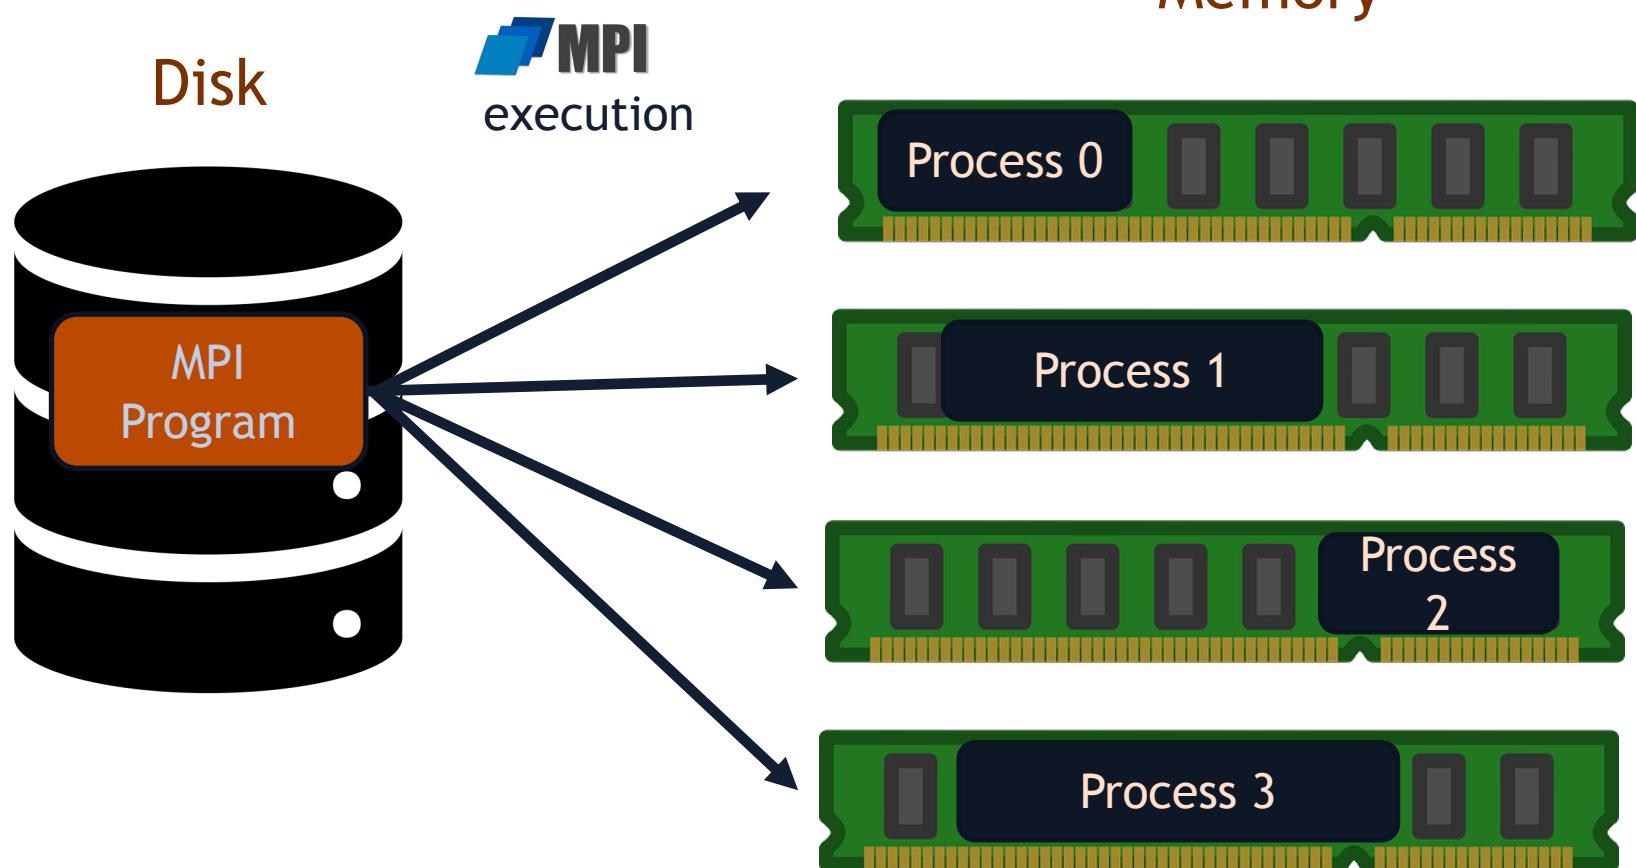


# MPI code compilation and execution

YourCode.<c/cpp/f90>



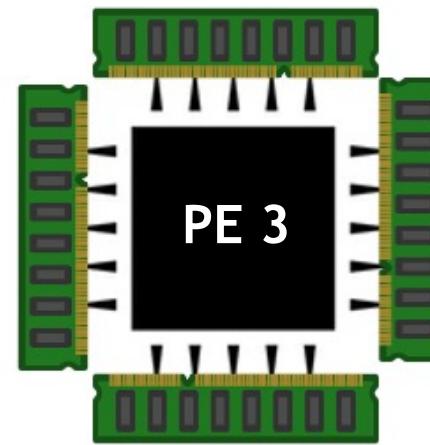
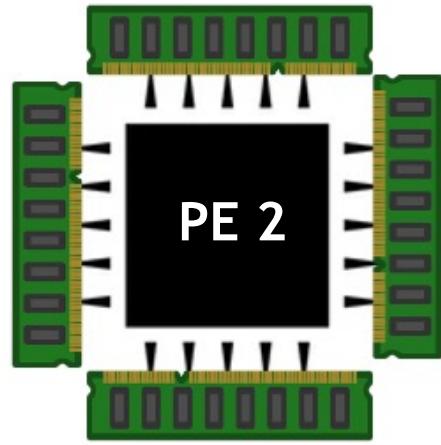
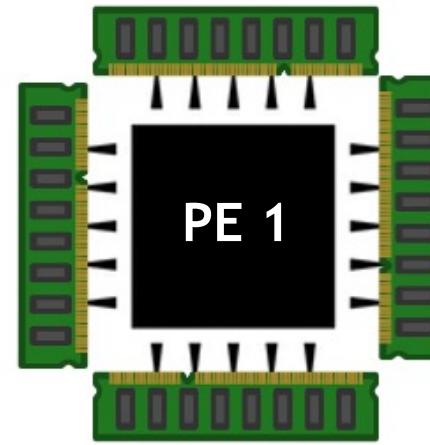
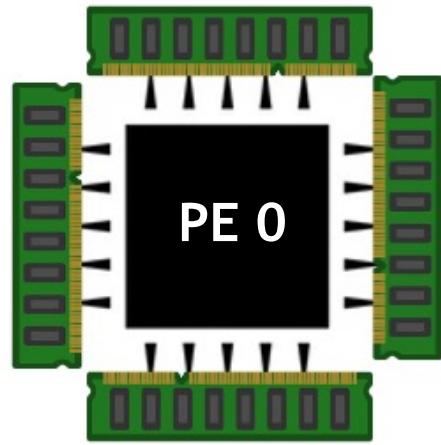
MPI compilation



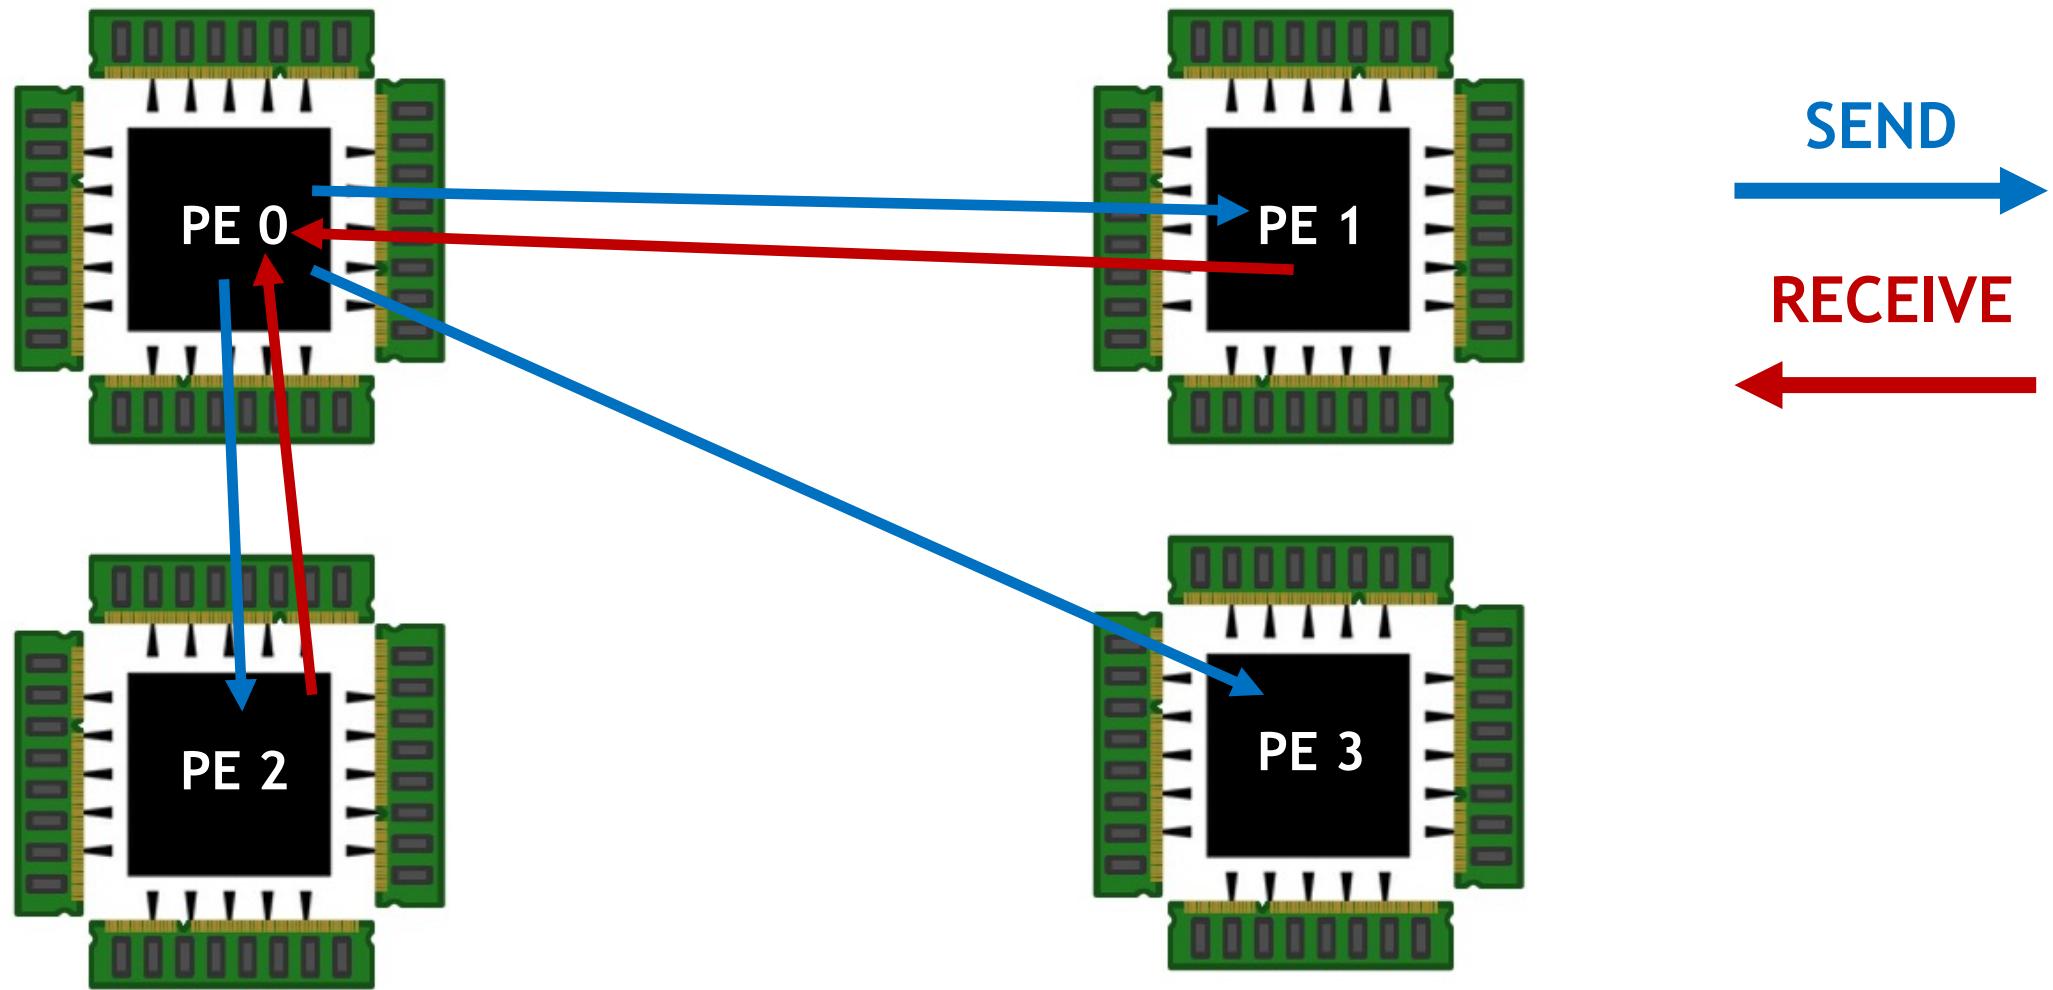
Instructions can be tailored to  
each and all processes with  
MPI library routines



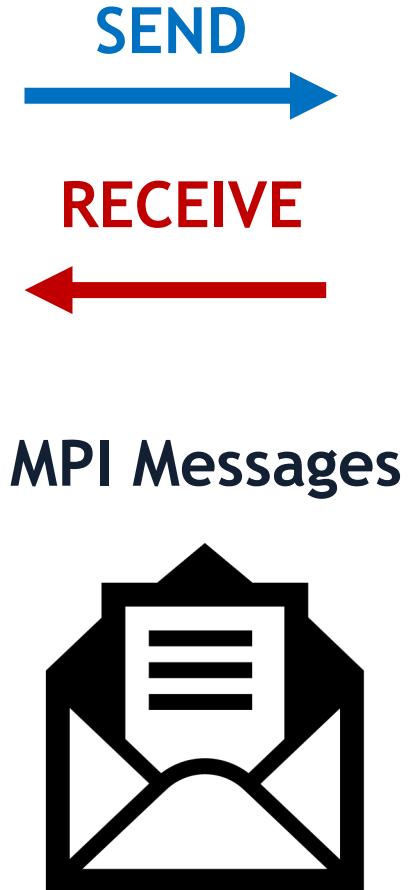
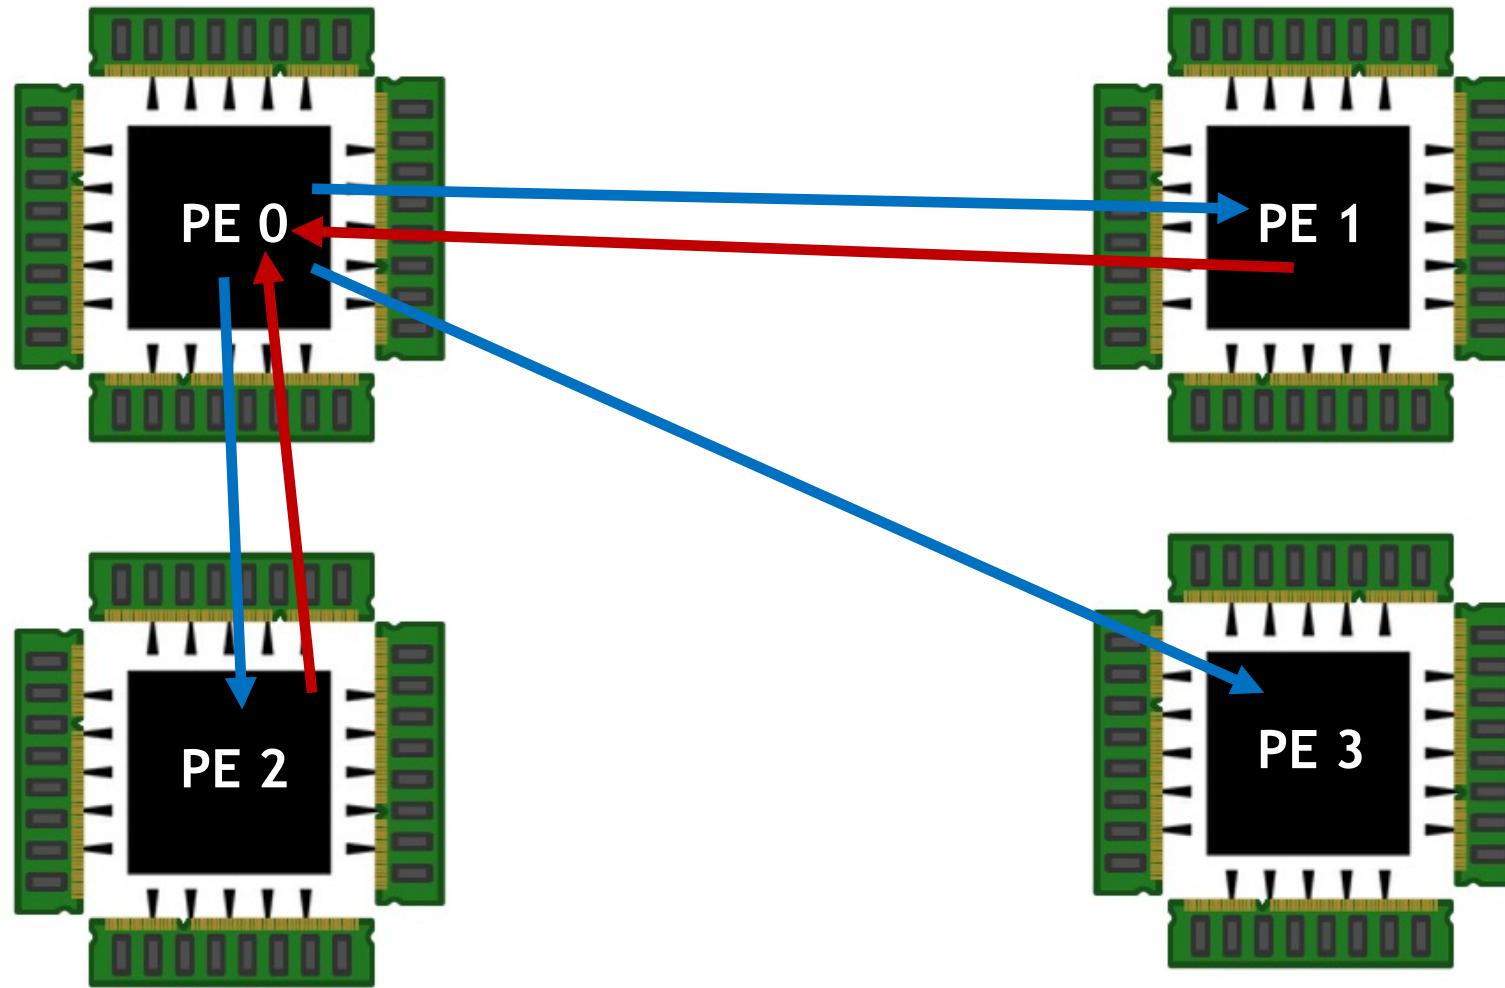
# MPI communication



# MPI communication



# MPI communication



# MPI messages



# MPI messages



# MPI messages



## ENVELOPE

Destination address, return address, and information needed to ensure correct delivery

- Source
- Destination
- Communicator
- Tag

# MPI messages



## BODY

What is being transferred (payload, content)

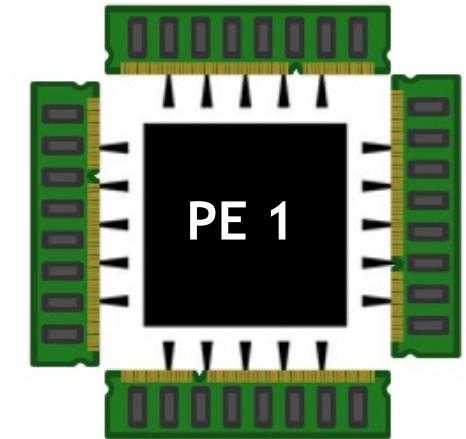
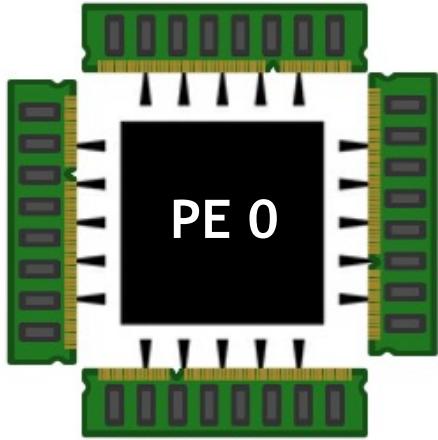
- Buffer (the data)
- Datatype (type of data in the message)
- Count (message size)

## ENVELOPE

Destination address, return address, and information needed to ensure correct delivery

- Source
- Destination
- Communicator
- Tag

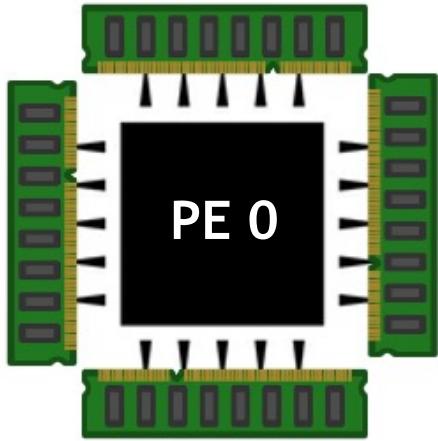
# MPI communication protocol



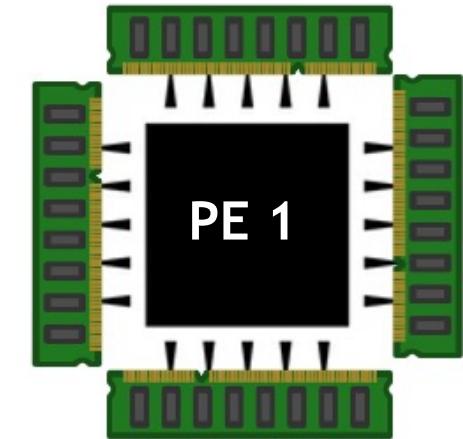
# MPI communication protocol



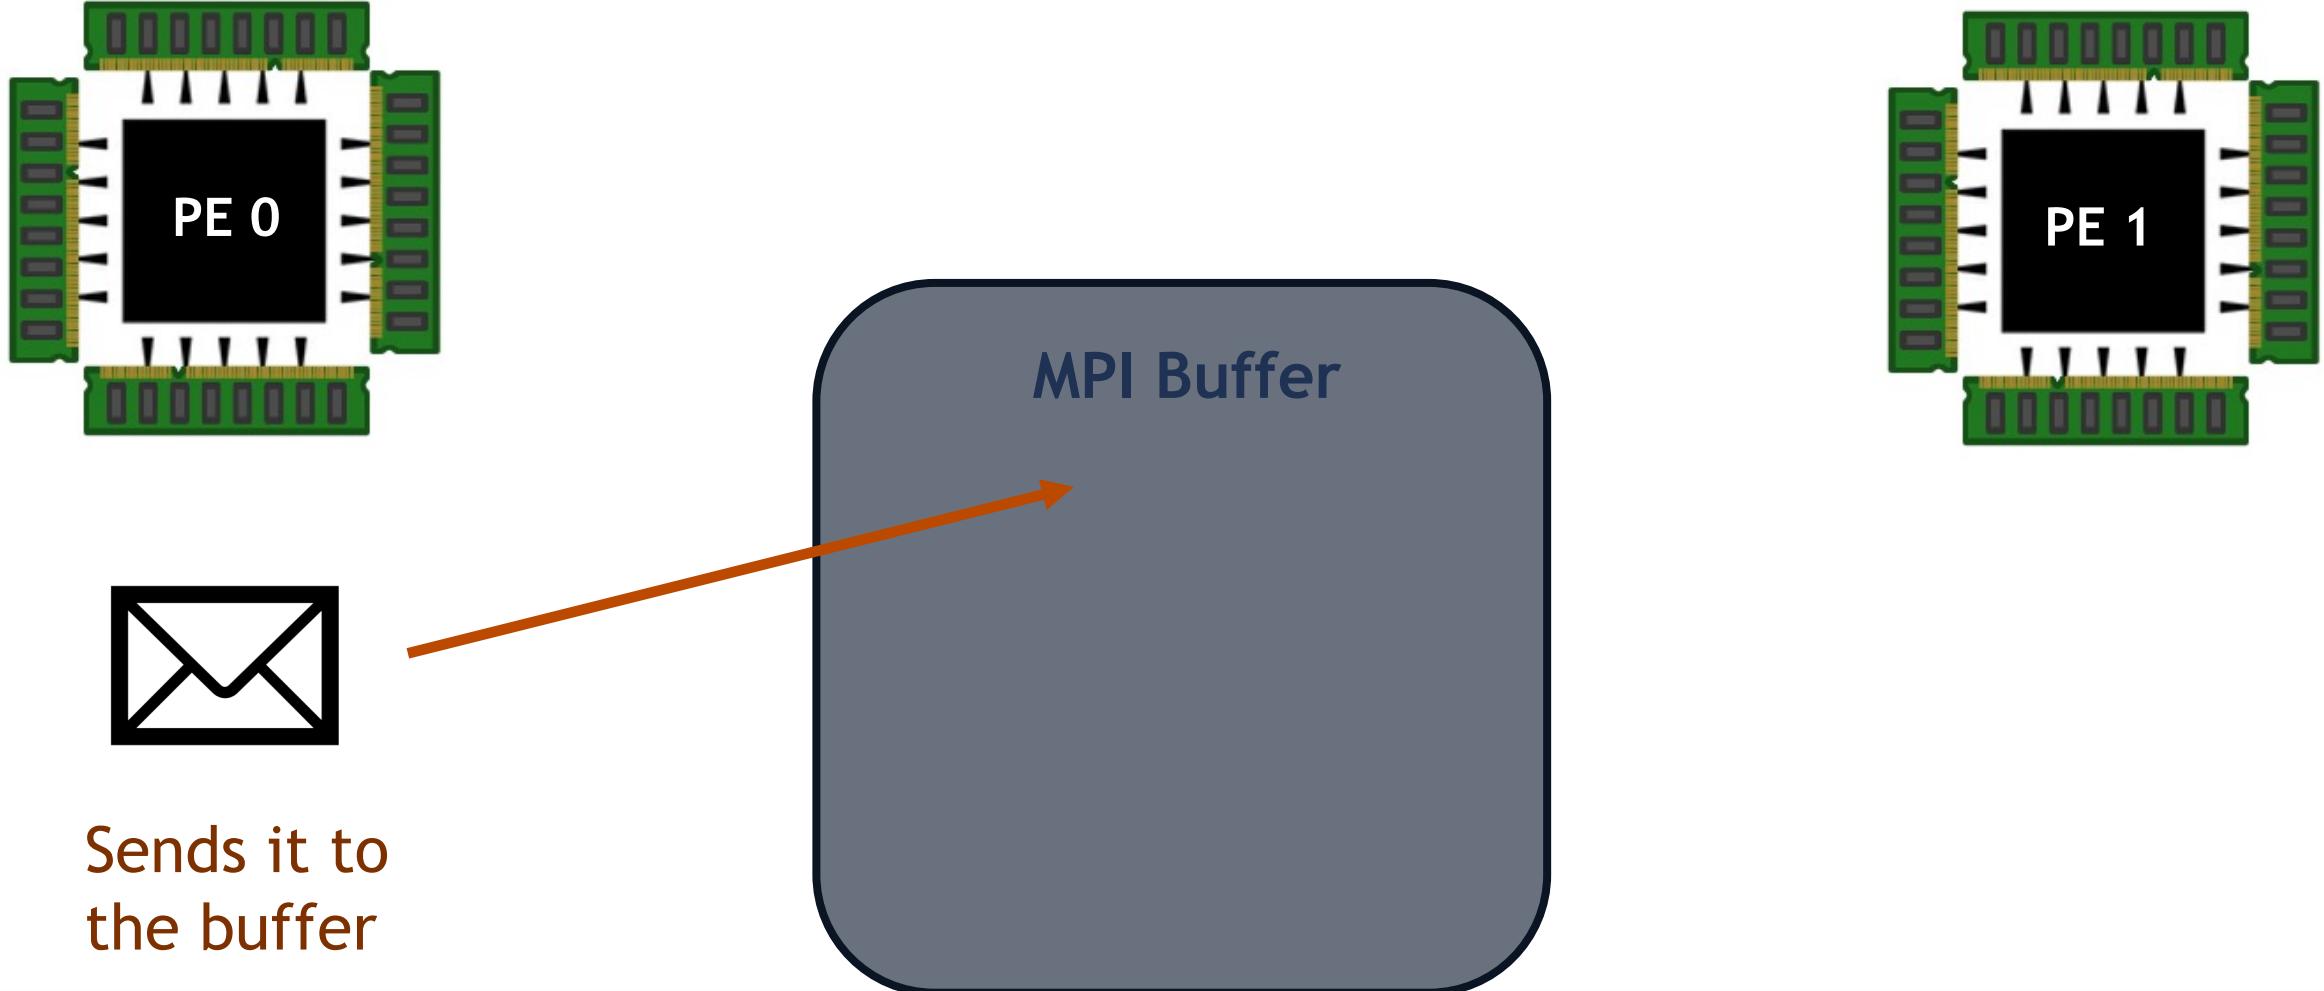
# MPI communication protocol



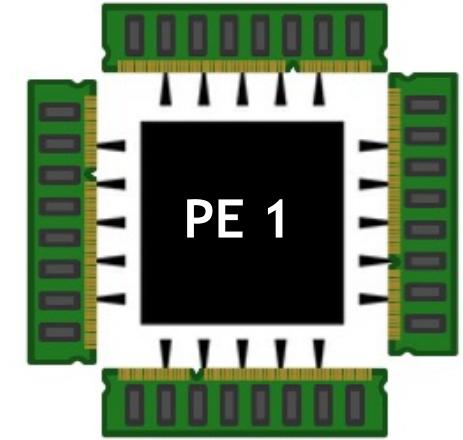
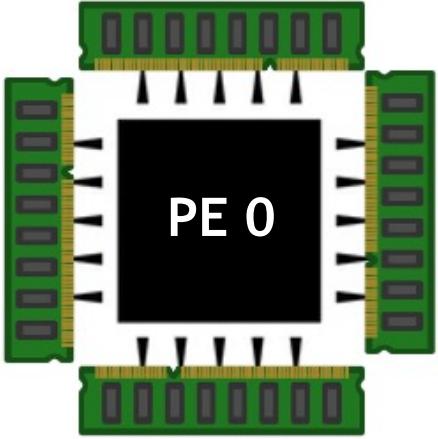
PE 0 writes  
a message



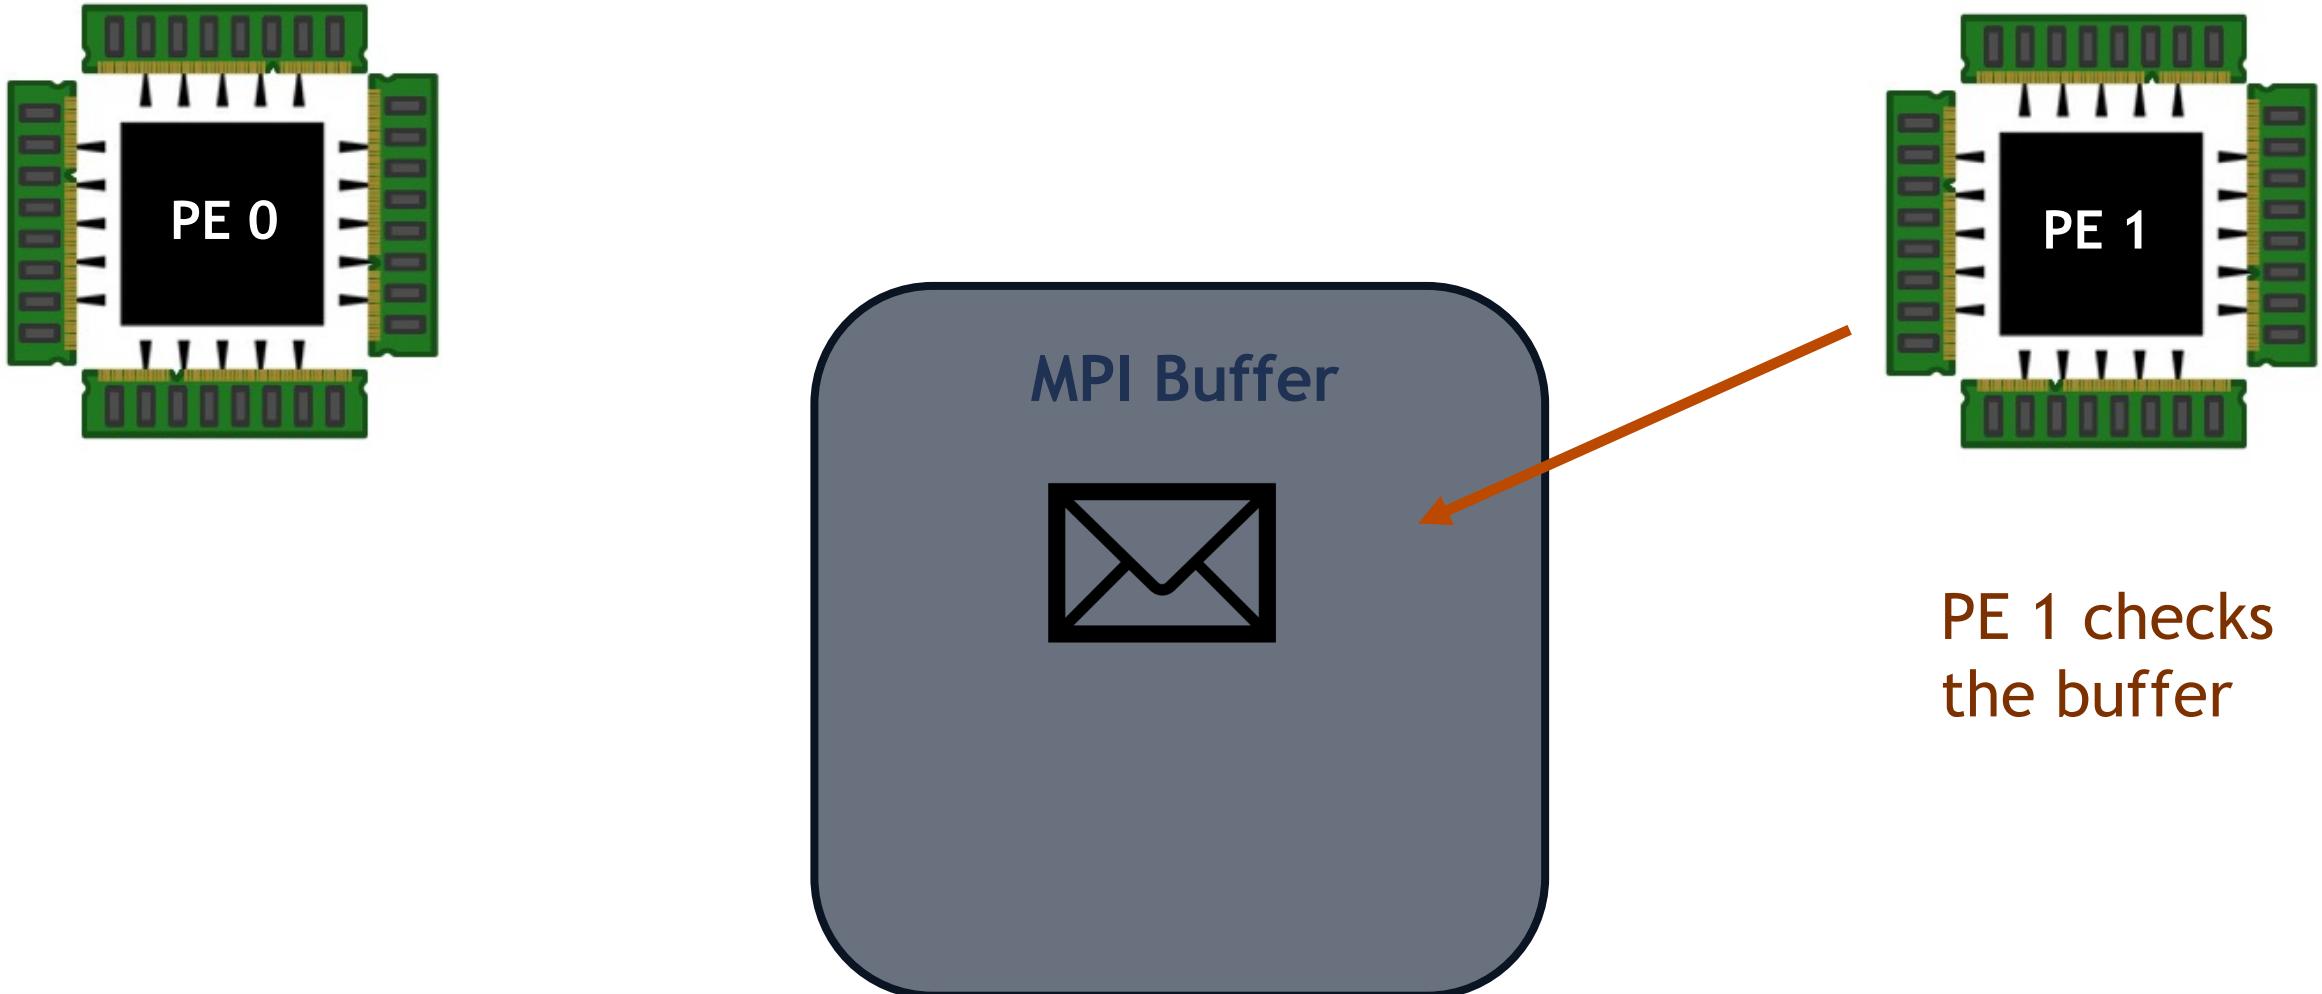
# MPI communication protocol



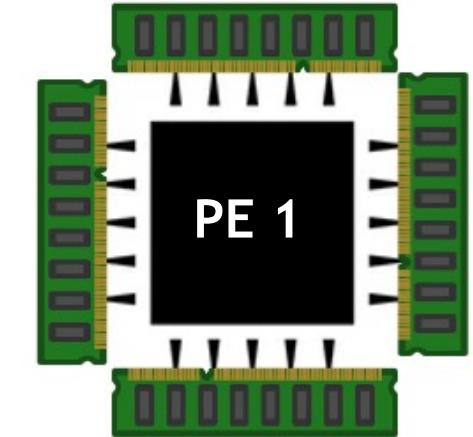
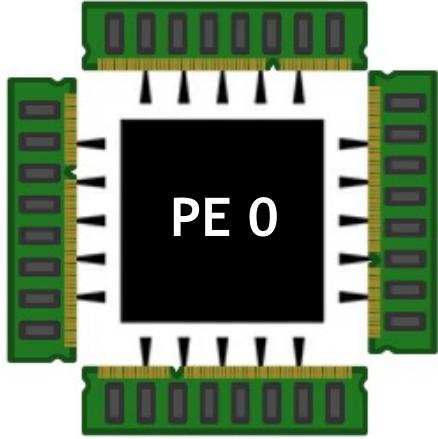
# MPI communication protocol



# MPI communication protocol



# MPI communication protocol



PE 1 receives  
the message

# MPI program structure

```
#include <iostream>
#include <mpi.h>

using namespace std;

int main()
{
    int my_rank;
    int mpierr; // this will hold error codes for MPI calls

    // start MPI
    mpierr = MPI_Init(NULL, NULL);

    // get my rank
    mpierr = MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);

    // print message
    cout << "Hello from PE " << my_rank << endl;

    // terminate MPI
    mpierr = MPI_Finalize();

    // goodbye
    return 0;
}
```



# MPI program structure

```
#include <iostream>
#include <mpi.h> ←
using namespace std;

int main()
{
    int my_rank;
    int mpierr; // this will hold error codes for MPI calls

    // start MPI
    mpierr = MPI_Init(NULL, NULL);

    // get my rank
    mpierr = MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);

    // print message
    cout << "Hello from PE " << my_rank << endl;

    // terminate MPI
    mpierr = MPI_Finalize();

    // goodbye
    return 0;
}
```

## 1. Include MPI header file



# MPI program structure

```
#include <iostream>
#include <mpi.h> ←
using namespace std;

int main()
{
    int my_rank; ←
    int mpierr; // this will hold error codes for MPI calls

    // start MPI
    mpierr = MPI_Init(NULL, NULL);

    // get my rank
    mpierr = MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);

    // print message
    cout << "Hello from PE " << my_rank << endl;

    // terminate MPI
    mpierr = MPI_Finalize();

    // goodbye
    return 0;
}
```

1. Include MPI header file
2. Variable declarations



# MPI program structure

```
#include <iostream>
#include <mpi.h> ←
using namespace std;

int main()
{
    int my_rank; ←
    int mpierr; // this will hold error codes for MPI calls

    // start MPI
    mpierr = MPI_Init(NULL, NULL); ←

    // get my rank
    mpierr = MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);

    // print message
    cout << "Hello from PE " << my_rank << endl;

    // terminate MPI
    mpierr = MPI_Finalize();

    // goodbye
    return 0;
}
```

1. Include MPI header file

2. Variable declarations

3. Initialize MPI environment



# MPI program structure

```
#include <iostream>
#include <mpi.h>

using namespace std;

int main()
{
    int my_rank; ← 1. Include MPI header file
    int mpierr; // this will hold error codes for MPI calls ← 2. Variable declarations

    // start MPI
    mpierr = MPI_Init(NULL, NULL); ← 3. Initialize MPI environment

    // get my rank
    mpierr = MPI_Comm_rank(MPI_COMM_WORLD, &my_rank); ← 4. Do computation, MPI calls

    // print message
    cout << "Hello from PE " << my_rank << endl;

    // terminate MPI
    mpierr = MPI_Finalize();

    // goodbye
    return 0;
}
```

1. Include MPI header file

2. Variable declarations

3. Initialize MPI environment

4. Do computation, MPI calls



# MPI program structure

```
#include <iostream>
#include <mpi.h>

using namespace std;

int main()
{
    int my_rank;
    int mpierr; // this will hold error codes for MPI calls

    // start MPI
    mpierr = MPI_Init(NULL, NULL);

    // get my rank
    mpierr = MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);

    // print message
    cout << "Hello from PE " << my_rank << endl;

    // terminate MPI
    mpierr = MPI_Finalize();

    // goodbye
    return 0;
}
```

1. Include MPI header file
2. Variable declarations
3. Initialize MPI environment
4. Do computation, MPI calls
5. Close MPI communications



# MPI library

The MPI Library has a myriad of functions that provide complete control over

- Communication procedures
- Task and load distribution
- Reductions and more complex operations

<https://www.mpi-forum.org/docs/mpi-4.0/mpi40-report.pdf>



# MPI library

The MPI Library has a myriad of functions that provide complete control over

- Communication procedures
  - Task and load distribution
  - Reductions and more complex

## MPI Function Index

The underlined page numbers refer to the function definitions.

control over procedures distribution more complex operations

<https://www mpi-forum.org/docs/mpi-4.0/mpi40-report.pdf>

MPI\_WAITANY, 54, 75, 76, 82, 519, 635, 636, 640  
MPI\_WAITSOME, 75, 80, 81, 82, 519, 635–637, 640  
MPI\_WIN\_ALLOCATE, 552, 556, 557, 559, 565, 567, 572, 605, 805, 806, 1047, 1049  
MPI\_Win\_allocate\_c, 556, 1046  
MPI\_WIN\_ALLOCATE\_CPTR, 557, 1049  
MPI\_WIN\_ALLOCATE\_SHARED, 339, 552, 558, 559, 561, 565, 567, 605, 806, 1047–1049  
MPI\_Win\_allocate\_shared\_c, 558, 1046  
MPI\_WIN\_ALLOCATE\_SHARED\_CPTR, 559, 1049  
MPI\_WIN\_ATTACH, 562, 563, 564, 565, 605  
MPI\_WIN\_F2C, 842  
MPI\_WIN\_CALL\_ERRHANDLER, 475, 477  
MPI\_WIN\_COMPLETE, 565, 594, 598, 599–602, 610, 617  
MPI\_WIN\_CREATE, 520, 552, 553, 555, 557, 559, 563–565, 567, 609  
MPI\_Win\_create\_c, 553, 1046  
MPI\_WIN\_CREATE\_DYNAMIC, 472, 552, 562, 563, 564, 566, 567, 610  
MPI\_WIN\_CREATE\_ERRHANDLER, 460, 463, 464, 875, 877, 1056  
MPI\_WIN\_CREATE\_KEYVAL, 365, 372, 379, 849, 874, 877, 1060  
MPI\_WIN\_DELETE\_ATTR, 365, 375, 379  
MPI\_WIN\_DETACH, 562, 564, 566  
MPI\_WIN\_DUP\_FN, 373, 865, 1051  
MPI\_WIN\_F2C, 842  
MPI\_WIN\_FENCE, 565, 574, 594, 596, 597, 598, 607, 608, 610, 611, 614, 619, 831  
MPI\_WIN\_FLUSH, 560, 584, 586, 605, 606, 610, 625, 626  
MPI\_WIN\_FLUSH\_ALL, 584, 586, 606, 610  
MPI\_WIN\_FLUSH\_LOCAL, 584, 606, 610  
MPI\_WIN\_FLUSH\_LOCAL\_ALL, 584, 607, 610  
MPI\_WIN\_FREE, 374, 495, 502, 547, 565, 566  
MPI\_WIN\_FREE\_KEYVAL, 365, 374, 379  
MPI\_WIN\_GET\_ATTR, 365, 375, 379, 566, 800, 849, 852  
MPI\_WIN\_GET\_ERRHANDLER, 460, 464, 1062  
MPI\_WIN\_GET\_GROUP, 567  
MPI\_WIN\_GET\_INFO, 567, 568, 569, 1047, 1054  
MPI\_WIN\_GET\_NAME, 386  
MPI\_WIN\_LOCK, 554, 565, 594, 602, 603–605, 607, 609, 611, 614–616  
MPI\_WIN\_LOCK\_ALL, 554, 594, 603, 604, 607, 609, 611, 616, 625  
MPI\_WIN\_NULL\_COPY\_FN, 373, 865, 1051  
MPI\_WIN\_NULL\_DELETE\_FN, 373, 865, 1051  
MPI\_WIN\_POST, 565, 594, 598, 599, 600–602, 604, 607, 608, 611, 617, 619  
MPI\_WIN\_SET\_ATTR, 365, 374, 379, 566, 800, 849, 853  
MPI\_WIN\_SET\_ERRHANDLER, 460, 464  
MPI\_WIN\_SET\_INFO, 568, 1047, 1054  
MPI\_WIN\_SET\_NAME, 385  
MPI\_WIN\_SHARED\_QUERY, 559, 560, 806, 1049  
MPI\_Win\_shared\_query\_c, 560, 1046  
MPI\_WIN\_SHARED\_QUERY\_CPTR, 562, 1049  
MPI\_WIN\_START, 565, 594, 598, 599–602, 607, 608, 617, 624  
MPI\_WIN\_SYNC, 607, 611–614, 617, 624, 626  
MPI\_WIN\_TEST, 600, 601  
MPI\_WIN\_UNLOCK, 565, 586, 594, 603, 605, 610, 611, 614, 615  
MPI\_WIN\_UNLOCK\_ALL, 586, 594, 603, 604, 610, 611, 614, 625  
MPI\_WIN\_WAIT, 565, 594, 600, 601, 602, 604, 610, 611, 614, 617, 618  
MPI\_WTICK, 26, 478, 760  
MPI\_WTIME, 16, 26, 454, 477, 478, 728, 747, 760  
mpexec, 489, 490, 492, 515, 1048  
mpirun, 515  
PMPL, 725, 800  
PMPL\_AINT\_ADD, 26  
PMPL\_AINT\_DIFF, 26  
PMPLISEND, 800, 803  
PMPLWTICK, 26  
PMPL\_WTIME, 26



# MPI functions syntax

The good thing is that their syntax is quite standardized...

C/C++

Fortran



# MPI functions syntax

The good thing is that their syntax is quite standardized...

C/C++

**MPI\_**

Fortran

**MPI\_**

- Indicates that an MPI function/routine is called



# MPI functions syntax

The good thing is that their syntax is quite standardized...

C/C++

`MPI_FunctionName`

Fortran

`MPI_FUNCTIONNAME`

- Indicates that an MPI function/routine is called
- **Typically a description of what the function does**



# MPI functions syntax

The good thing is that their syntax is quite standardized...

C/C++

`MPI_FunctionName( )`

Fortran

`MPI_FUNCTIONNAME( )`

- Indicates that an MPI function/routine is called
- Typically a description of what the function does
- Arguments to be passed and collected



# MPI functions syntax

The good thing is that their syntax is quite standardized...

C/C++

```
MPI_FunctionName( BODY , ENVELOPE )
```

Fortran

```
MPI_FUNCTIONNAME( BODY , ENVELOPE )
```

- Indicates that an MPI function/routine is called
- Typically a description of what the function does
- Arguments to be passed and collected
  - Envelope
  - Body



# MPI call: Send and Receive

C/C++

```
MPI_FunctionName( BODY , ENVELOPE )
```

```
MPI_FunctionName( BODY , ENVELOPE )
```

Fortran

```
MPI_FUNCTIONNAME( BODY , ENVELOPE )
```

```
MPI_FUNCTIONNAME( BODY , ENVELOPE )
```



# MPI call: Send and Receive

C/C++

```
MPI_Send( BODY , ENVELOPE)
```

```
MPI_Recv( BODY , ENVELOPE)
```

Fortran

```
MPI_SEND( BODY , ENVELOPE )
```

```
MPI_RECV( BODY , ENVELOPE )
```



# MPI call: Send and Receive

C/C++

```
MPI_Send( BODY , ENVELOPE)
```

```
MPI_Recv( BODY , ENVELOPE)
```

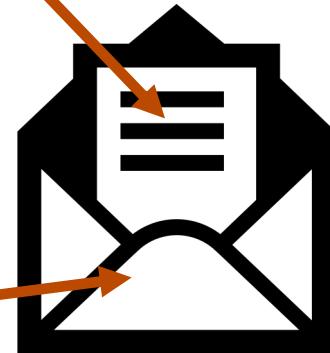
Fortran

```
MPI_SEND( BODY , ENVELOPE )
```

```
MPI_RECV( BODY , ENVELOPE )
```

BODY

- Buffer
- Count
- Data type



ENVELOPE

- Source
- Destination
- Tag
- Communicator

# MPI call: Send and Receive

## C/C++

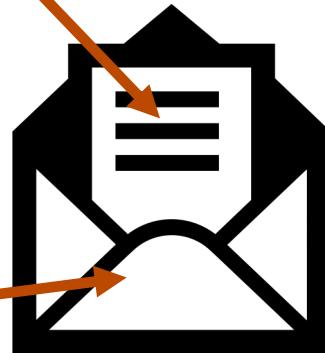
```
MPI_Send(void *buf,      )  
  
MPI_Recv(void *buf,      )
```

## Fortran

```
MPI_SEND(BUF,      )  
  
MPI_RECV(BUF,      )
```

## BODY

- Buffer
- Count
- Data type



## ENVELOPE

- Source
- Destination
- Tag
- Communicator

# MPI call: Send and Receive

## C/C++

```
MPI_Send(void *buf, int count, )
```

```
MPI_Recv(void *buf, int count, )
```

## Fortran

```
MPI_SEND(BUF, COUNT, )
```

```
MPI_RECV(BUF, COUNT, )
```

## BODY

- Buffer
- Count
- Data type



## ENVELOPE

- Source
- Destination
- Tag
- Communicator



# MPI call: Send and Receive

## C/C++

```
MPI_Send(void *buf, int count, MPI_Datatype dtype, )  
MPI_Recv(void *buf, int count, MPI_Datatype dtype, )
```

## Fortran

```
MPI_SEND(BUF, COUNT, DTYPE, )  
MPI_RECV(BUF, COUNT, DTYPE, )
```

### BODY

- Buffer
- Count
- Data type



### ENVELOPE

- Source
- Destination
- Tag
- Communicator

# MPI call: Send and Receive

## C/C++

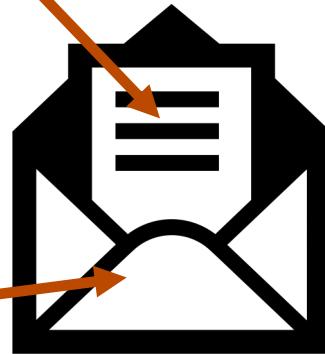
```
MPI_Send(void *buf, int count, MPI_Datatype dtype, )  
MPI_Recv(void *buf, int count, MPI_Datatype dtype, int src,)
```

## Fortran

```
MPI_SEND(BUF, COUNT, DTYPE, )  
MPI_RECV(BUF, COUNT, DTYPE, SOURCE, )
```

## BODY

- Buffer
- Count
- Data type



## ENVELOPE

- Source
- Destination
- Tag
- Communicator

# MPI call: Send and Receive

## C/C++

```
MPI_Send(void *buf, int count, MPI_Datatype dtype, int dst, )
```

```
MPI_Recv(void *buf, int count, MPI_Datatype dtype, int src, )
```

## Fortran

```
MPI_SEND(BUF, COUNT, DTYPE, DEST, )
```

```
MPI_RECV(BUF, COUNT, DTYPE, SOURCE, )
```

## BODY

- Buffer
- Count
- Data type



## ENVELOPE

- Source
- Destination
- Tag
- Communicator



# MPI call: Send and Receive

## C/C++

```
MPI_Send(void *buf, int count, MPI_Datatype dtype, int dst,  
int tag, )
```

```
MPI_Recv(void *buf, int count, MPI_Datatype dtype, int src,  
int tag, )
```

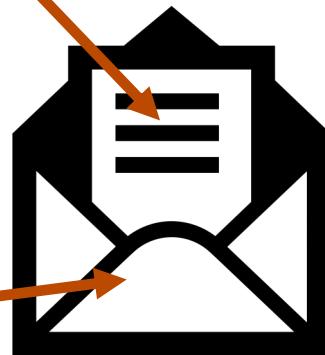
## Fortran

```
MPI_SEND(BUF, COUNT, DTYPE, DEST, TAG, )
```

```
MPI_RECV(BUF, COUNT, DTYPE, SOURCE, TAG, )
```

## BODY

- Buffer
- Count
- Data type



## ENVELOPE

- Source
- Destination
- Tag
- Communicator



# MPI call: Send and Receive

## C/C++

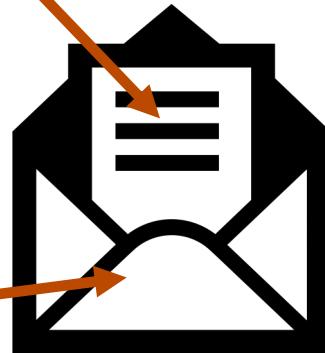
```
MPI_Send(void *buf, int count, MPI_Datatype dtype, int dst,  
int tag, MPI_Comm comm);  
  
MPI_Recv(void *buf, int count, MPI_Datatype dtype, int src,  
int tag, MPI_Comm comm, )
```

## Fortran

```
MPI_SEND(BUF, COUNT, DTYPE, DEST, TAG, COMM, )  
  
MPI_RECV(BUF, COUNT, DTYPE, SOURCE, TAG, COMM, )
```

## BODY

- Buffer
- Count
- Data type



## ENVELOPE

- Source
- Destination
- Tag
- Communicator



# MPI call: Send and Receive

## C/C++

```
MPI_Send(void *buf, int count, MPI_Datatype dtype, int dst,  
int tag, MPI_Comm comm);  
  
MPI_Recv(void *buf, int count, MPI_Datatype dtype, int src,  
int tag, MPI_Comm comm, MPI_Status *status);
```

## Fortran

Status about the message that was received

```
MPI_SEND(BUF, COUNT, DTYPE, DEST, TAG, COMM, )
```

```
MPI_RECV(BUF, COUNT, DTYPE, SOURCE, TAG, COMM, STATUS, )
```

## BODY

- Buffer
- Count
- Data type



## ENVELOPE

- Source
- Destination
- Tag
- Communicator

# MPI call: Send and Receive

## C/C++

```
int ierr = MPI_Send(void *buf, int count, MPI_Datatype  
dtype, int dst, int tag, MPI_Comm comm);  
  
int ierr = MPI_Recv(void *buf, int count, MPI_Datatype  
dtype, int src, int tag, MPI_Comm comm, MPI_Status *status);
```

## Fortran

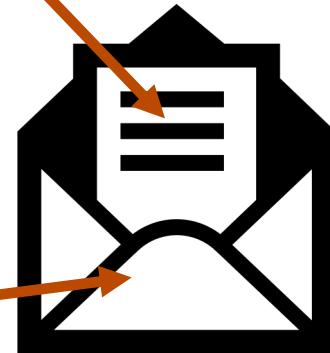
Error code when routine/function returns

```
MPI_SEND(BUF, COUNT, DTYPE, DEST, TAG, COMM, IERR)
```

```
MPI_RECV(BUF, COUNT, DTYPE, SOURCE, TAG, COMM, STATUS, IERR)
```

## BODY

- Buffer
- Count
- Data type



## ENVELOPE

- Source
- Destination
- Tag
- Communicator

# MPI call: Send and Receive

## C/C++

```
int ierr = MPI_Send(void *buf, int count, MPI_Datatype  
dtype, int dst, int tag, MPI_Comm comm);  
  
int ierr = MPI_Recv(void *buf, int count, MPI_Datatype  
dtype, int src, int tag, MPI_Comm comm, MPI_Status *status);
```

## Fortran

```
MPI_SEND(BUF, COUNT, DTYPE, DEST, TAG, COMM, IERR)
```

```
MPI_RECV(BUF, COUNT, DTYPE, SOURCE, TAG, COMM, STATUS, IERR)
```



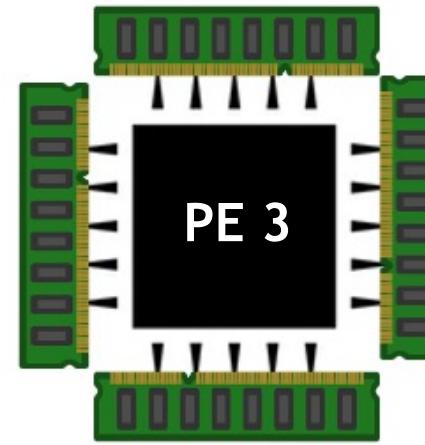
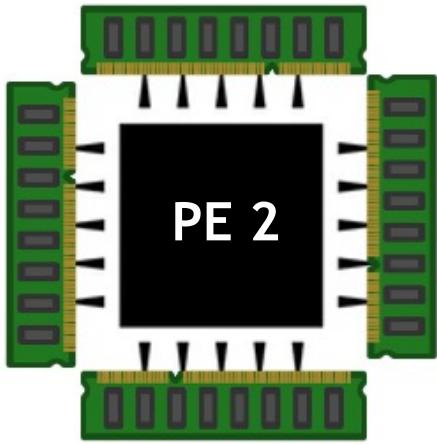
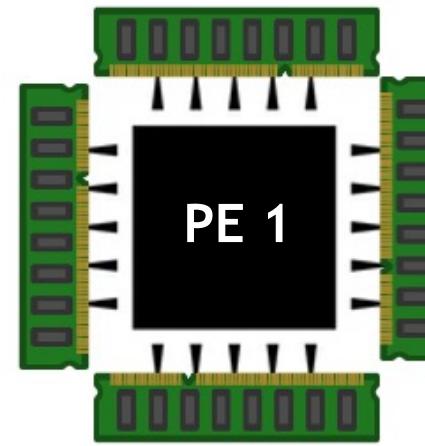
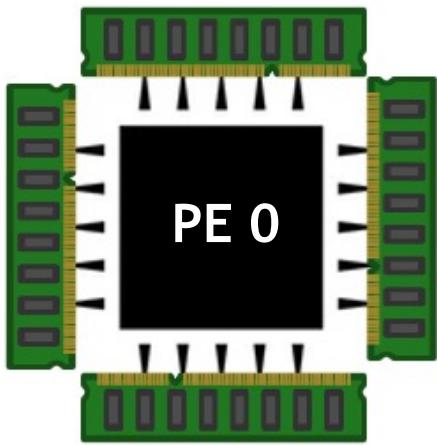
# MPI collectives

Collective communications are a great way to get started with MPI. They are:

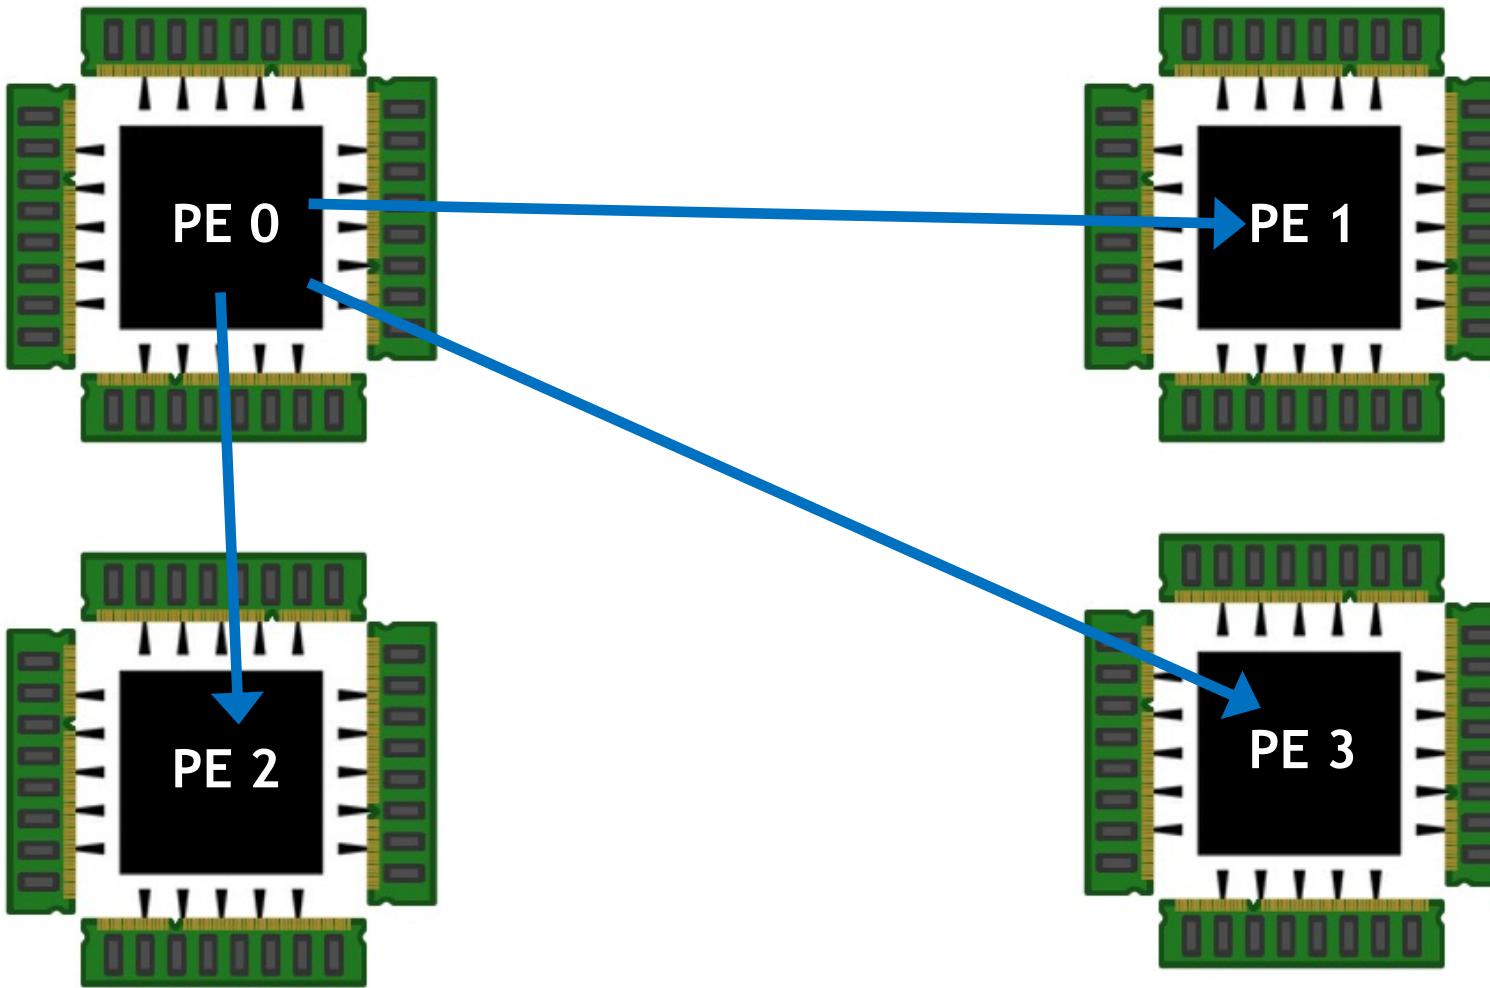
- Simple to understand
- Robust in execution
- Enclose a lot of functionality that solves a lot of common problems
- Avoid deadlocks (the beginner's nightmare)



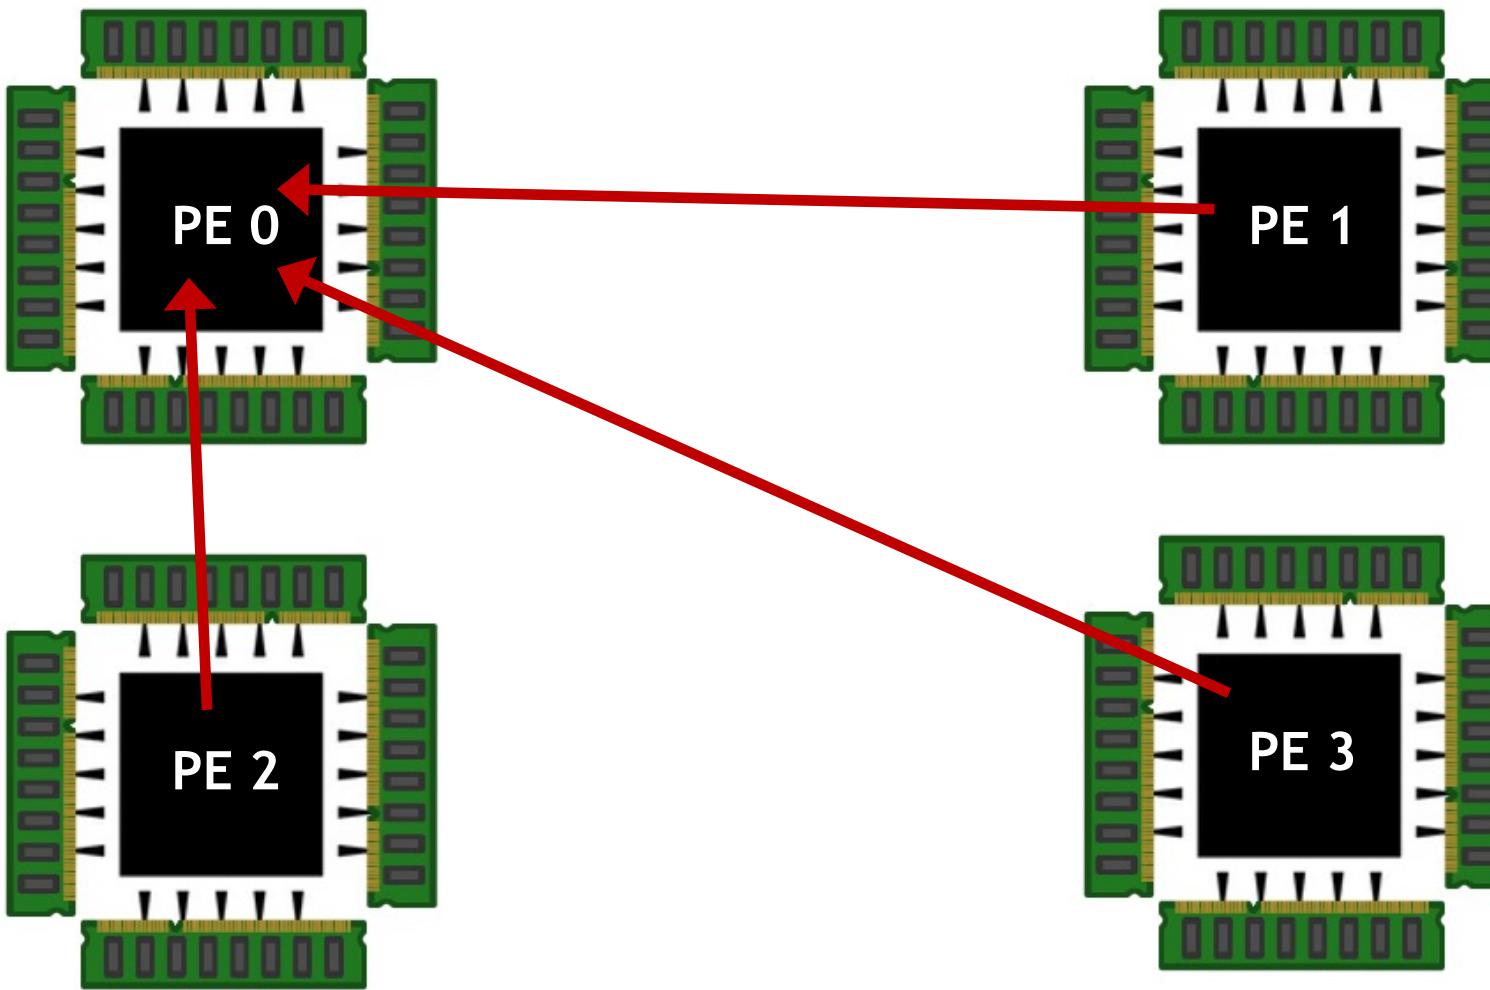
# MPI collectives



# MPI collectives: one-to-all



# MPI collectives: all-to-one



# MPI Broadcast

Copies data from the memory of the root process to the same memory locations for the other processes within the communicator



# MPI Broadcast

Copies data from the memory of the root process to the same memory locations for the other processes within the communicator

## C/C++

```
int ierr = MPI_Bcast( );
```

## Fortran

```
MPI_BCAST( IERR )
```



# MPI Broadcast

Copies data from the memory of the root process to the same memory locations for the other processes within the communicator

## C/C++

```
int ierr = MPI_Bcast(void *send_buf, );
```

Starting address of the data that is being sent

## Fortran

```
MPI_BCAST(SEND_BUF, IERR)
```



# MPI Broadcast

Copies data from the memory of the root process to the same memory locations for the other processes within the communicator

## C/C++

```
int ierr = MPI_Bcast(void *send_buf, int send_count, );
```

Number of elements to be sent

## Fortran

```
MPI_BCAST(SEND_BUF, SEND_COUNT, IERR)
```



# MPI Broadcast

Copies data from the memory of the root process to the same memory locations for the other processes within the communicator

C/C++

```
int ierr = MPI_Bcast(void *send_buf, int send_count, MPI_Datatype dtype, );
```

The data type of the elements being sent

Fortran

```
MPI_BCAST(SEND_BUF, SEND_COUNT, DTYPE, IERR)
```



# MPI Broadcast

Copies data from the memory of the **root process** to the same memory locations for the other processes within the communicator

## C/C++

```
int ierr = MPI_Bcast(void *send_buf, int send_count, MPI_Datatype dtype,  
                     int root_id, );
```

The ID of the root process

## Fortran

```
MPI_BCAST(SEND_BUF, SEND_COUNT, DTYPE, ROOT_ID, IERR)
```



# MPI Broadcast

Copies data from the memory of the root process to the same memory locations for the other processes within the communicator

## C/C++

```
int ierr = MPI_Bcast(void *send_buf, int send_count, MPI_Datatype dtype,  
                     int root_id, MPI_Comm comm);
```

The communicator (group of processes)

## Fortran

```
MPI_BCAST(SEND_BUF, SEND_COUNT, DTYPE, ROOT_ID, COMM, IERR)
```



# MPI Broadcast

Copies data from the memory of the root process to the same memory locations for the other processes within the communicator

## C/C++

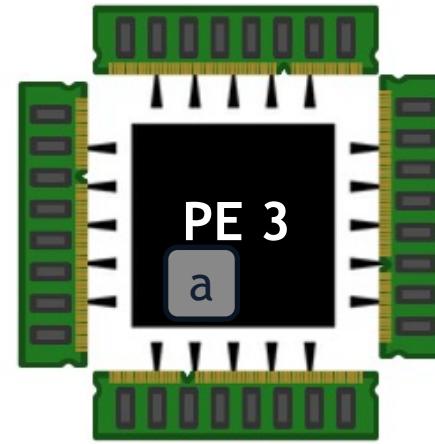
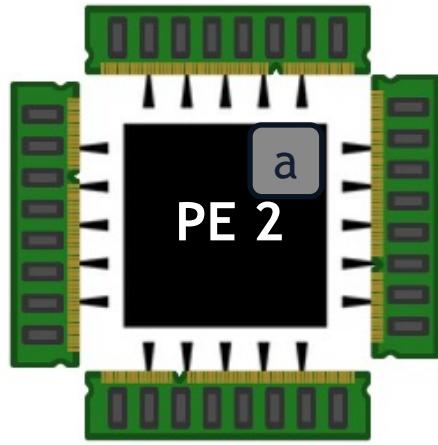
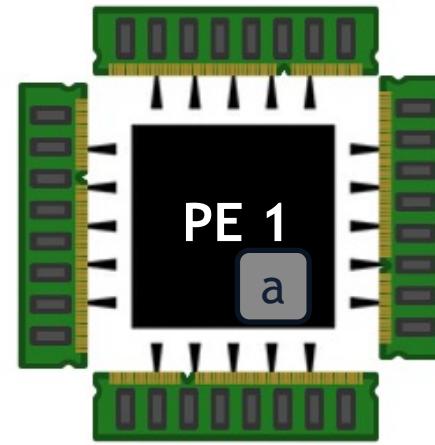
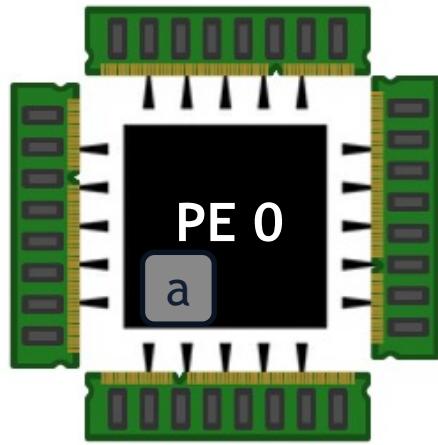
```
int ierr = MPI_Bcast(void *send_buf, int send_count, MPI_Datatype dtype,  
                      int root_id, MPI_Comm comm);
```

## Fortran

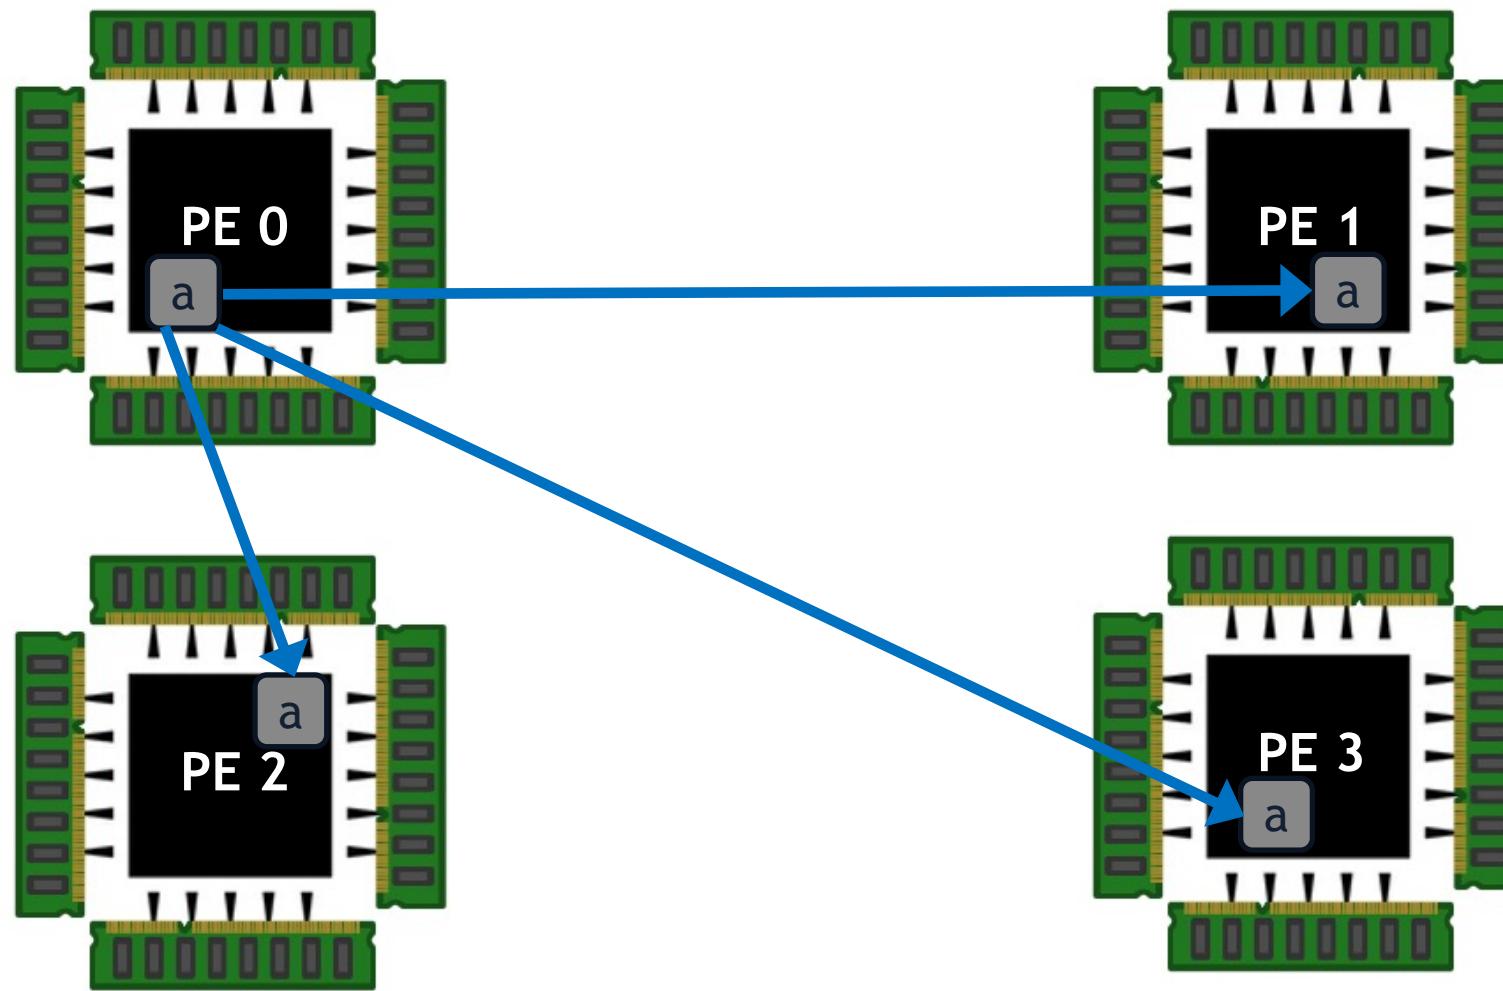
```
MPI_BCAST(SEND_BUF, SEND_COUNT, DTYPE, ROOT_ID, COMM, IERR)
```



# MPI Broadcast



# MPI Broadcast



# MPI Reduce

Collects data from each process, reduces it to a single value via a specified operation, stores the reduced result on the root process



# MPI Reduce

Collects data from each process, reduces it to a single value via a specified operation, stores the reduced result on the root process

## C/C++

```
int ierr = MPI_Reduce( );
```

## Fortran

```
MPI_REDUCE( IERR )
```



# MPI Reduce

Collects data from each process, reduces it to a single value via a specified operation, stores the reduced result on the root process

## C/C++

```
int ierr = MPI_Reduce(void *send_buf, );
```

Starting address of the data that is being sent

## Fortran

```
MPI_REDUCE(SEND_BUF, IERR)
```



# MPI Reduce

Collects data from each process, reduces it to a single value via a specified operation, stores the reduced result on the root process

## C/C++

```
int ierr = MPI_Reduce(void *send_buf, void *recv_buf, );
```

Address of where the resulting operation  
is going to be sent to

## Fortran

```
MPI_REDUCE(SEND_BUF, RECV_BUF, IERR)
```



# MPI Reduce

Collects data from each process, reduces it to a single value via a specified operation, stores the reduced result on the root process

## C/C++

```
int ierr = MPI_Reduce(void *send_buf, void *recv_buf, int count, );
```

Number of elements in the send buffer

## Fortran

```
MPI_REDUCE(SEND_BUF, RECV_BUF, COUNT, IERR)
```



# MPI Reduce

Collects data from each process, reduces it to a single value via a specified operation, stores the reduced result on the root process

## C/C++

```
int ierr = MPI_Reduce(void *send_buf, void *recv_buf, int count,  
MPI_Datatype dtype, );
```

Data type of elements in the send buffer

## Fortran

```
MPI_REDUCE(SEND_BUF, RECV_BUF, COUNT, DTYPE, IERR)
```



# MPI Reduce

Collects data from each process, reduces it to a single value via a specified operation, stores the reduced result on the root process

## C/C++

```
int ierr = MPI_Reduce(void *send_buf, void *recv_buf, int count,  
MPI_Datatype dtype, MPI_Op op, );
```

The reduction operation to be performed

## Fortran

```
MPI_REDUCE(SEND_BUF, RECV_BUF, COUNT, DTYPE, OP, IERR)
```



# MPI Reduce

Collects data from each process, reduces it to a single value via a specified operation, stores the reduced result on the **root process**

## C/C++

```
int ierr = MPI_Reduce(void *send_buf, void *recv_buf, int count,  
MPI_Datatype dtype, MPI_Op op, int root_id, );
```

The ID of the root process

## Fortran

```
MPI_REDUCE(SEND_BUF, RECV_BUF, COUNT, DTYPE, OP, ROOT_ID, IERR)
```



# MPI Reduce

Collects data from each process, reduces it to a single value via a specified operation, stores the reduced result on the root process

## C/C++

```
int ierr = MPI_Reduce(void *send_buf, void *recv_buf, int count,  
MPI_Datatype dtype, MPI_Op op, int root_id, MPI_Comm comm);
```

The communicator (group of processes)

## Fortran

```
MPI_REDUCE(SEND_BUF, RECV_BUF, COUNT, DTYPE, OP, ROOT_ID, COMM, IERR)
```



# MPI Reduce

Collects data from each process, reduces it to a single value via a specified operation, stores the reduced result on the root process

## C/C++

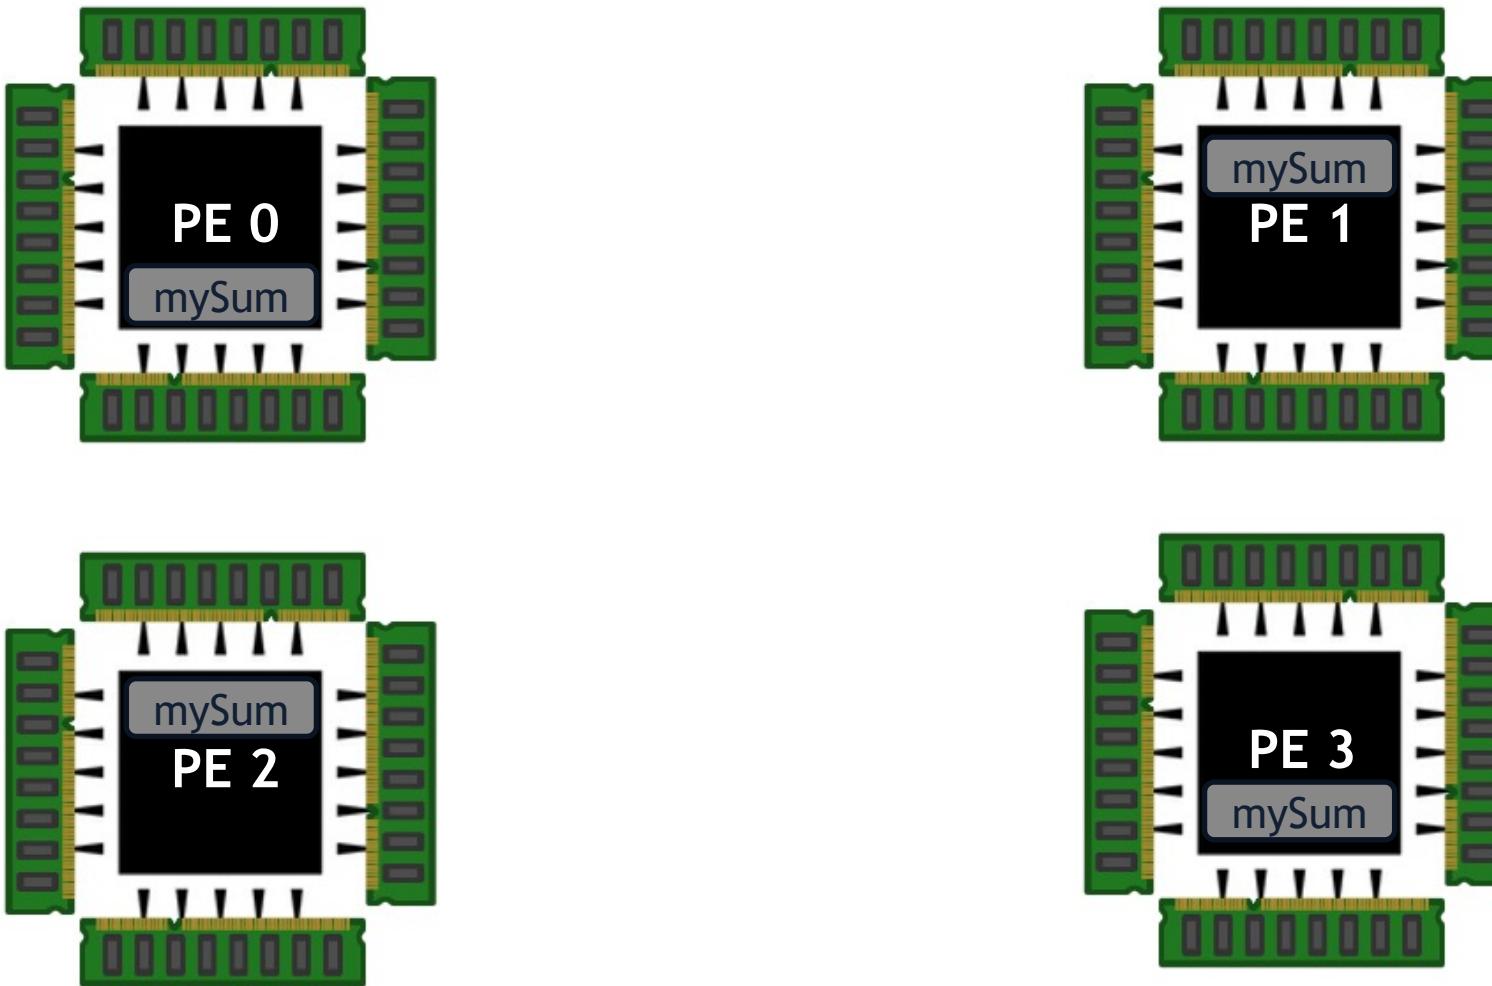
```
int ierr = MPI_Reduce(void *send_buf, void *recv_buf, int count,  
MPI_Datatype dtype, MPI_Op op, int root_id, MPI_Comm comm);
```

## Fortran

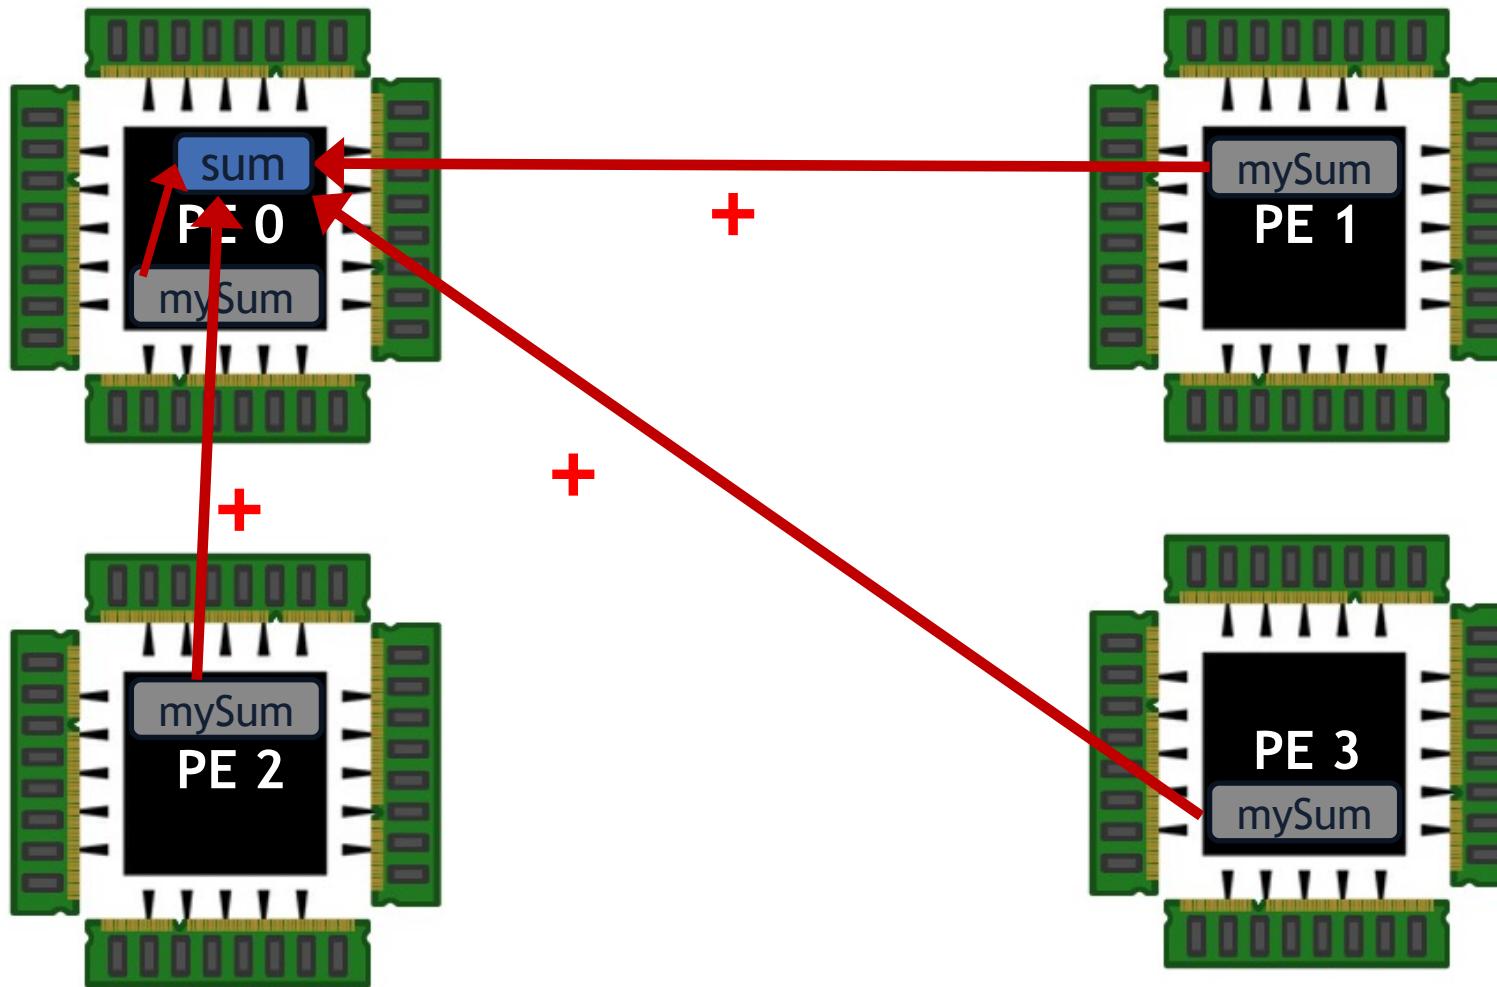
```
MPI_REDUCE(SEND_BUF, RECV_BUF, COUNT, DTYPE, OP, ROOT_ID, COMM, IERR)
```



# MPI Reduce



# MPI Reduce



# MPI library: Datatypes

<https://www mpi-forum.org/docs/mpi-4.0/mpi40-report.pdf> Section 3.2.2

## C/C++

MPI datatype	C datatype
MPI_CHAR	char (treated as printable character)
MPI_SHORT	signed short int
MPI_INT	signed int
MPI_LONG	signed long int
MPI_LONG_LONG_INT	signed long long int
MPI_LONG_LONG (as a synonym)	signed long long int
MPI_SIGNED_CHAR	signed char (treated as integral value)
MPI_UNSIGNED_CHAR	unsigned char (treated as integral value)
MPI_UNSIGNED_SHORT	unsigned short int
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_UNSIGNED_LONG_LONG	unsigned long long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_WCHAR	wchar_t (defined in <stddef.h> (treated as printable character)

## Fortran

MPI datatype	Fortran datatype
MPI_INTEGER	INTEGER
MPI_REAL	REAL
MPI_DOUBLE_PRECISION	DOUBLE PRECISION
MPI_COMPLEX	COMPLEX
MPI_LOGICAL	LOGICAL
MPI_CHARACTER	CHARACTER(1)
MPI_BYTE	
MPI_PACKED	



# MPI library: Datatypes

<https://www mpi-forum.org/docs/mpi-4.0/mpi40-report.pdf> Section 3.2.2

## C/C++

MPI datatype	C datatype
MPI_CHAR	char (treated as printable character)
MPI_SHORT	signed short int
MPI_INT	signed int
MPI_LONG	signed long int
MPI_LONG_LONG_INT	signed long long int
MPI_LONG_LONG (as a synonym)	signed long long int
MPI_SIGNED_CHAR	signed char (treated as integral value)
MPI_UNSIGNED_CHAR	
MPI_UNSIGNED_SHORT	
MPI_UNSIGNED	
MPI_UNSIGNED_LONG	
MPI_UNSIGNED_LONG_LONG	
MPI_FLOAT	
MPI_DOUBLE	
MPI_LONG_DOUBLE	
MPI_WCHAR	
MPI_C_BOOL	_Bool
MPI_INT8_T	int8_t
MPI_INT16_T	int16_t
MPI_INT32_T	int32_t
MPI_INT64_T	int64_t
MPI_UINT8_T	uint8_t
MPI_UINT16_T	uint16_t
MPI_UINT32_T	uint32_t
MPI_UINT64_T	uint64_t
MPI_C_COMPLEX	float _Complex
MPI_C_FLOAT_COMPLEX (as a synonym)	float _Complex
MPI_C_DOUBLE_COMPLEX	double _Complex
MPI_C_LONG_DOUBLE_COMPLEX	long double _Complex
MPI_BYTE	
MPI_PACKED	

## Fortran

MPI datatype	Fortran datatype
MPI_INTEGER	INTEGER
MPI_REAL	REAL
MPI_DOUBLE_PRECISION	DOUBLE PRECISION
MPI_COMPLEX	COMPLEX
MPI_LOGICAL	LOGICAL
MPI_CHARACTER	CHARACTER(1)
MPI_BYTE	
MPI_PACKED	



# MPI library: reduce operations

<https://www mpi-forum.org/docs/mpi-4.0/mpi40-report.pdf> Section 6.9.1

Name	Meaning
MPI_MAX	maximum
MPI_MIN	minimum
MPI_SUM	sum
MPI_PROD	product
MPI_LAND	logical and
MPI_BAND	bit-wise and
MPI_LOR	logical or
MPI_BOR	bit-wise or
MPI_LXOR	logical exclusive or (xor)
MPI_BXOR	bit-wise exclusive or (xor)
MPI_MAXLOC	max value and location
MPI_MINLOC	min value and location



# MPI library: communicators

`MPI_Comm MPI_COMM_WORLD`



Every process in  
your running  
application

You can construct custom communicators if you wish...

<https://www mpi-forum.org/docs/mpi-4.0/mpi40-report.pdf> Chapter 7

# MPI library: ubiquitous functions

These are the ones that you will **always** use in an MPI program

C/C++

Fortran



# MPI library: ubiquitous functions

These are the ones that you will **always** use in an MPI program

C/C++

```
int ierr = MPI_Init(&argc, &argv)
```

Fortran

```
MPI_INIT(IERR)
```

Initializes the MPI environment



# MPI library: ubiquitous functions

These are the ones that you will **always** use in an MPI program

C/C++

```
int ierr = MPI_Init(&argc, &argv)  
  
int ierr = MPI_Finalize()
```

Fortran

```
MPI_INIT(IERR)  
  
MPI_FINALIZE(IERR)
```

Terminates the MPI environment



# MPI library: ubiquitous functions

These are the ones that you will **always** use in an MPI program

## C/C++

```
int ierr = MPI_Init(&argc, &argv)
```

```
int ierr = MPI_Finalize()
```

```
int ierr = MPI_Comm_size(MPI_Comm comm,  
&nprocs)
```

## Fortran

```
MPI_INIT(IERR)
```

```
MPI_FINALIZE(IERR)
```

```
MPI_COMM_SIZE(COMM, NPROCS, IERR)
```

Gives the number of processes  
in the communicator



# MPI library: ubiquitous functions

These are the ones that you will **always** use in an MPI program

## C/C++

```
int ierr = MPI_Init(&argc, &argv)
```

```
int ierr = MPI_Finalize()
```

```
int ierr = MPI_Comm_size(MPI_Comm comm,  
&nprocs)
```

```
int ierr = MPI_Comm_rank(MPI_Comm comm,  
&my_id)
```

## Fortran

```
MPI_INIT(IERR)
```

```
MPI_FINALIZE(IERR)
```

```
MPI_COMM_SIZE(COMM, NPROCS, IERR)
```

```
MPI_COMM_RANK(COMM, MY_ID, IERR)
```

Gives the ID of each process in  
the communicator



# MPI library: ubiquitous functions

These are the ones that you will **always** use in an MPI program

## C/C++

```
int ierr = MPI_Init(&argc, &argv)  
  
int ierr = MPI_Finalize()  
  
int ierr = MPI_Comm_size(MPI_Comm comm,  
    &nprocs)  
  
int ierr = MPI_Comm_rank(MPI_Comm comm,  
    &my_id)
```

## Fortran

```
MPI_INIT(IERR)  
  
MPI_FINALIZE(IERR)  
  
MPI_COMM_SIZE(COMM, NPROCS, IERR)  
  
MPI_COMM_RANK(COMM, MY_ID, IERR)
```



# Outline

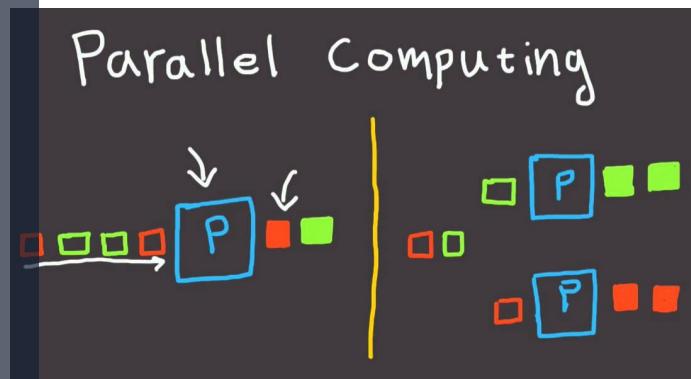
## Part A: Parallel Computing Quick Review (10 min)

- Task and Data parallelism
- Parallel computers
- Threads, processes and more



## Part B: MPI Intro (60 min)

- Core concepts
- Message passing
- Execution



## Part C: Exercises (45+ min)

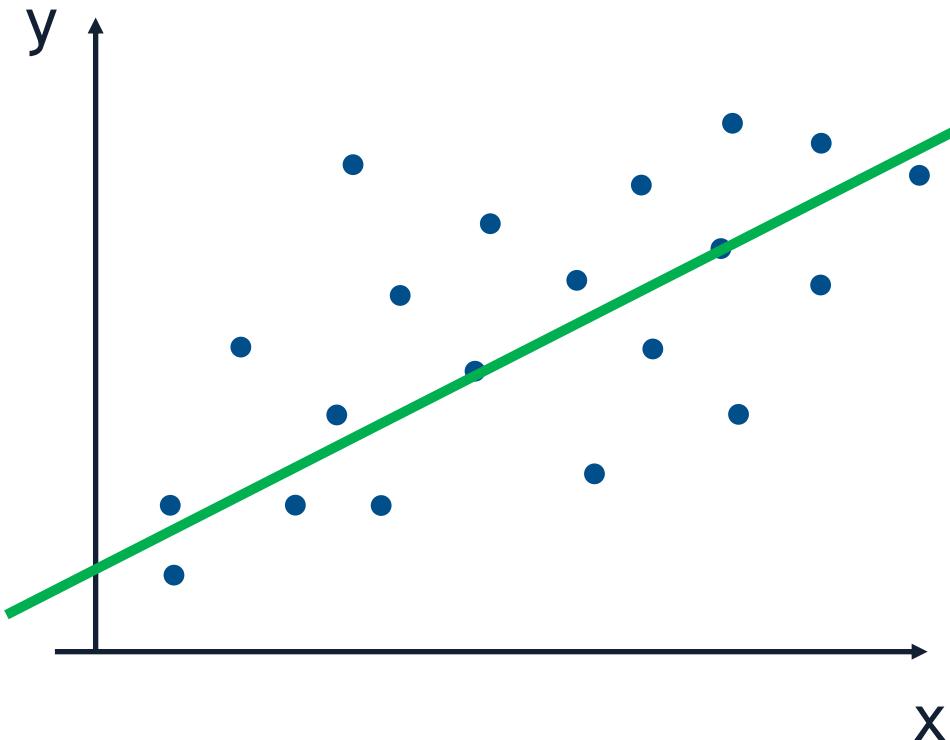
- Ex 1: Hello World
- Ex 2: Domain Decomposition
- Ex 3: Parallel Linear Regression

# Exercise 1 – MPI Hello World!



[1. module load intel/18.0
2. Parallelize the Hello World serial code and make each PE print their ID with the hello message](https://github.com/babreu-ncsa/IntroToMPI/tree/uiuc-icc/Exercises>Hello</a></p></div><div data-bbox=)

# Linear Regression 101



$x$ : inputs (features, control variable)  
 $y$ : outputs (observations, response variable)

Model

$$y = a \cdot x + b$$

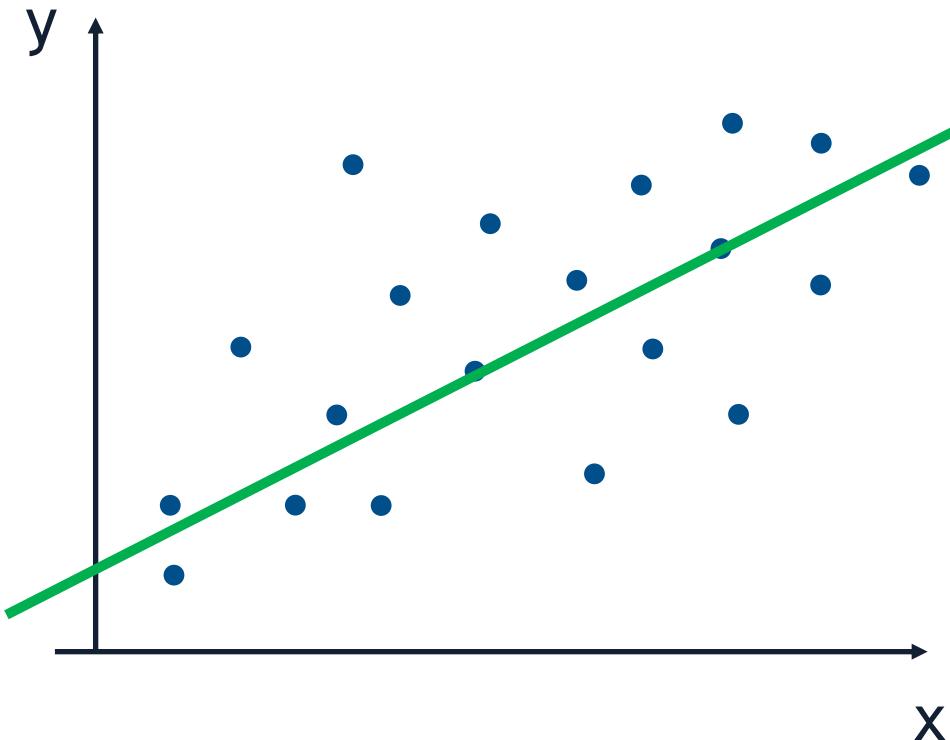
Parameters

Fitting

Find the “best” (a, b) pair

Inference

# Linear Regression 101



$x$ : inputs (features, control variable)  
 $y$ : outputs (observations, response variable)

Model

$$y = a \cdot x + b$$

Parameters

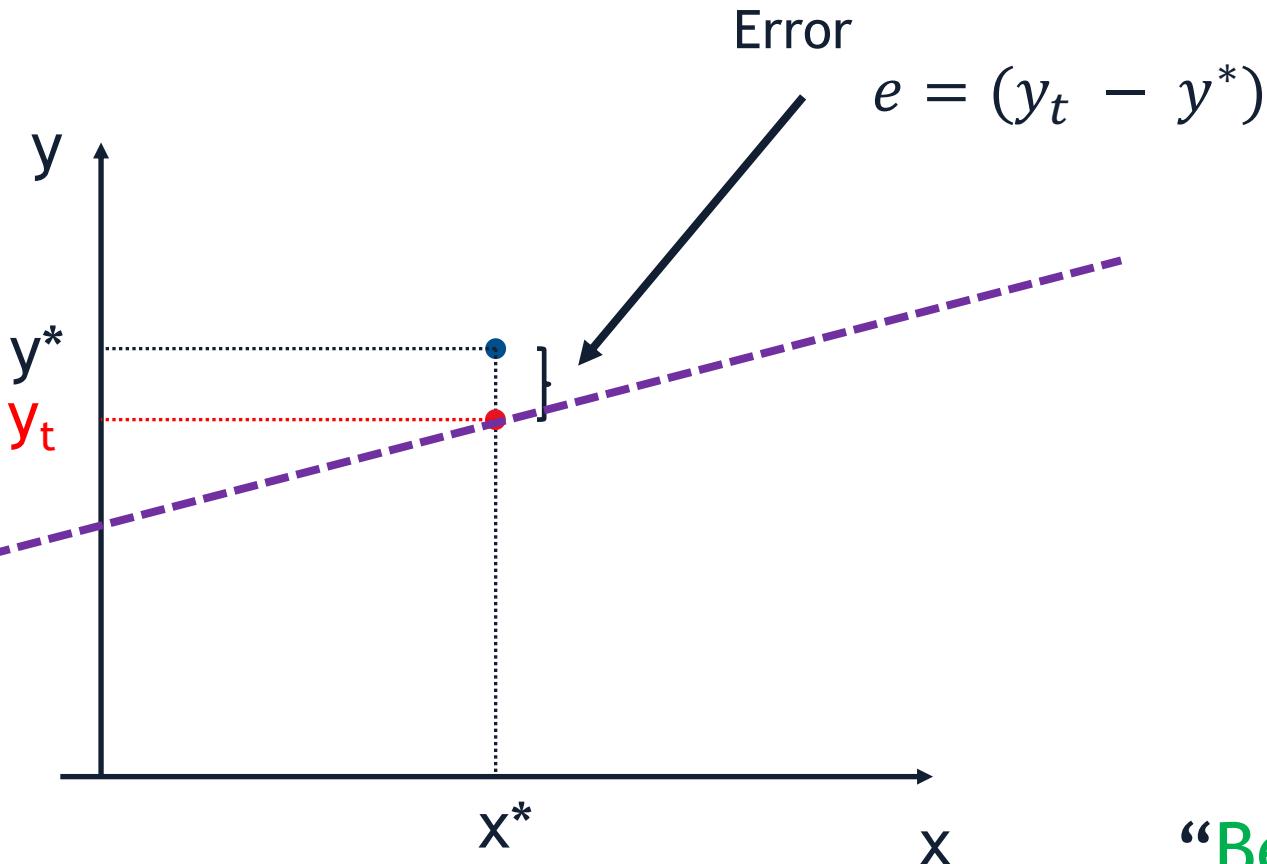
Fitting

Find the “best” (a, b) pair

?

Inference

# Linear Regression 101



$$y_t = a_t \cdot x^* + b_t$$

**Mean Squared Error (MSE)**

$$MSE = \frac{1}{N} \sum_{i=1}^N e_i^2$$

“Best”  $(a, b)$   $\rightarrow$  Smallest MSE

# LinReg with MPI



<https://github.com/babreu-ncsa/IntroToMPI/tree/uiuc-icc/Exercises/LinearRegression>

# Serial LinReg

```
int main()
{
    // parameter map
    int na = 10, nb = 10;      // number of points for each parameter in grid space
    double da = 0.1, db = 0.1; // grid spacing in each direction
    vector<double> a, b;      // parameters in each direction
```



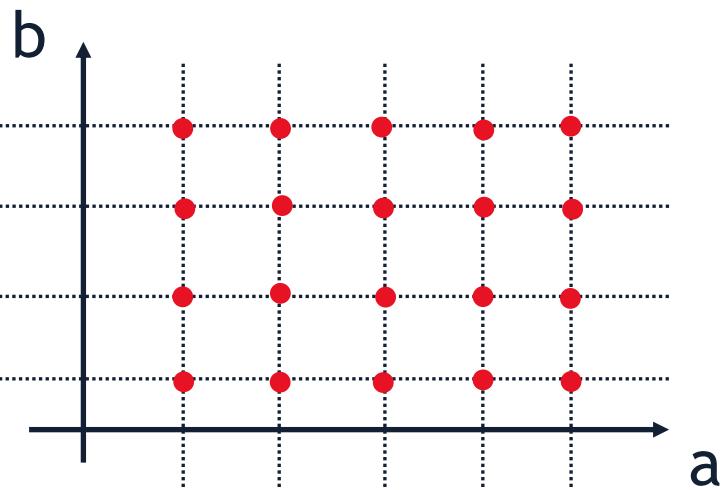
# Serial LinReg

```
int main()
{
    // parameter map
    int na = 10, nb = 10;      // number of points for each parameter in grid space
    double da = 0.1, db = 0.1; // grid spacing in each direction
    vector<double> a, b;      // parameters in each direction
```



# Serial LinReg

```
int main()
{
    // parameter map
    int na = 10, nb = 10;      // number of points for each parameter in grid space
    double da = 0.1, db = 0.1; // grid spacing in each direction
    vector<double> a, b;      // parameters in each direction
```



```
// 1. Build parameter map - square grid in (a,b) space
for (i = 0; i < na; i++)
{
    a.push_back((i + 1) * da);
}
for (j = 0; j < nb; j++)
{
    b.push_back((j + 1) * db);
}
```



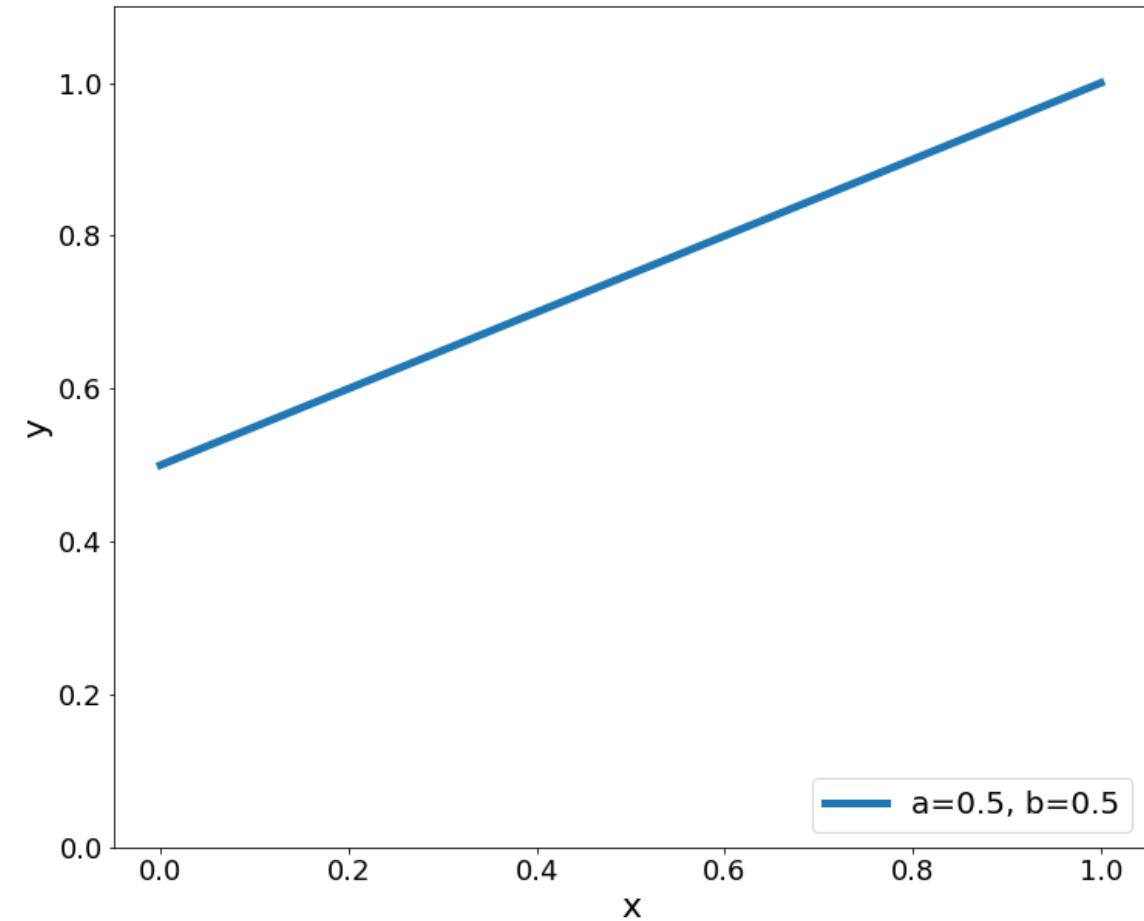
# Serial LinReg

```
// target straight line
int n = 1 << 23;           // number of "data" points
vector<double> x;         // control variable
double dx;                 // control variable spacing
vector<double> y;         // response variable
double at = 0.5, bt = 0.5; // target parameters
```



# Serial LinReg

```
// target straight line
int n = 1 << 23;           // number of "data" points
vector<double> x;          // control variable
double dx;                  // control variable spacing
vector<double> y;          // response variable
double at = 0.5, bt = 0.5; // target parameters
```



# Serial LinReg

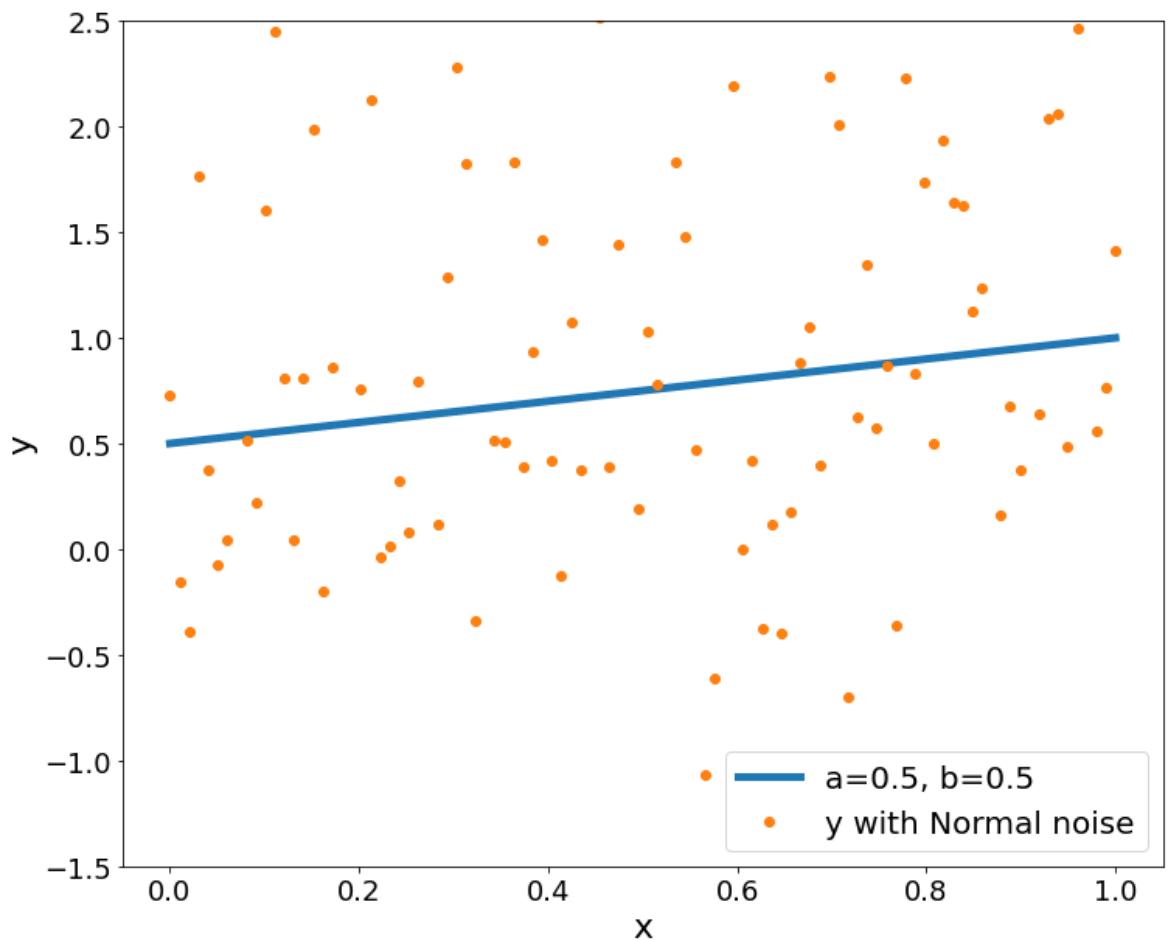
```
// 2. Build points to fit (dataset)
dx = 1.0 / double(n); // we'll make x go from 0 to 1
for (i = 0; i < n; i++)
{
    xp = i * dx;
    yp = at * xp + bt;
    rand1 = (double)rand() / RAND_MAX;
    rand2 = (double)rand() / RAND_MAX;
    z = sqrt(-2.0 * log(rand1)) * cos(2.0 * pi * rand2);
    yp = yp + z;
    x.push_back(xp);
    y.push_back(yp);
}
```



# Serial LinReg

```
// 2. Build points to fit (dataset)
dx = 1.0 / double(n); // we'll make x go from 0 to 1
for (i = 0; i < n; i++)
{
    xp = i * dx;
    yp = at * xp + bt;
    rand1 = (double)rand() / RAND_MAX;
    rand2 = (double)rand() / RAND_MAX;
    z = sqrt(-2.0 * log(rand1)) * cos(2.0 * pi * rand2);
    yp = yp + z;
    x.push_back(xp);
    y.push_back(yp);
}
```

[https://colab.research.google.com/drive/124oGew9s\\_GA3vKkLjr2-pRepEk3MdKze?usp=sharing](https://colab.research.google.com/drive/124oGew9s_GA3vKkLjr2-pRepEk3MdKze?usp=sharing)



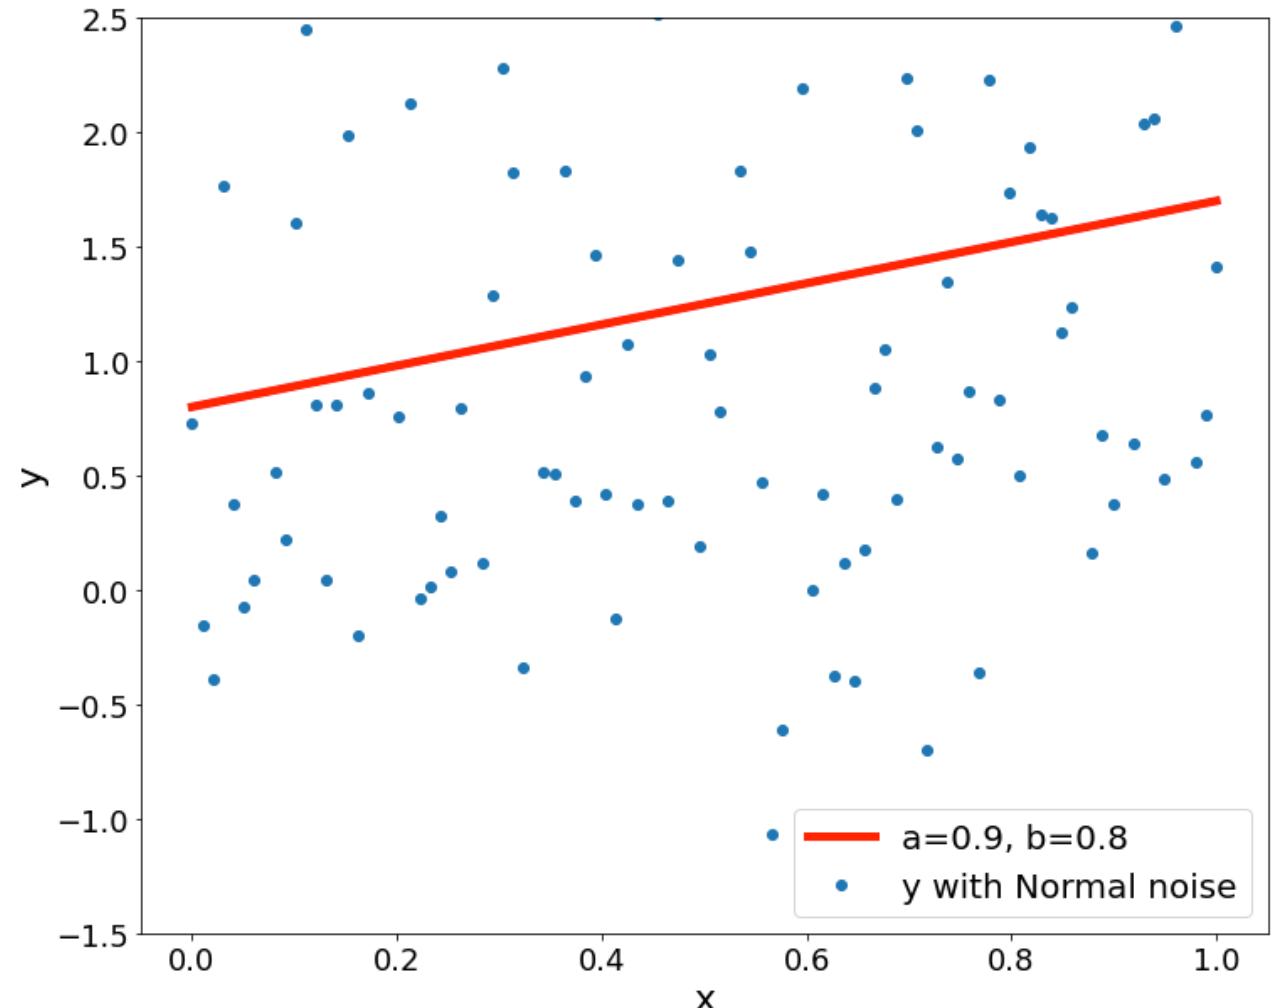
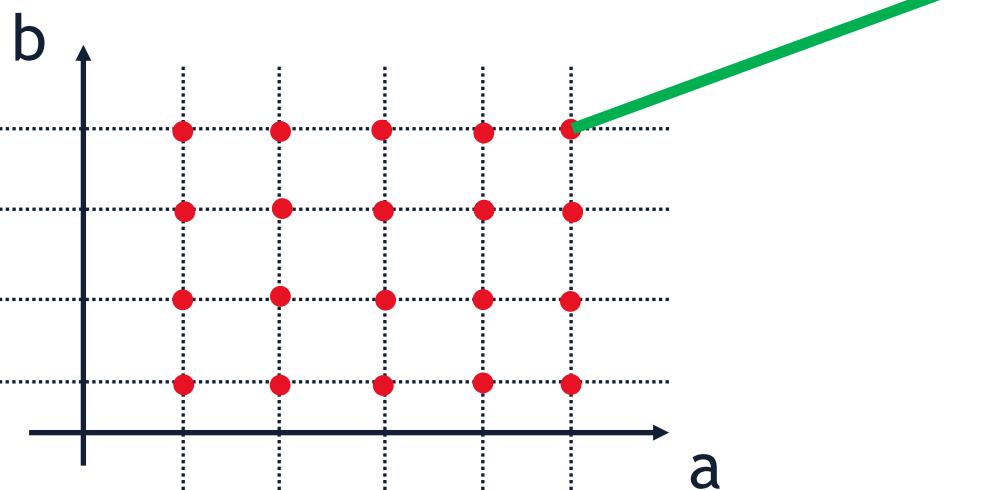
# Serial LinReg

```
// 3. Explore parameter space
for (i = 0; i < na; i++)
{
    as = a[i];
    for (j = 0; j < nb; j++)
    {
        bs = b[j];
        rss = 0.0;
        for (k = 0; k < n; k++)
        {
            ys = as * x[k] + bs;
            rss = rss + pow((ys - y[k]), 2.0);
        }
        mse.push_back(rss / n);
        cout << "(a,b) = (" << a[i] << "," << b[j] << ")"
    }
}
```



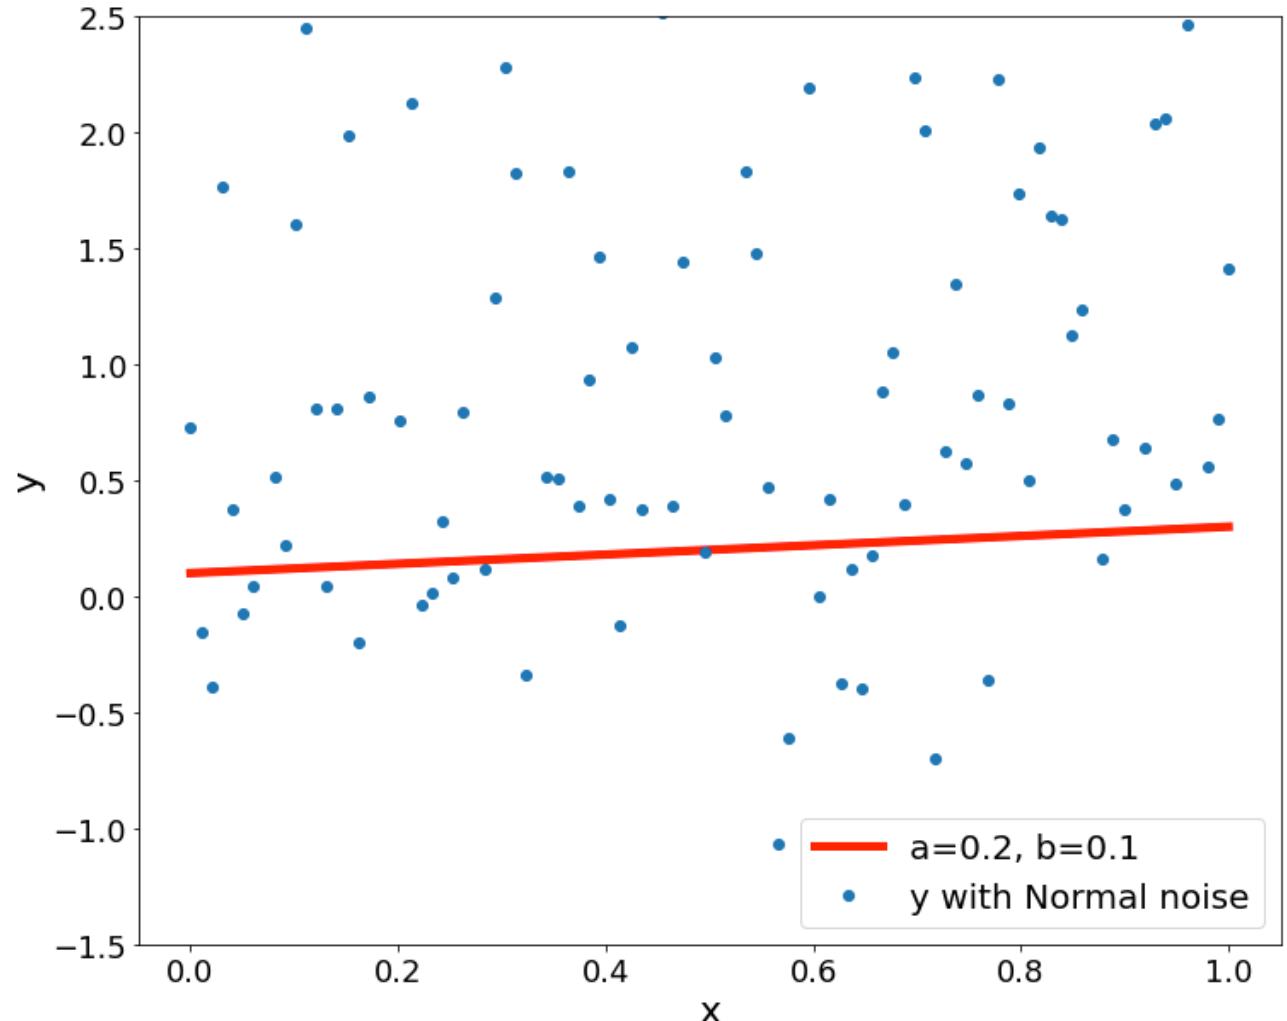
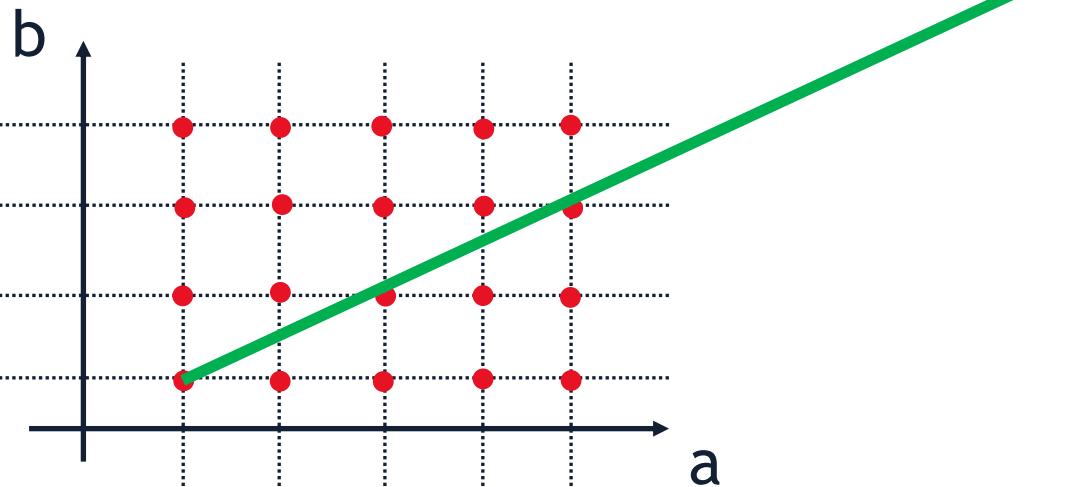
# Serial LinReg

```
// 3. Explore parameter space
for (i = 0; i < na; i++)
{
    as = a[i];
    for (j = 0; j < nb; j++)
    {
        bs = b[j];
        rss = 0.0;
        for (k = 0; k < n; k++)
        {
            ys = as * x[k] + bs;
            rss = rss + pow((ys - y[k]), 2.0);
        }
        mse.push_back(rss / n);
        cout << "(a,b) = (" << a[i] << "," << b[j] << ")"
    }
}
```



# Serial LinReg

```
// 3. Explore parameter space
for (i = 0; i < na; i++)
{
    as = a[i];
    for (j = 0; j < nb; j++)
    {
        bs = b[j];
        rss = 0.0;
        for (k = 0; k < n; k++)
        {
            ys = as * x[k] + bs;
            rss = rss + pow((ys - y[k]), 2.0);
        }
        mse.push_back(rss / n);
        cout << "(a,b) = (" << a[i] << "," << b[j] << ")"
    }
}
```



# Serial LinReg

```
// 4. Look for best combination of (a,b)
counter = 0;
best_mse = mse[0];
best_i = 0;
best_j = 0;
for (i = 0; i < na; i++)
{
    for (j = 0; j < nb; j++)
    {
        if (mse[counter] < best_mse)
        {
            best_mse = mse[counter];
            best_i = i;
            best_j = j;
        }
        counter++;
    }
}
```



# Serial LinReg

```
// 4. Look for best combination of (a,b)
counter = 0;
best_mse = mse[0];
best_i = 0;
best_j = 0;
for (i = 0; i < na; i++)
{
    for (j = 0; j < nb; j++)
    {
        if (mse[counter] < best_mse)
        {
            best_mse = mse[counter];
            best_i = i;
            best_j = j;
        }
        counter++;
    }
}
```

$$MSE = \frac{1}{N} \sum_{i=1}^N e_i^2$$

“Best” (a, b)  Smallest MSE

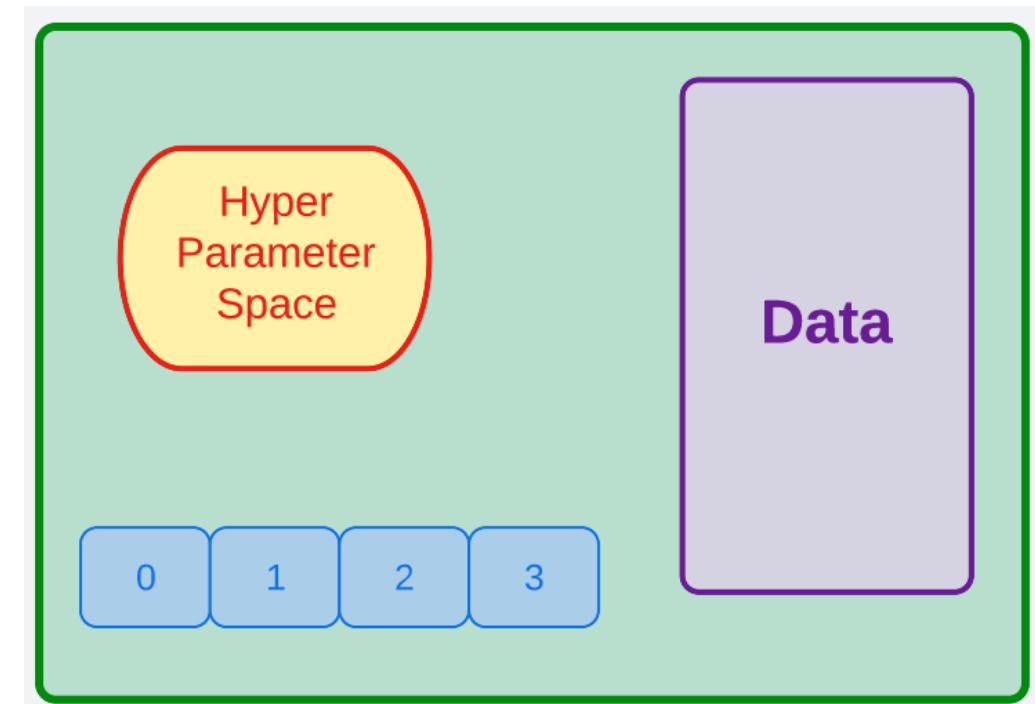
# Exercise 2 – Domain decomposition

It's time to parallelize the Linear Regression using MPI!

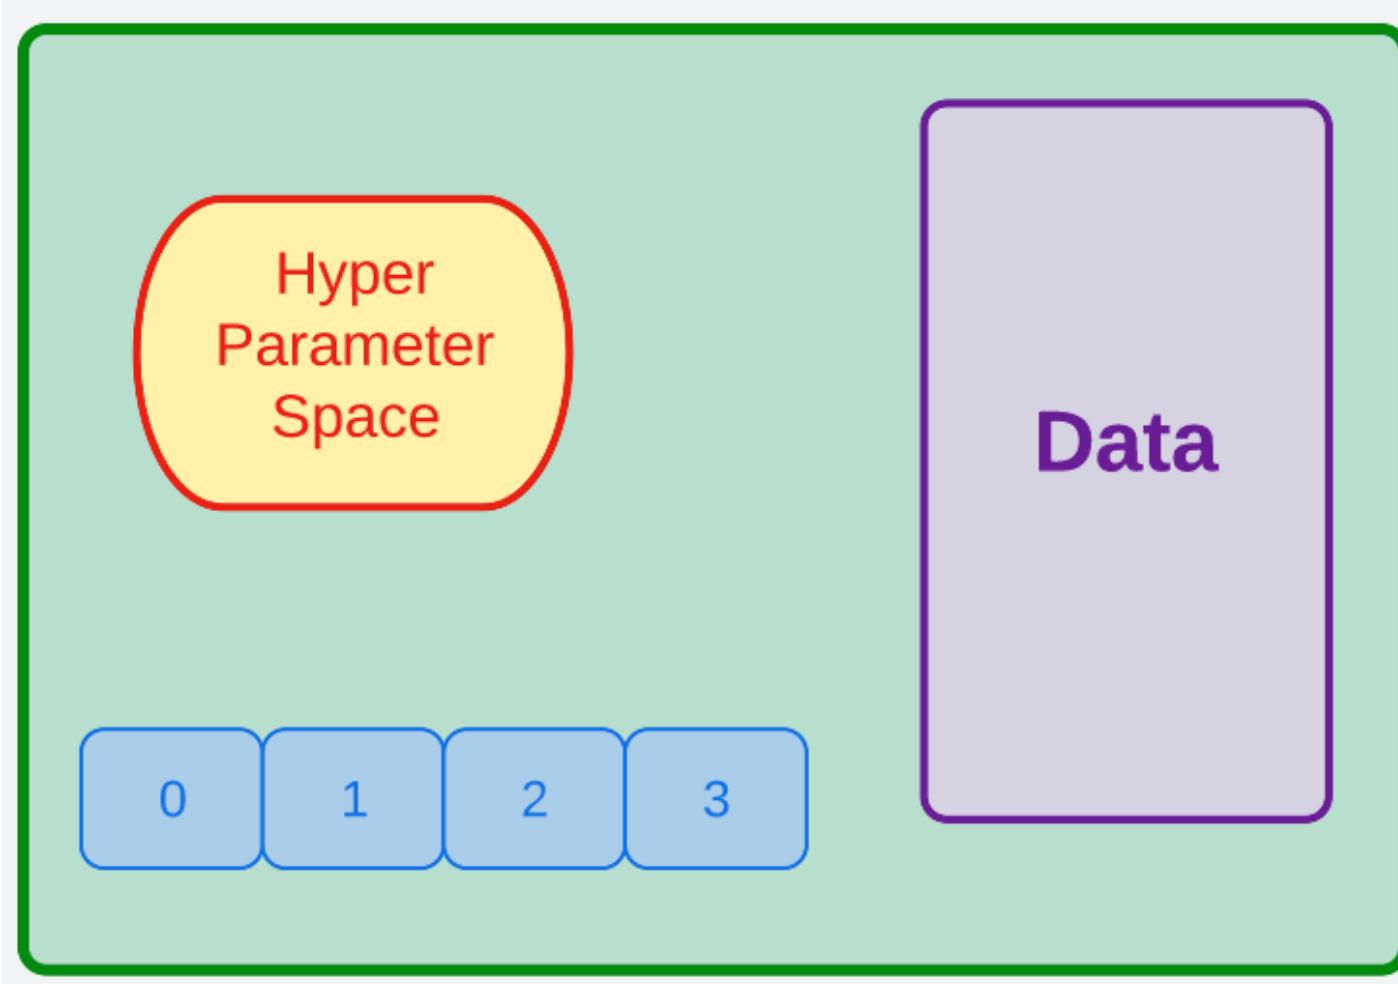
The first step to parallelization is called domain decomposition

- What each process will be doing?
- How do I distribute the tasks to each of them?
- How do I slice my entire dataset so that each process works on a particular part of it?
- Do I want my code to just run fast, or do I want to solve bigger and bigger problems?

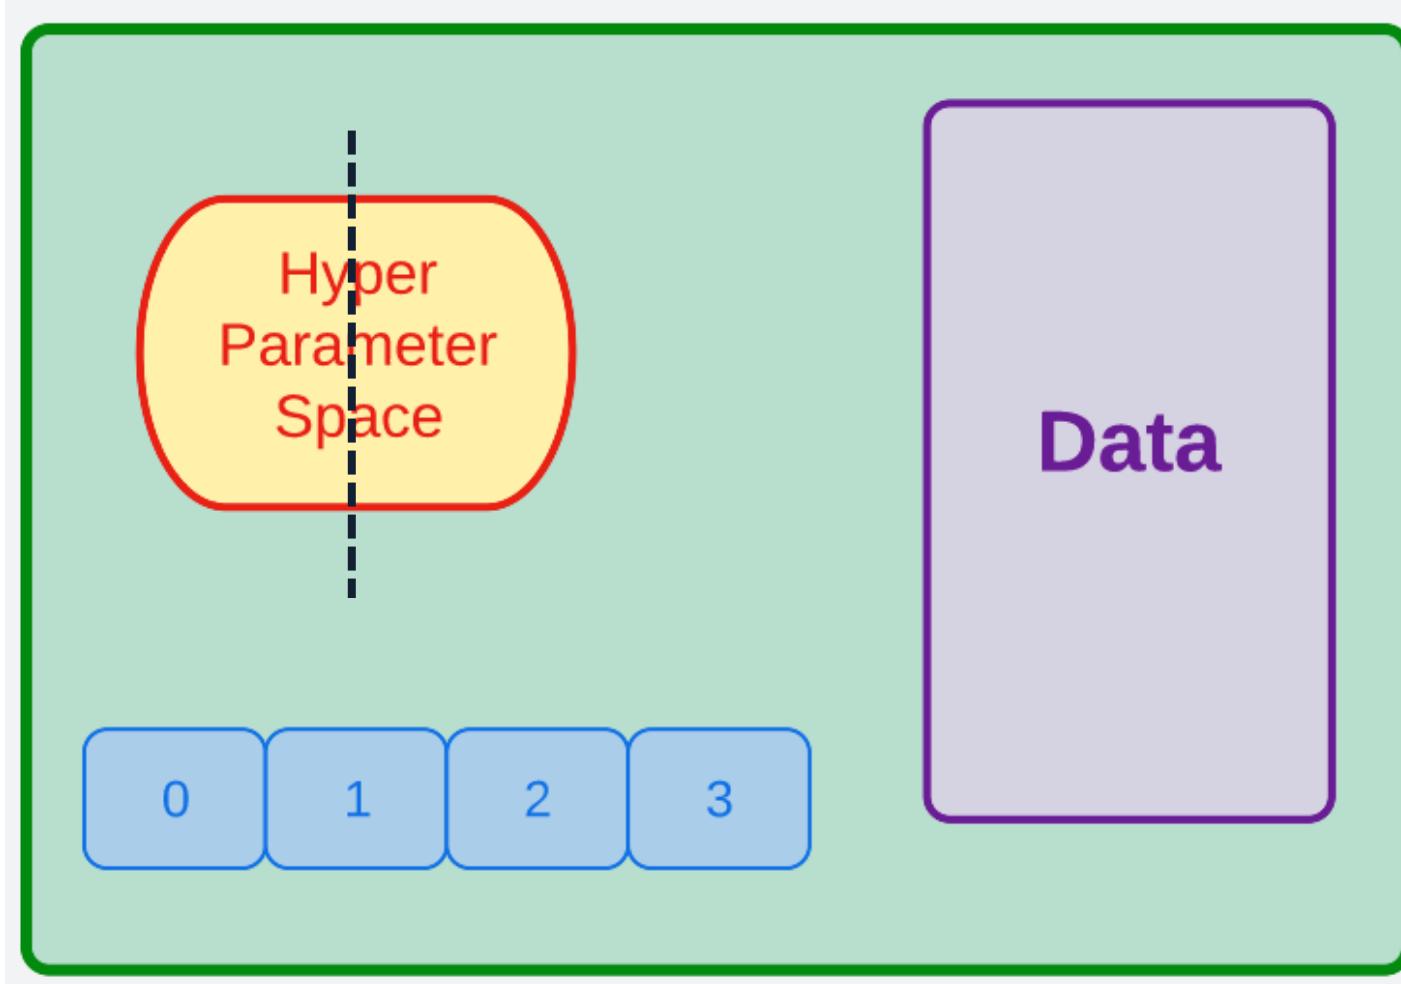
[https://docs.google.com/presentation/d/13skCf6Wf\\_cB0bbF1Claz\\_z1ZOQSDQ5KxKqCA9Q5XOo/edit?usp=sharing](https://docs.google.com/presentation/d/13skCf6Wf_cB0bbF1Claz_z1ZOQSDQ5KxKqCA9Q5XOo/edit?usp=sharing)



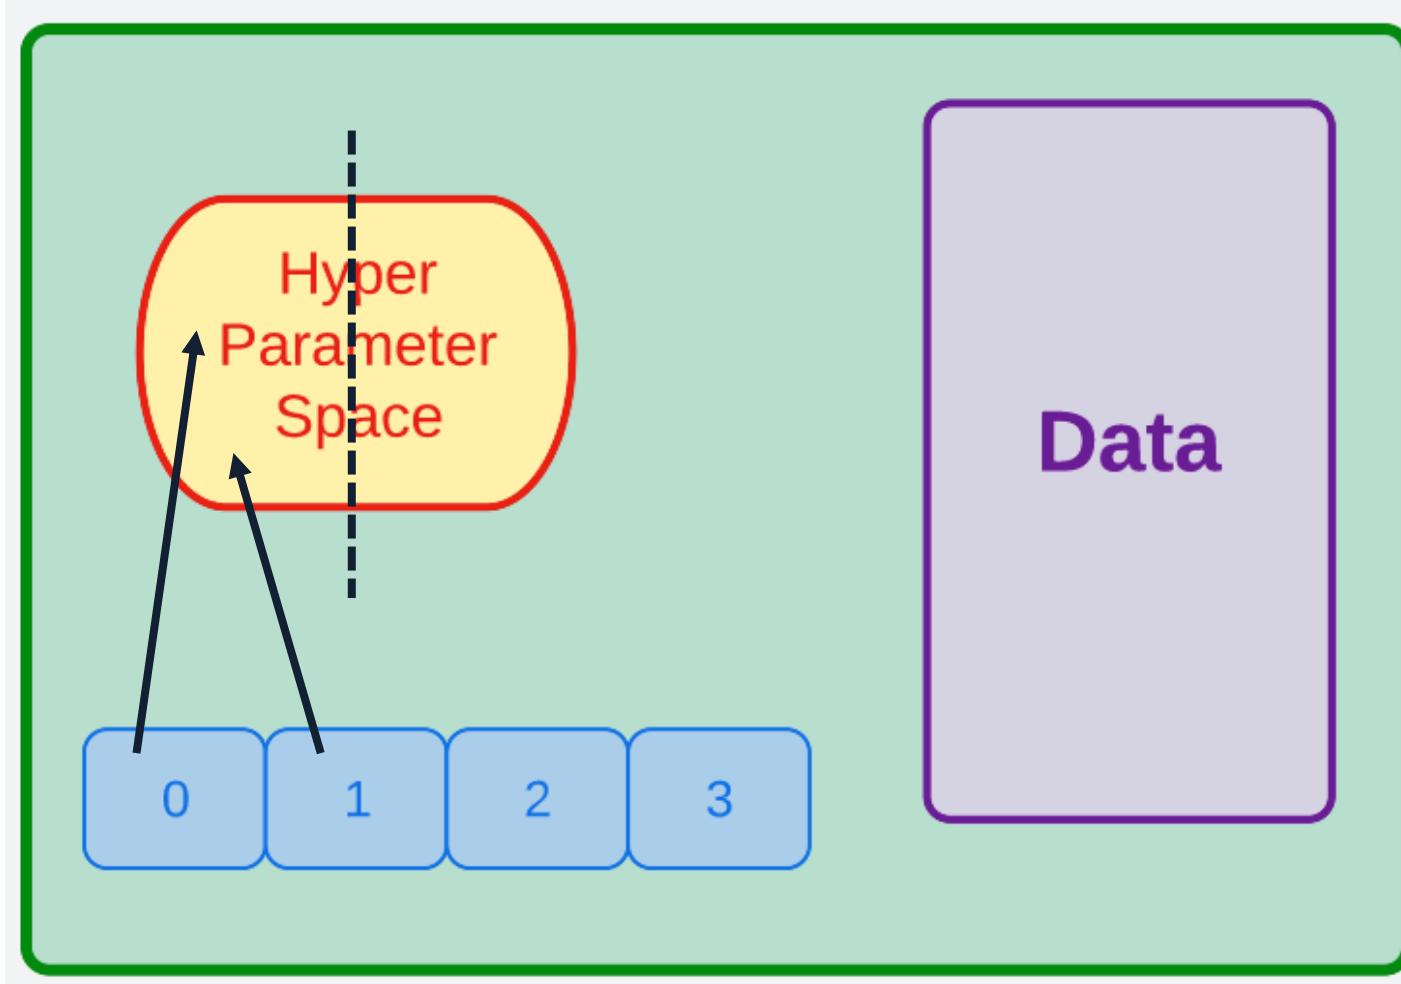
# Decomposition example



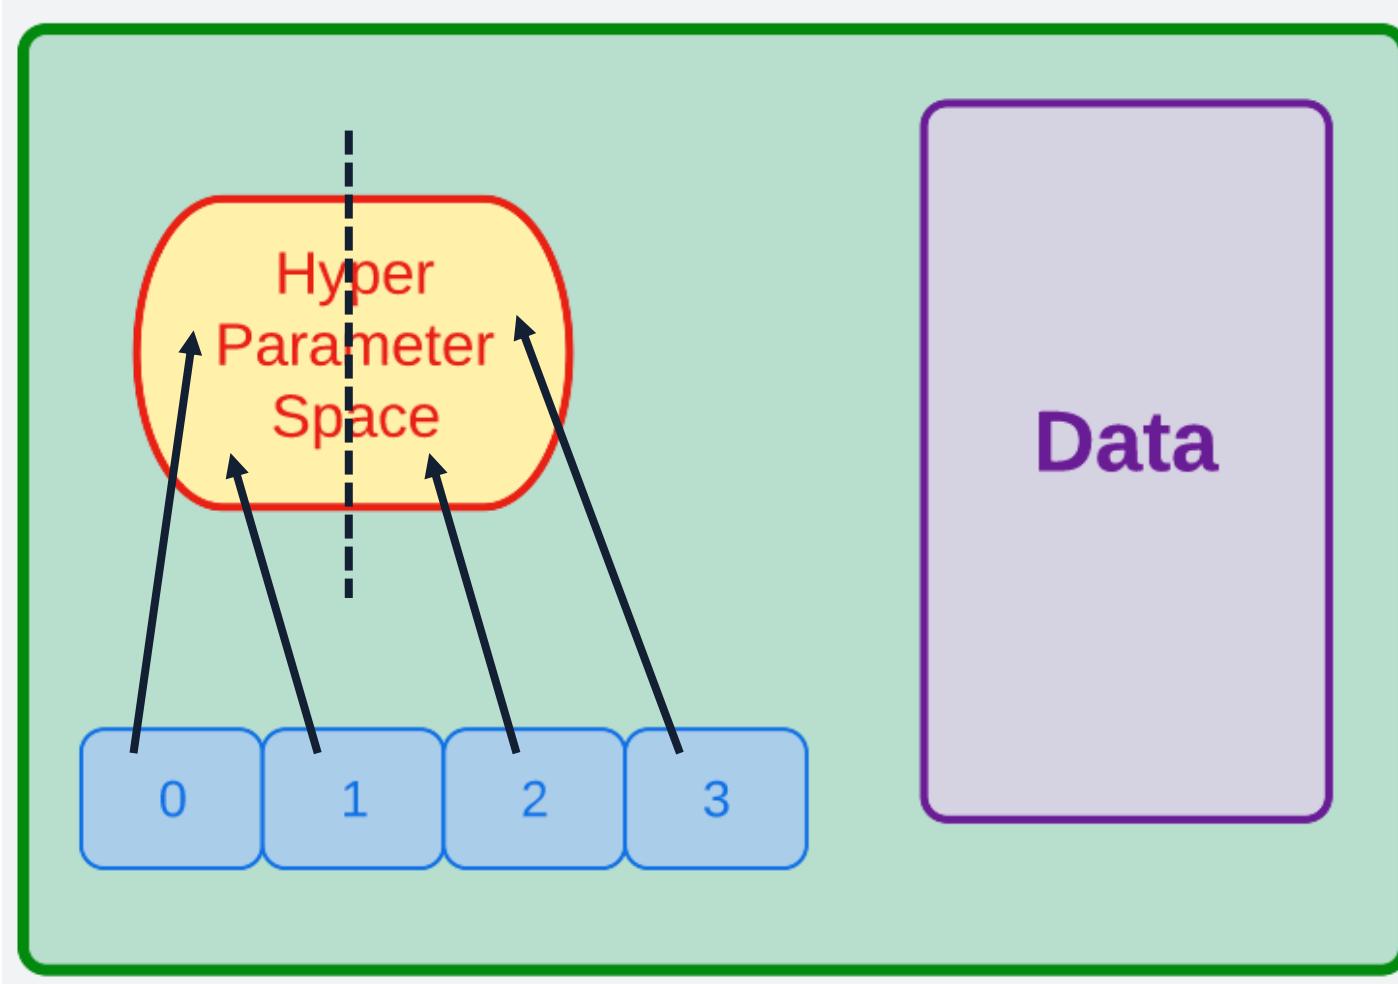
# Decomposition example



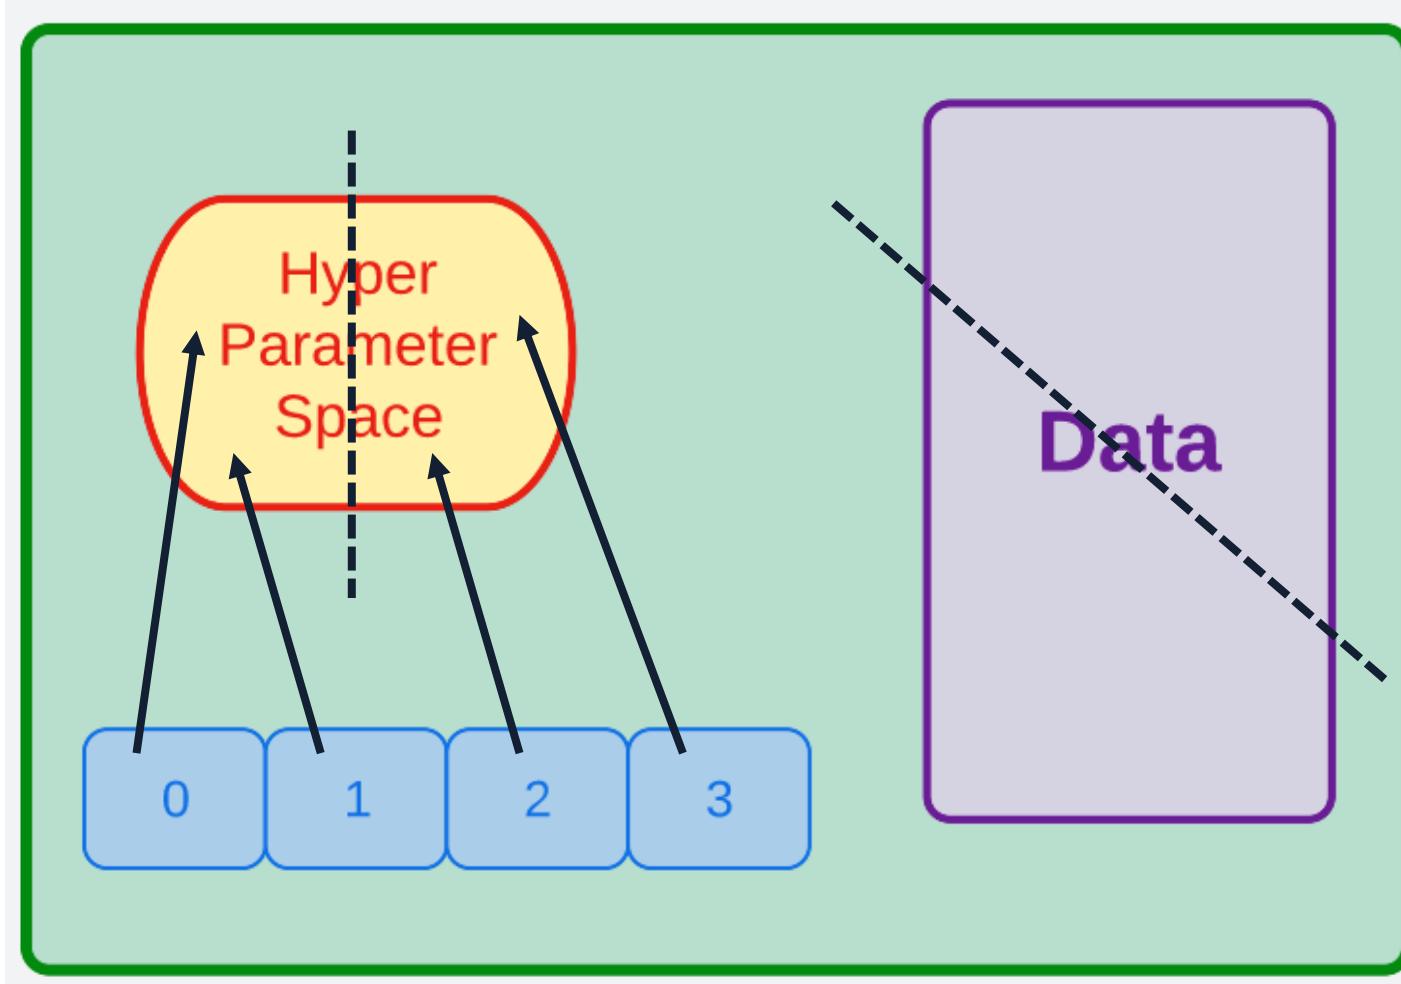
# Decomposition example



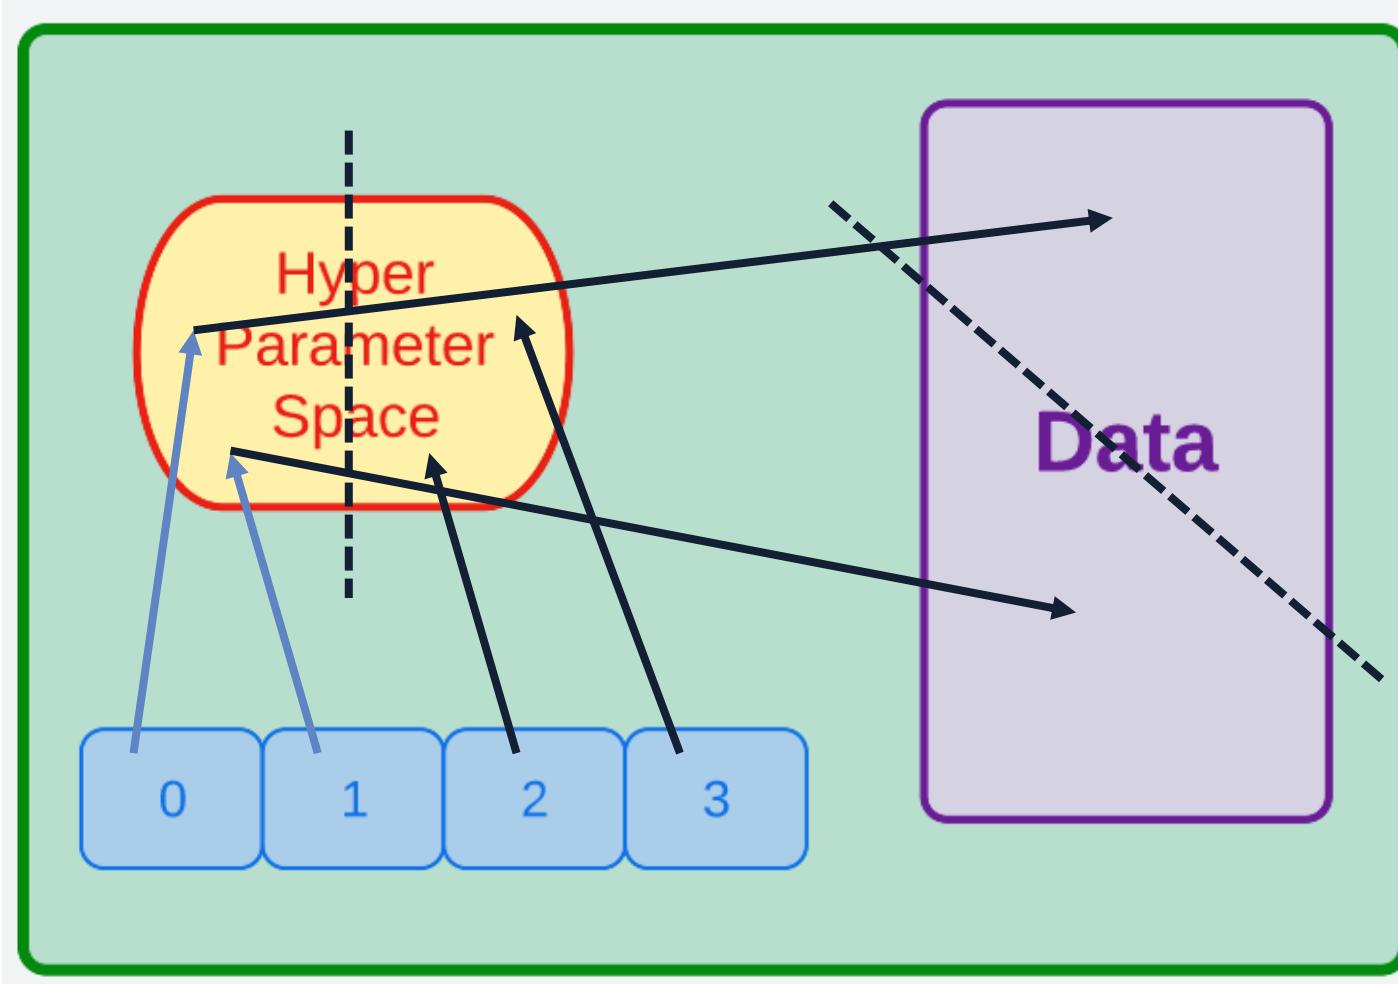
# Decomposition example



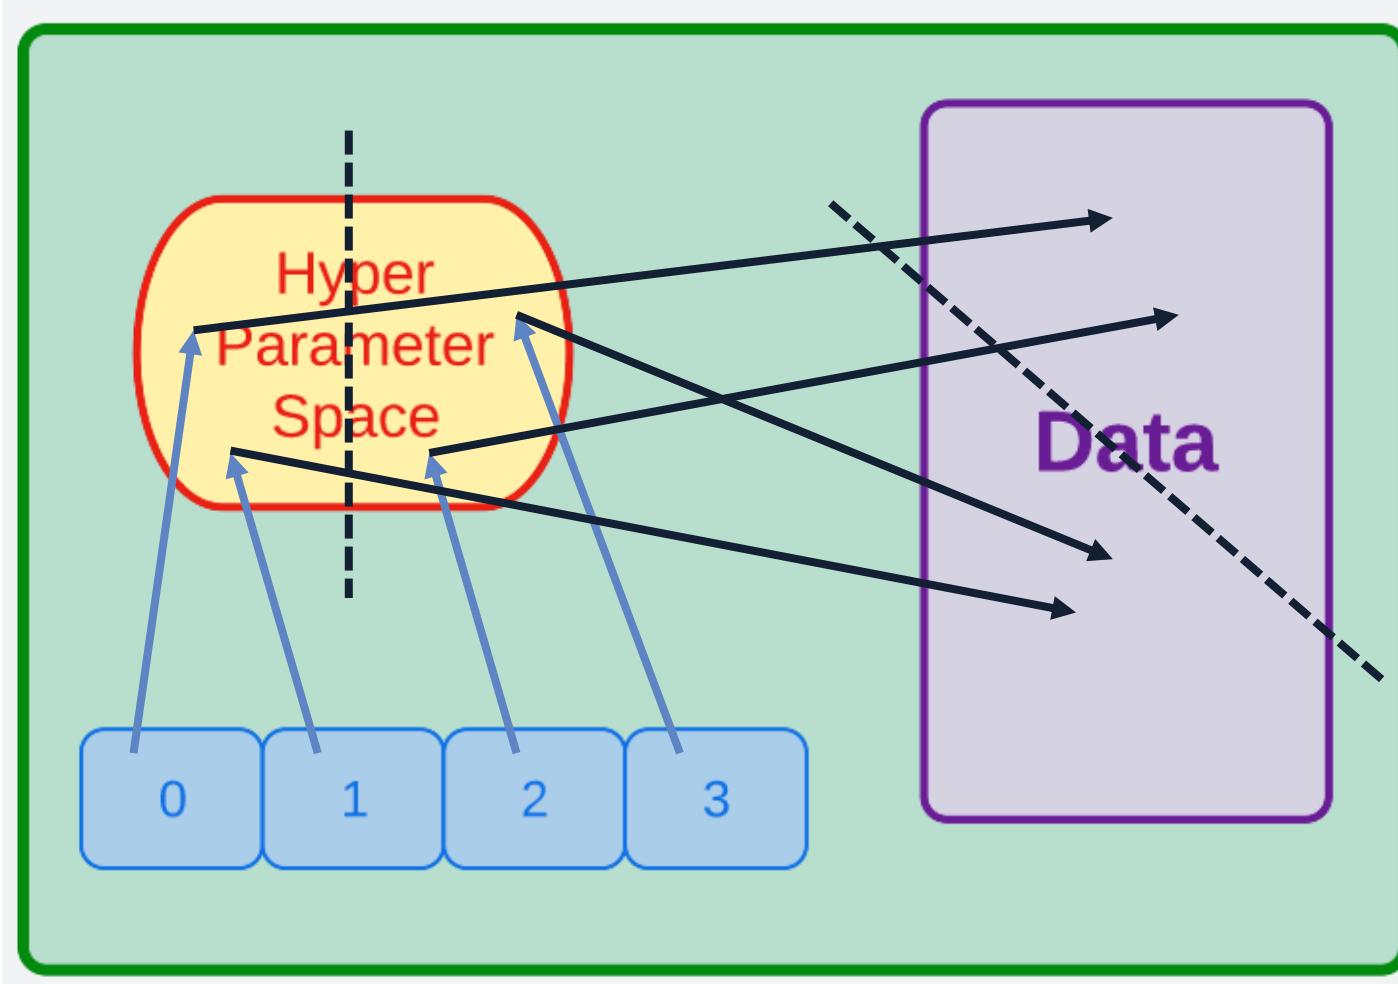
# Decomposition example



# Decomposition example



# Decomposition example



# Exercise 3 – MPI Linear Regression

1. Implement your domain decomposition using the MPI library functions
2. Check for consistency of your results
3. Perform a scaling analysis

# Outline

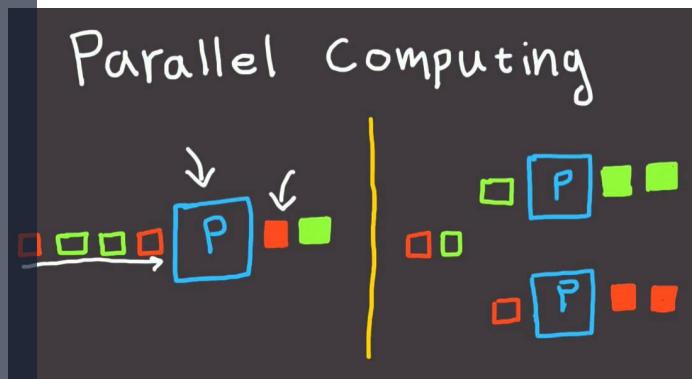
## Part A: Parallel Computing Quick Review (10 min)

- Task and Data parallelism
- Parallel computers
- Threads, processes and more



## Part B: MPI Intro (60 min)

- Core concepts
- Message passing
- Execution



## Part C: Exercises (45+ min)

Ex 1: Hello World

Ex 2: Domain Decomposition

Ex 3: Parallel Linear Regression

# Summary

Parallel computing is essential for modern scientific research computing

Computers today execute multiple instructions on multiple pieces of data (MIMD)

MPI allows for parallelization that can employ thousands of processes;  
redesigning an existing serial application is necessary

MPI collective communications are a good starting point; Point-to-point  
communication are often necessary for fully scalable parallelism



# Final assignment (optional)

Write a simple code that is related to the scientific research problem you are interested in (you may already have it!)

Line up the steps to parallelize it: perform domain decomposition, map the time of required communications

Parallelize it using MPI library functions

Share your results! Let's add it to the collaborative GitHub folder!

<https://github.com/babreu-ncsa/parallelization>



# Useful links

GitHub Parallelization repo

<https://github.com/babreu-ncsa/parallelization>

GitHub repo with this session's exercises

<https://github.com/babreu-ncsa/IntroToMPI>

HPC Training Moodle - Introduction to OpenMP course

<https://www.hpc-training.org/xsede/moodle/course/view.php?id=72>

MPI related contents

<https://www mpi-forum.org/docs/>

<https://mpitutorial.com/tutorials/>

Google Drive folder with this session's material

<https://drive.google.com/drive/folders/17WliPmA3vZ-PjXdiVlH1rx3WQtDCIZZV?usp=sharing>



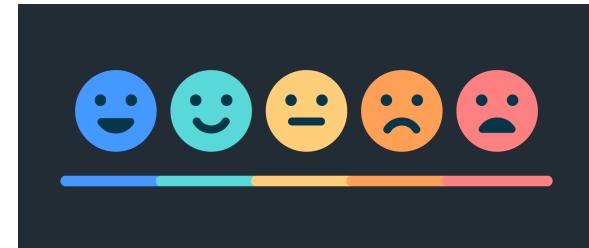
# Feed us back, please!

Feedback is vital for our continuous improvement as instructors/lecturers, and for our organization.

Please take a couple of minutes to answer this anonymous survey. Any sort of comments are extremely helpful and all of them are going to be carefully taken under consideration.

Thank you!

<https://go.ncsa.illinois.edu/IntroToMPIFeedback>





# Q&A

**Bruno Abreu, Ph.D.**

Research Scientist

[babreu@illinois.edu](mailto:babreu@illinois.edu)

[https://go.ncsa.illinois.edu/  
IntroToMPIFeedback](https://go.ncsa.illinois.edu/IntroToMPIFeedback)



National Center for  
Supercomputing Applications

UNIVERSITY OF ILLINOIS URBANA-CHAMPAIGN