

Linux tools for text processing

COMPLECS Team

San Diego Supercomputer Center

<https://bit.ly/COMPLECS>

<https://github.com/sdsc-complecs/Linux-text-tools>

About COMPLECS

COMPLECS (COMPrehensive Learning for end-users to Effectively utilize CyberinfraStructure) is a new SDSC program where training will cover non-programming skills needed to effectively use supercomputers. Topics include parallel computing concepts, Linux tools and bash scripting, security, batch computing, how to get help, data management and interactive computing.

*COMPLECS is supported by
NSF award 2320934.*



Table of contents

- Introduction
- head / tail
- paste
- nl
- sort (and shuf)
- split
- grep
- sed
- awk
- Case studies
- Using LLMs
- Conclusions

This session is intended for anyone who ...

... needs to post-process text files to extract run times, results or information about the hardware used for the calculations. The skills taught here will also be useful when constructing workflows, preparing high-throughput computing workloads or pre-processing data to get it into the correct format.

The Linux environment provides a wide range of tools for sorting, splitting and manipulating files along with tools for extracting content based on location (line number) or value.

In this sessions we'll cover the tools split, head, tail, awk, sed, grep and paste. Some of these tools have many options, so we'll just look at the basics to keep things more manageable.

Can't I just do this all by hand? Do I really need to learn yet another tool (or tools)?

While text manipulation tasks *could* be done by hand, the process can be

- Time-consuming
- Tedious
- Error prone (!)

This is especially true if they need to be done many times, such as extracting a value from the potentially thousands of files generated during a set of parameter-sweep calculations.

An obvious solution is to automate the process. Sometimes this requires the development of relatively complex parsers that are beyond the ability of non-programmers to write, but often simple Linux utilities are sufficient.

Hands on examples

The accompanying GitHub repo contains examples showing usage of each tool. You have the option of executing the Jupyter notebook to walk through each exercise or downloading and executing scripts from the command line.

The best way to master these tools is to try new things

- Deliberately introduce errors (e.g., leave out quotes, commas or semicolons) to see what happens
- Extend the examples to introduce new features
- Chain the tools together to create more complex workflows
- Apply the tools to your own use cases

Table of contents

- Introduction
- head / tail
- paste
- nl
- sort (and shuf)
- split
- grep
- sed
- awk
- Case studies
- Using LLMs
- Conclusions

head / tail – output first / last part of file

By default, head / tail output the first / last 10 lines of file

```
$ head file.txt
```

```
Line 1  
Line 2  
Line 3  
Line 4  
Line 5  
Line 6  
Line 7  
Line 8  
Line 9  
Line 10
```

```
$ tail file.txt
```

```
Line 11  
Line 12  
Line 13  
Line 14  
Line 15  
Line 16  
Line 17  
Line 18  
Line 19  
Line 20
```


head / tail – output first / last part of file

As expected, the user can override the default and specify number of output lines

```
$ head -n 7 file.txt  
Line 1  
Line 2  
Line 3  
Line 4  
Line 5  
Line 6  
Line 7
```

```
$ tail -n 7 file.txt  
Line 14  
Line 15  
Line 16  
Line 17  
Line 18  
Line 19  
Line 20
```

head / tail – output first / last part of file

Can also specify that all but the last NUM lines be output (head -n -NUM) or all lines from NUM onward are output (tail -n +NUM)

```
$ head -n -15 file.txt  
Line 1  
Line 2  
Line 3  
Line 4  
Line 5
```

```
$ tail -n +16 file.txt  
Line 16  
Line 17  
Line 18  
Line 19  
Line 20
```

Table of contents

- Introduction
- head / tail
- **paste**
- nl
- sort (and shuf)
- split
- grep
- sed
- awk
- Case studies
- Using LLMs
- Conclusions

paste – merge lines of files

The paste command combines two or more files line-by-line. By default, tab is used as delimiter, but this can be overridden using -d option. Much more convenient than cutting and pasting into a spreadsheet.

```
$ cat fruits.txt  
apple  
banana  
grape
```

```
$ cat colors.txt  
red  
yellow  
purple
```

```
$ paste fruits.txt colors.txt  
apple red  
banana yellow  
grape purple
```

```
$ cat fruits.txt  
apple  
banana  
grape
```

```
$ cat colors.txt  
red  
yellow  
purple
```

```
$ paste -d ',' fruits.txt colors.txt  
apple,red  
banana,yellow  
grape,purple
```

Table of contents

- Introduction
- head / tail
- paste
- nl
- sort (and shuf)
- split
- grep
- sed
- awk
- Case studies
- Using LLMs
- Conclusions

nl – add line numbers to a file

The nl command adds line numbers to a file. It's not an essential command since the same task could be done using other tools, but it is extremely convenient.

```
$ nl fruits.txt # default
  1 apple
  2 banana
  3 grape

$ nl -n ln fruits.txt # left justified line numbers
1   apple
2   banana
3   grape

# the following are more awkward ways of doing the same thing
$ seq `wc -l < fruits.txt` > x ; paste x fruits.txt
$ awk '{print NR, $0}' fruits.txt
```

Table of contents

- Introduction
- head / tail
- paste
- nl
- sort (and shuf)
- split
- grep
- sed
- awk
- Case studies
- Using LLMs
- Conclusions

sort – sort lines of text files

Without arguments, sort is done by lexicographical order. Be aware that the handling of uppercase and lowercase letters will depend on the LC_ALL environment variable

```
$ cat unsorted1.txt
Banana
apple
Grape
pear
apple
peach
Banana
Grape
orange
```

```
$ sort unsorted1.txt
apple
apple
Banana
Banana
Grape
Grape
orange
peach
pear
```

```
$ export LC_ALL=C
$ sort unsorted1.txt
Banana
Banana
Grape
Grape
apple
apple
orange
peach
pear
```


sort – sort lines of text files

Can also choose the field to be used for sorting using the -k option, with the -n option used to sort by numeric value. In the second example, sorting numerically is probably the desired behavior.

```
$ cat unsorted2.txt
banana 17  taco
apple   6  pizza
grape  23  hotdog
pear   35  pretzel
peach  12  popcorn
```

```
$ sort -k3 unsorted2.txt
grape  23  hotdog
apple   6  pizza
peach  12  popcorn
pear   35  pretzel
banana 17  taco
```

```
$ sort -k2 unsorted2.txt
peach  12  popcorn
banana 17  taco
grape  23  hotdog
pear   35  pretzel
apple   6  pizza
```

```
$ sort -k2 -n unsorted2.txt
apple   6  pizza
peach  12  popcorn
banana 17  taco
grape  23  hotdog
pear   35  pretzel
```

sort – sort lines of text files

Sort can be restricted to listing unique values by specifying -u option

```
$ cat unsorted1.txt  
Banana  
apple  
Grape  
pear  
apple  
peach  
Banana  
Grape  
orange
```

```
$ sort unsorted1.txt  
apple  
apple  
Banana  
Banana  
Grape  
Grape  
orange  
peach  
pear
```

```
$ sort -u unsorted1.txt  
apple  
Banana  
Grape  
orange  
peach  
pear
```

sort – sort lines of text files

Use multiple ‘-k’ options for multicolumn sorts. Typically use ‘-k[n],[n]’ to restrict sort to only the data in that column. Include ‘r’ and ‘n’ for reverse and numeric sorts.

```
$ sort -k3,3 -k1,1 -k2,2 unsorted3.txt
apple green Canada
apple red Canada
banana green Canada
banana red Canada
apple green Mexico
apple red Mexico
banana green Mexico
banana red Mexico
apple green USA
apple red USA
banana green USA
banana red USA
```

```
$ sort -k2r,2 -k1,1 -k3,3 unsorted3.txt
apple red Canada
apple red Mexico
apple red USA
banana red Canada
banana red Mexico
banana red USA
apple green Canada
apple green Mexico
apple green USA
banana green Canada
banana green Mexico
banana green USA
```

sort – sort lines of text files

Sort has an option to randomly shuffle the lines in a file but be careful. Because of the way it's implemented, identical lines will always appear next to each other. The Linux shuf utility is usually a better option. If you want a random permutation of integers, consider piping output from seq into shuf.

```
$ cat abc.txt  
a  
a  
a  
b  
b  
b  
c  
c  
c
```

```
$ sort -R abc.txt  
b  
b  
b  
c  
c  
c  
a  
a  
a
```

```
$ shuf abc.txt  
b  
a  
c  
a  
c  
b  
a  
b  
c
```

```
$ seq 9 | shuf  
9  
4  
1  
8  
6  
3  
7  
5  
2
```

Table of contents

- Introduction
- head / tail
- paste
- nl
- sort (and shuf)
- **split**
- grep
- sed
- awk
- Case studies
- Using LLMs
- Conclusions

split – split a file into pieces

By default, split breaks a file into chunks of 1000 lines, using naming convention xaa, xab, xac, ... Options available to specify different chunk sizes, prefixes and numbering

```
$ split genome.fasta
$ wc -l *
  9991 genome.fasta
   1000 xaa
   1000 xab
   1000 xac
   1000 xad
   1000 xae
   1000 xaf
   1000 xag
   1000 xah
   1000 xai
    991 xaj
 19982 total
```

If file cannot be divided evenly, the last output file gets whatever is left

```
$ $ split -l 2000 genome.fasta genome_
$ wc -l *
  2000 genome_aa
   2000 genome_ab
   2000 genome_ac
   2000 genome_ad
   1991 genome_ae
   9991 genome.fasta
 19982 total
```

split – split a file into pieces

For files that have a well-defined structure, such as FASTA files used for genomic data, the default record separator (newline character) may lead to undesired behavior. In this case, we want to keep the annotation and sequence together

```
$ head -1 genome_aa
>sp|Q61151|2A5E_MOUSE Serine/threonine-protein phosphatase 2A 56 ...

$ head -1 genome_ab
SYTTYMKEEVDRYRITIGNKTCVFEEKENDPSVMRSPSAGKLIQYIVEDGGHVFAGQCYAE

$ head -1 genome_ac
PLGGGREVWFGFHQSVRPAMWNMMLNIDVSATAFYRAQPIIEFMCEVLDIQNINEQTKPL
```

split – split a file into pieces

Fortunately, we can specify a different record separator using the -t option to control how the file is split. Be careful that delimiter only occurs where expected and note that it is not included in output (we'll show later how to fix this)

```
$ split -l 200 -t '>' genome.fasta genome_  
  
$ head -1 genome_aa  
>sp|Q61151|2A5E_MOUSE Serine/threonine-protein phosphatase 2A 56 ...  
  
$ head -1 genome_ab  
sp|O35674|ADA19_MOUSE Disintegrin and metalloproteinase domain ...  
  
$ head -1 genome_ac  
sp|Q91WF3|ADCY4_MOUSE Adenylate cyclase type 4 OS=Mus musculus ...
```


Table of contents

- Introduction
- head / tail
- paste
- nl
- sort (and shuf)
- split
- **grep**
- sed
- awk
- Case studies
- Using LLMs
- Conclusions

grep – print lines that match a pattern

At its simplest, grep returns all lines in a file containing the search string. Upper and lowercase letters are different characters, and you'll need the -i flag if you want your matches to be case insensitive

```
$ cat file1.txt
apple banana pear
cherry banana peach
orange lemon lime
apple Banana pear
cherry Banana peach
pear kunquat grape
grape guava lime
pear peach kiwi
```

```
$ grep banana file1.txt
apple banana pear
cherry banana peach
```

```
$ grep Banana file1.txt
apple Banana pear
cherry Banana peach
```

```
$ grep -i banana file1.txt
apple banana pear
cherry banana peach
apple Banana pear
cherry Banana peach
```

grep – print lines that match a pattern

grep provides options for listing lines that occur before or after the matching line. This is extremely useful when the desired content can vary but occurs in a known location relative to matching pattern.

```
$ grep -B 3 kumquat file1.txt # print three lines before matching line
orange lemon lime
apple Banana pear
cherry Banana peach
pear kumquat grape

$ grep -A 2 kumquat file1.txt # print two lines after matching line
pear kumquat grape
grape guava lime
pear peach kiwi
```

grep – print lines that match a pattern

grep can do an inverted match to print lines that do not match the pattern

```
$ grep lime * # contain lime
file1.txt:orange lemon lime
file1.txt:grape guava lime
file2.txt:orange lemon lime
file2.txt:pear lime grape
file2.txt:grape guava lime
```

```
$ grep -v lime * # do not contain lime
file1.txt:apple banana pear
file1.txt:cherry banana peach
file1.txt:apple Banana pear
file1.txt:cherry Banana peach
file1.txt:pear kumquat grape
file1.txt:pear peach kiwi
file2.txt:apple banana pear
file2.txt:cherry banana peach
file2.txt:apple Banana pear
file2.txt:cherry Banana peach
file2.txt:pear peach kiwi
```

grep – print lines that match a pattern

Other options include printing the names of files that contain or don't contain the match and listing counts of matching lines

```
$ grep lime *
file1.txt:orange lemon lime
file1.txt:grape guava lime
file2.txt:orange lemon lime
file2.txt:pear lime grape
file2.txt:grape guava lime

$ grep -c lime * # number of matches
file1.txt:2
file2.txt:3
```

```
$ grep -l kumquat * # contain match
file1.txt

$ grep -L kumquat * # don't contain match
file2.txt
```

grep – print lines that match a pattern

Up to this point, our pattern has been a string literal (e.g., lime or banana), but grep can recognize more complex patterns that use character classes, anchors and even regular expressions

```
$ grep pear file1.txt      # All lines containing pear
apple banana pear
apple Banana pear
pear kumquat grape
pear peach kiwi

$ grep '^pear' file1.txt  # Lines starting with pear (uses ^ anchor)
pear kumquat grape
pear peach kiwi


$ grep 'pear$' file1.txt  # Lines ending with pear (uses $ anchor)
apple banana pear
apple Banana pear
```

Table of contents

- Introduction
- head / tail
- paste
- nl
- sort (and shuf)
- split
- grep
- **sed**
- awk
- Case studies
- Using LLMs
- Conclusions

sed – stream editor for filtering and transforming text

sed is an extremely powerful editor that operates on a stream of data. We could easily spend days exploring sed and entire books have been written about it, but we'll just look at simple examples demonstrating two commonly used features.



<https://www.gnu.org/software/sed/manual/sed.html>

<http://sed.sourceforge.net/>

<http://sed.sourceforge.net/grabbag/tutorials/>

sed – stream editor for filtering and transforming text

sed can be used to apply a substitution to each line of a file. While *some* of these operations can easily be done using a global replace in a traditional editor (vi, emacs), sed has the advantage that it can be scripted. In addition, since it works on a stream of data, the entire file does not need to fit into memory.

```
sed 's/replace-first-occurrence-of-this/with-this/' [input-file]  
sed 's/replace-every-occurrence-of-this/with-this/g' [input-file]
```

sed – stream editor for filtering and transforming text

```
# First instance in each line
$ sed s'/pear/XXXX/' file1.txt
XXXX apple banana grape
apple banana grape XXXX
XXXX apple pear grape
apple XXXX banana pear
XXXX apple banana pear
```

```
# All instances in each line
$ sed s'/pear/XXXX/g' file1.txt
XXXX apple banana grape
apple banana grape XXXX
XXXX apple XXXX grape
apple XXXX banana XXXX
XXXX apple banana XXXX
```

```
# Instances at start of each line
$ sed s'^pear/XXXX/' file1.txt
XXXX apple banana grape
apple banana grape pear
XXXX apple pear grape
apple pear banana pear
XXXX apple banana pear
```

```
# Instances at end of each line
$ sed s'/pear$/XXXX/' file1.txt
pear apple banana grape
apple banana grape XXXX
pear apple pear grape
apple pear banana XXXX
pear apple banana XXXX
```

sed – stream editor for filtering and transforming text

sed can also print out specific lines. There are a lot of options, and we'll just focus on a few things that aren't easy to do using other tools

```
$ sed -n '5'p file2.txt    # print fifth line  
This is line 5
```

```
$ sed -n '3,5'p file2.txt # print lines 3-5  
This is line 3  
This is line 4  
This is line 5
```

```
$ sed -n '2~3'p file2.txt # print every third line starting at line 2  
This is line 2  
This is line 5  
This is line 8
```

sed – stream editor for filtering and transforming text

Like all other commands that write to stdout, sed output can be directed to a file. There's just one gotcha – since sed streams through the file rather than storing in memory, redirecting to original file doesn't work. Write to temp file, or even better use the -i option to edit in place

```
$ head -2 file1.txt > file3.txt
$ sed s'/pear/XXXX/' file3.txt > file3.txt
$ cat file3.txt
[no output]


$ head -2 file1.txt > file3.txt
$ sed -i s'/pear/XXXX/' file3.txt
$ cat file3.txt
XXXX apple banana grape
apple banana grape XXXX
```

Table of contents

- Introduction
- head / tail
- paste
- nl
- sort (and shuf)
- split
- grep
- sed
- **awk**
- Case studies
- Using LLMs
- Conclusions

awk – pattern scanning and processing language

awk is a powerful text processing tool named after its three creators Alfred Aho, Peter Weinberger, and Brian Kernighan. There are a lot of capabilities, and we'll limit ourselves to a few of the simple and commonly used ones.



<https://www.gnu.org/software/gawk/manual/gawk.html>

https://en.wikibooks.org/wiki/An_Awk_Primer/Awk_Command-Line_Examples

awk – pattern scanning and processing language

awk programs consist of one or more *rules* containing an *option* and/or an *action**

```
pattern {action}
```

Short programs can be executed from the command line as follows

```
awk 'pattern {action}' file [file2]
```

Multiple actions can be separated by semicolons and there are special actions that are executed once at the start and once at the end of execution

```
awk 'BEGIN {action} pattern {action1; action 2} END {action}' file1 [file2]
```

** awk is a full-fledged programming language with arrays, control statements and other advanced features. We're going to ignore all that for now.*

awk – pattern scanning and processing language

The most common use of awk is to print selected fields from each line of a file. Awk processes one line at a time, splitting into fields referenced by position (\$1, \$2, ...). The last field can also be accessed as \$NF, where NF is the number of fields.

```
$ cat people.txt
Bob dog meatloaf banana blue
Cindy hamster pizza kiwi purple
Mary cat sushi apple green
Susan fish pizza grape yellow

$ awk '{print $1, $3, $NF}' people.txt
Bob meatloaf blue
Cindy pizza purple
Mary sushi green
Susan pizza yellow
```


awk – pattern scanning and processing language

Using awk's predefined NF variable is particularly useful if you don't know the number of fields or are working with ragged data (each line has different field count).

```
$ cat ragged.txt
field1 field2 field3
field1 field2
field1 field2 field3 field4
field1
field1 field2 field3 field4 field5

awk '{print $NF}' ragged.txt
field3
field2
field4
field1
field5
```

awk – pattern scanning and processing language

awk can also take a search pattern and include other text in the output

```
$ awk '{print $1 " favorite food is " $3}' people.txt
Bob favorite food is meatloaf
Cindy favorite food is pizza
Mary favorite food is sushi
Susan favorite food is pizza

$ awk '/pizza/ {print $1, $3}' people.txt # limit to lines matching pizza
Cindy pizza
Susan pizza
```

awk – pattern scanning and processing language

By default, awk uses whitespace as the delimiter, but this can be changed using the -F option. This is useful when working with comma separated value (CSV) files

```
$ cat states.txt
California,Sacramento,Valley Quail,Poppy
Texas,Austin,Mockingbird,Bluebonnet
Florida,Tallahassee,Mockingbird,Orange Blossom
New York,Albany,Eastern Bluebird,Rose
Pennsylvania,Harrisburgh,Ruffed Grouse,Mountain Laurel

$ awk -F ',' '{print $1 " state bird is " $3}' states.txt
California state bird is Valley Quail
Texas state bird is Mockingbird
Florida state bird is Mockingbird
New York state bird is Eastern Bluebird
Pennsylvania state bird is Ruffed Grouse
```

awk – pattern scanning and processing language

awk provides basic math capabilities: square root, log, exp, trig functions and random numbers, etc. Note that multiple actions are separated by semicolons

```
$ awk '{sum=$1+$2; prod=$1*$2; print $1, $2, sum, prod}' numbers.txt
1.1 2.2 3.3 2.42
3.3 4.4 7.7 14.52
```

```
$ awk '{sin1=sin($1); cos2=cos($2); print $1, $2, sin1, cos2}' numbers.txt
1.1 2.2 0.891207 -0.588501
3.3 4.4 -0.157746 -0.307333
```

```
$ awk '{log1=log($1); sqrt2=sqrt($2); print $1, $2, log1, sqrt2}' numbers.txt
1.1 2.2 0.0953102 1.48324
3.3 4.4 1.19392 2.09762
```

awk – pattern scanning and processing language

1. By default, awk applies operations to each line. BEGIN and END allow you to execute additional code just once at the start or end of the file processing
2. Commands can be broken across multiple lines
3. C-style formatted printing available using printf, explicitly need newline (\n)

```
$ $ awk 'BEGIN {PI=3.14159; print "  x      y      sin(x)  cos(y)"}
> {sin1=sin($1*PI); cos2=cos($2*PI);
> printf "%4.1f %4.1f %7.4f %7.4f\n",
> $1, $2, sin1, cos2}' numbers.txt
  x      y      sin(x)  cos(y)
1.1    2.2 -0.3090    0.8090
3.3    4.4 -0.8090    0.3090
```

awk – pattern scanning and processing language

An especially useful application of awk is to calculate sums, averages and other quantities over columns of numbers*

```
$ awk '{s1+=$1; s2+=$2; print} END {print s1/NR, s2/NR}' numbers.txt
1.1 2.2
3.3 4.4
7.7 5.1
5.5 0.4
4.4 3.025
```

Notes: NR is awk predefined variable containing number of lines; print without any arguments prints the entire line

** Much easier and more convenient than cutting and pasting into a spreadsheet*

Table of contents

- Introduction
- head / tail
- paste
- nl
- sort (and shuf)
- split
- grep
- sed
- awk
- Case studies
- Using LLMs
- Conclusions

Case studies

Most of our examples have been simple and maybe even a little contrived, involving things like lists of fruits. In the remainder of this tutorial, we give some more complex real-life examples illustrating how multiple tools can be used together.

This would be a good time to think about your own workloads and how you could use these tools to simplify, automate or streamline your research

Case study – High Performance LINPACK (HPL)

When deploying a new supercomputer, we run thousands of HPL jobs to ensure that all nodes are functioning correctly and delivering expected performance. This requires extracting node name, run time and completion code from the output files.

```
vgr-10-01
...
=====
T/V              N      NB      P      Q              Time              Gflops
-----
WR12R2R4         207800   384      6      8             4386.47             1.3638e+03
HPL_pdgesv() start time Wed Apr 13 21:55:29 2022

HPL_pdgesv() end time   Wed Apr 13 23:08:35 2022
...
-----
||Ax-b||_oo/(eps*(||A||_oo*||x||_oo+||b||_oo)*N)= 1.00164594e-03 ..... PASSED
=====
...
```

Case study – High Performance LINPACK (HPL)

Bash script illustrating iteration over files, backticks to capture command output and use of multiple utilities (head, grep, awk, wc) to extract required info from output files

```
echo "node      time      P F"          # Print header
for file in `ls output.hpl*`          # Iterate over output files
do
    node=`head -1 $file`              # Node name from first line
    t=`grep WR12R2R4 $file | awk '{print $6}'` # Time from 6th field of WR12R2R4 line
    npass=`grep PASSED $file| wc -l`    # Number of passed tests
    nfail=`grep FAILED $file| wc -l`    # Number of failed tests
    echo $node $t $nfail $npass
done
```

HPL output taken from real machine testing, with failure conditions added for illustrative purposes only.

Case study – High Performance LINPACK (HPL)

Bash script illustrating iteration over files, backticks to capture command output and use of multiple utilities (head, grep, awk, wc) to extract required info from output files

```
echo 'node      time      P F'
for file in `ls output.hpl*`
do
  node=`head -1 $file`
  t=`grep WR12R2R4 $file | awk '{print $6}'`
  npass=`grep PASSED $file| wc -l`
  nfail=`grep FAILED $file| wc -l`
  echo $node $t $nfail $npass
done
```



```
$ ./process-hpl.sh
node      time      P F
vgr-10-30 4729.98 0 1
vgr-10-24 4751.87 0 1
vgr-10-11 4824.54 0 1
vgr-10-11 4808.07 1 0
vgr-10-11 4804.13 0 1
vgr-10-08 4883.66 0 1
vgr-10-37 4851.29 0 1
vgr-10-07 4697.88 1 0
```

Case study – BERT natural language processing

In a recent deployment of specialized AI hardware, we ran BERT benchmarks to ensure that all eight of the cards in each node were working correctly and delivering consistent performance. Required extracting node name, card ID, timestamps, percentage correct answers and achieved rate from output files

```
vgr-10-02
...
[1] AIP (hl1) 0000:1b:00.0
...
TBEFORE 1646939242
...
Correct answers: 1466 out of 1724 = 85.03 %
Elapsed time (sec) including enqueue: total - 1.0239 per sample - 0.000593907
Achieved sentences/sec: 1683.77
...
TAFTER 1646940484
```

Case study – BERT natural language processing

Bash script illustrating iteration over files, backticks to capture command output and use of multiple utilities (head, tail, grep, awk, wc) to extract required info

```
for file in `ls -1 out*`
do
  node=`head -1 $file | awk '{print $1}'`
  card=`grep 'AIP (' $file | awk '{print $4}'`
  tstart=`grep TBEFORE $file | awk '{print $2}'`
  tend=`grep AFTER $file | awk '{print $2}'`
  corr=`grep 'Correct answers' $file | tail -1 | awk '{print $8}'` # Last instance
  ach=`grep 'Achieved sentences' $file | tail -1 | awk '{print $3}'` # Last instance
  t=$((expr $tend - $tstart)) # Do math to calculate tend - tstart
  echo $node $card $corr $ach $t
done
```

Case study – BERT natural language processing

Bash script illustrating iteration over files, backticks, shell arithmetic and use of multiple utilities (head, tail, grep, awk, wc) to extract required info

```
for file in `ls -l out*`
do
  node=`head -1 $file | awk '{print $1}'`
  card=`grep 'AIP (' $file | awk '{print $4}'`
  tstart=`grep TBEFORE $file | awk '{print $2}'`
  tend=`grep AFTER $file | awk '{print $2}'`
  corr=`grep 'Correct answers' $file | tail -1 | awk '{print $8}'` # Last instance
  ach=`grep 'Achieved sentences' $file | tail -1 | awk '{print $3}'` # Last instance
  t=$((expr $tend - $tstart)) # Do math to calculate tend - tstart
  echo $node $card $corr $ach $t
done
```

```
$ ./process-bert.sh
vgr-10-02 0000:1b:00.0 85.03 1683.77 1242
vgr-10-01 0000:b1:00.0 85.03 1656.87 1342
vgr-10-01 0000:89:00.0 85.03 1687.34 1357
```

Case study – correctly splitting FASTA file

In the original example where we split a FASTA file, we noted that the delimiter was lost in the first annotation line. We can clean this up by using sed to modify lines that begin with 'sp|', which occurs only in annotations, to instead begin with '>sp|'

```
split -l 2000 -t '>' genome.fasta genome_  
for file in genome_*  
do  
    sed 's/^sp|/>sp|/' $file > temp  
    mv temp $file  
done
```

```
$ head -1q genome_*  
>sp|Q61151|2A5E_MOUSE Serine/threonine-protein phosphatase 2A 56 kDa ...  
>sp|O35674|ADA19_MOUSE Disintegrin and metalloproteinase domain-containing ...  
>sp|Q91WF3|ADCY4_MOUSE Adenylate cyclase type 4 OS=Mus musculus GN=Adcy4 PE=1 ...  
>sp|Q8K209|AGRG1_MOUSE Adhesion G-protein coupled receptor G1 OS=Mus musculus ...  
>sp|Q68G58|APEX2_MOUSE DNA-(apurinic or apyrimidinic site) lyase 2 OS=Mus ...
```

Table of contents

- Introduction
- head / tail
- paste
- nl
- sort (and shuf)
- split
- grep
- sed
- awk
- Case studies
- **Using LLMs**
- Conclusions

Using large language models (LLMs)

Large language models can help, but it's still valuable to know a bit about these tools and their capabilities. Without some basic knowledge, it's hard to know what question to ask.



You

Can you give me an awk one-liner to print the second and fourth fields from a file?



ChatGPT

Certainly! Here's an `awk` one-liner to print the second and fourth fields from a file:

bash

Copy code

```
awk '{print $2, $4}' yourfile.txt
```

Replace `'yourfile.txt'` with the actual name of your file. This command will print the



You

Thanks, now can you add the condition that this is only done for lines where the second field is greater than the fourth field?



ChatGPT

Certainly! Here's the modified `awk` one-liner with the added condition:

bash

Copy code

```
awk '$2 > $4 {print $2, $4}' yourfile.txt
```

This command now checks if the second field is greater than the fourth field before printing the values. Replace `'yourfile.txt'` with the actual name of your file. Only lines meeting the specified condition will be printed.

Using large language models (LLMs)

Large language models can help, but it's still valuable to know a bit about these tools and their capabilities. Without some basic knowledge, it's hard to know what question to ask.



You

Can you give me a Linux command to sort a file by the second field numerically and then by the first field in reverse alphabetical order?



ChatGPT

Certainly! You can use the `sort` command with the `-k` option to achieve this.

Here's the command:

bash

Copy code

```
sort -t ' ' -k2,2n -k1,1r yourfile.txt
```

Explanation:

- `-t ' '` specifies that the fields are separated by a space.
- `-k2,2n` sorts by the second field numerically.
- `-k1,1r` sorts by the first field in reverse alphabetical order.

Replace `yourfile.txt` with the actual name of your file.



Table of contents

- Introduction
- head / tail
- paste
- nl
- sort (and shuf)
- split
- grep
- sed
- awk
- Case studies
- Using LLMs
- **Conclusions**

Conclusions

- Linux provides a powerful suite of tools for text manipulation. They can improve your productivity AND reduce the likelihood of making errors.
- Don't be discouraged if you didn't understand everything. We covered a lot of material in a pretty short time. Work through the examples at your own pace until you feel like you've truly mastered the content.
- We've barely scratched the surface, especially for the more complex tools such as awk and sed. While you can go much deeper, we probably covered most of what you need to know.
- Large language models (LLMs) can help, especially for the more complex tools. We recommend though that you at least know the basics so that you can formulate meaningful prompts